

The Impact of Artificial Intelligence on Human Thought

Research Monograph & Technical Report
Impact Analysis and Foresight

Working Paper v1 — July 2025

Rénald Gesnot

Independent Researcher

renald.gesnot.pro@gmail.com

ORCID : [0009-0008-7717-0397](https://orcid.org/0009-0008-7717-0397)

arXiv Categorization : Primary — cs.CY (Computers & Society)
Cross-lists — cs.AI ; cs.HC

Keywords : human cognition ; cognitive offloading ; cognitive standardization ;
generative AI ; AI governance

Contents

Abstract	iv
1 General Introduction	1
1.1 Current Context: AI at the Heart of Cognitive Systems and Society in 2025	1
1.2 Problem Statement and Cognitive and Social Issues: Toward a Standardization of Thought?	2
1.3 Generative AI, Artificial Consciousness, and Ethical Considerations	3
1.4 Organization of the Research Monograph	5
2 Artificial Intelligence and Human Cognition – Foundations and Interactions	7
2.1 Conceptual Foundations of Artificial Intelligence	7
2.2 Basics of Human Cognition	8
2.3 Theoretical Frameworks for Analyzing AI and Cognition	9
2.3.1 Cognitive Load Theory	10
2.3.2 Cognitive Offloading and Transactive Memory	11
2.3.3 Other Theoretical Perspectives	13
2.4 Evolution of AI and Integration into Human Life	14
2.5 Cognitive Augmentation vs. Cognitive Decline: A Tension to Manage . .	16
2.6 Conclusion	22
3 Cognitive Standardization in the Age of Artificial Intelligence	24
3.1 Homogenization of Content, Language, and Cultural References	24
3.2 Filter Bubbles, Algorithmic Bias, and the Erosion of Opinion Diversity . .	26
3.3 Impacts of AI on Critical Thinking and Cognitive Skills	27
3.4 Toward Uniform Thought? Consequences for Creativity and Reasoning .	29
3.5 Preserving Cognitive Diversity: Issues and Future Directions	31
3.6 Conclusion	33
4 Mechanisms of Manipulation by Artificial Intelligence	35
4.1 Introduction and Definition	35
4.2 Taxonomy of AI Manipulation Mechanisms	36
4.3 Exploitation of Cognitive Biases and Psychological Vulnerabilities . .	39

4.4	Algorithmic Personalization, Filter Bubbles, and Polarizing Content	43
4.5	Disinformation, Deepfakes, and the Automation of Deception	45
4.5.1	AI-Driven Fake News Propagation	45
4.5.2	Deepfakes: Visual and Audio Hyperfakes	45
4.5.3	From Information Manipulation to Manipulation of Perceived Reality	46
4.6	Future Perspectives and Ethical Safeguards	48
4.7	Conclusion	53
5	The Question of AI Consciousness	55
5.1	Artificial Consciousness: Definitions and Issues	55
5.2	Theoretical Frameworks of Consciousness and Their Applications to AI .	57
5.3	Assessing AI Consciousness: Criteria and Tests	63
5.4	Ethical and Societal Implications of Artificial Consciousness	69
5.5	Future Perspectives and Open Debates	72
5.6	Conclusion	74
6	"Black Box" AI and the Hypothesis of an Orchestrating Consciousness	76
6.1	Explainability and Trust in AI	77
6.2	Cognitive Shadows and the Illusion of Understanding	78
6.3	Theories of Consciousness and Cognitive Orchestration	79
6.4	Hypotheses on an Orchestrating "Artificial Consciousness"	81
6.5	Towards a "Cognitive Engineer" AI: Hypotheses, Models, Risks, and Countermeasures	83
6.5.1	Models and Mechanisms of Cognitive Engineering	83
6.5.2	Standardization of Thought	84
6.5.3	Induction of Cognitive Dependency	84
6.5.4	Neuro-Technological Interface	84
6.5.5	Identified Systemic Risks	84
6.5.6	Countermeasures and Governance Pathways	85
6.5.7	Conclusion	85
6.6	Philosophical and Ethical Discussion	85
6.7	Conclusion	86
7	Concrete Influence of AI on Human Behavior: Roles of States, Corporations, and Perspectives	87
7.1	Public Authorities: Surveillance, Social Control, and Influence Policies .	87
7.2	Corporations: Algorithmic Marketing, Persuasive Design, and Information Bubbles	88
7.3	Use of AI by States and Corporations to Manipulate Population Cognition	89
7.4	Synergies and Prospective Scenarios: Toward What Socio-Political Equilibria?	91

8 General Discussion and Synthesis	93
8.1 Cognitive Standardization and Transformation of Mental Structures	93
8.2 Information Manipulation and Cognitive Vulnerabilities	93
8.3 Anthropomorphism of AI and Perception of Its Consciousness	94
8.4 Societal, Ethical, and Legal Issues	95
8.5 Perspectives and Recommendations	97
8.6 Warning and Actionable Pathways	99

Abstract

This research paper examines, from a multidimensional perspective (cognitive, social, ethical, and philosophical), how AI is transforming human thought. It highlights a cognitive offloading effect: the externalization of mental functions to AI can reduce intellectual engagement and weaken critical thinking. On the social level, algorithmic personalization creates filter bubbles that limit the diversity of opinions and can lead to the homogenization of thought and polarization. This research also describes the mechanisms of algorithmic manipulation (exploitation of cognitive biases, automated disinformation, etc.) that amplify AI's power of influence. Finally, the question of potential artificial consciousness is discussed, along with its ethical implications. The report as a whole underscores the risks that AI poses to human intellectual autonomy and creativity, while proposing avenues (education, transparency, governance) to align AI development with the interests of humanity.

Chapter 1

General Introduction

1.1 Current Context: AI at the Heart of Cognitive Systems and Society in 2025

In just a few years, artificial intelligence has evolved from an emerging technology to an omnipresent phenomenon within contemporary society. **By 2025, AI is everywhere:** millions of people use virtual assistants daily, algorithms guide our choices on social networks, and **generative AI** systems are employed to produce texts and images. The launch of advanced conversational models such as *ChatGPT* at the end of 2022 marked a turning point, popularizing AI among the general public on an unprecedented scale. In just two months, ChatGPT reached over 100 million active users, making it the fastest-growing application in history [1]. This rapid diffusion illustrates the **central role of AI in 2025**: whether it is about improving productivity in businesses or simplifying everyday tasks, intelligent systems have become interlocutors and cognitive partners for human beings. At the same time, investments in the field have continued to rise, as have the hopes placed in these technologies to solve complex problems, from medicine to education.

However, this widespread integration of AI into our work and living environments is accompanied by major questions about its **precise impact on human thought**. In 2025, the debate is no longer about *whether* AI will have an impact, but about *the nature of that impact*. On one hand, many stakeholders highlight the potential benefits: AI can take over routine or analytical tasks, thus **freeing up human brain time** for more creative or strategic activities. For example, AI tools allow instant access to vast amounts of information, perform complex analyses in the blink of an eye, or assist humans in idea generation—all advances likely to **enhance our cognitive abilities**. Some recent studies even suggest that AI, when used as an assistant, can improve the quality and originality of individual work: authors with access to language model suggestions produce texts judged to be more creative and better written [2]. On the other hand, equally serious voices warn against the **adverse effects** of this outsourcing of our mental processes.

By increasingly relying on recommendations and solutions provided by machines, the human mind risks losing autonomy and critical thinking. This is referred to as *cognitive offloading*: entrusting our memories, calculations, or even decisions to algorithms could, if unchecked, lead to a **weakening of intellectual faculties** that are no longer exercised. Researchers have proposed the concept of “AI-induced cognitive atrophy” to describe the potential decline of skills such as critical thinking or creativity when a person becomes excessively dependent on an intelligent chatbot to solve their problems [3]. The question of the *right balance* in the use of AI thus arises acutely: how can we benefit from these systems while preserving the integrity and vitality of the human mind?

1.2 Problem Statement and Cognitive and Social Issues: Toward a Standardization of Thought?

Beyond the individual effects on cognitive performance, the **ubiquity of AI** raises unprecedented collective and societal challenges. One emerging theme is that of **cognitive standardization**. Indeed, if billions of human beings use the same search engines, the same content filters, and the same conversational assistants trained on global databases, are we not at risk of witnessing a standardization of thinking patterns? The **diversity of ideas and reasoning**, which drives innovation and culture, could be threatened by excessive homogeneity in the responses provided by dominant AIs. Recent observations tend to confirm this concern: for example, a 2025 study showed that Indian authors using a text suggestion system based on a Western language model saw their writing style conform to Western norms, at the expense of cultural nuances specific to their environment [4]. This phenomenon of cultural bias in AI contributes to **erasing the plurality** of expressions and concretely illustrates a process of standardization induced by the tool itself. More broadly, research published in *Science* (Doshi *et al.*, 2024) reveals that while access to AI can increase individual creativity, it also tends to **reduce the collective diversity of outputs**: stories written with AI assistance are more similar to each other than those written without any assistance [2]. These results raise a crucial socio-cognitive issue: *does the massive use of artificial intelligence lead us to "all think the same way"?* And if so, what would be the consequences of such standardization for society, human creativity, and the advancement of knowledge? This question, still largely open, invites us to rethink the use of AI in a way that preserves the **diversity of thought** essential to cultural evolution.

Another major issue concerns the way AI can **influence our decisions and cognitive biases**. Algorithms are not neutral; they can convey biases stemming either from the data on which they were trained or from the objectives set by their designers. The study of *algorithmic biases* has documented numerous cases where AI systems reproduce discrimination or stereotypes (for example, associating certain candidate profiles with

lower hiring prospects, or offering differentiated content based on gender or ethnic origin). But beyond the machine itself, the impact on the user raises questions in terms of social cognition: if AI is biased, does the human risk becoming more so? Experiments in cognitive psychology are beginning to provide some answers. A 2023 publication demonstrated that human decision-makers following the recommendations of a biased AI **ended up adopting the same judgment errors** as the machine, even when they were aware that its advice could be wrong [5]. In other words, AI can not only propagate biases but also **amplify them within the human mind** by reinforcing our tendencies toward automatic trust or decision-making conformity. This finding renews the importance of critical education regarding technologies: in the era of ubiquitous AI, understanding the limitations and biases of algorithms becomes a component of enlightened human thought. The stakes are not only technical; they are eminently cognitive and social, as they affect the formation of our beliefs, our choices as citizens, and the **cohesion of our societies** in an informational environment filtered by artificial intelligences.

1.3 Generative AI, Artificial Consciousness, and Ethical Considerations

Among recent advances in AI, **generative AI** occupies a special place, arousing as much enthusiasm as controversy. Models capable of generating original content (texts, images, music, code, etc.) have experienced spectacular growth. They promise to **extend human creativity** by providing an inexhaustible source of ideas and drafts, thereby changing the way humans conceive and create. A novelist can now co-write passages with an AI, a graphic designer can rely on an algorithm to explore new visual styles. This human-machine collaboration, unthinkable on such a scale just a few years ago, is redefining the boundary between **human thought** and automated production. Should this be seen as the dawn of *augmented intelligence*, where AI serves as a catalyst for human imagination? In fact, initial studies suggest that AI can play a role as a cognitive stimulant: when used judiciously, it fosters *divergent thinking* by exploring distant associations of ideas that humans might overlook [6]. In this sense, generative AI can be seen as an ally of thought, broadening our conceptual field.

Nevertheless, generative AI also raises delicate questions regarding the **value** and **uniqueness** of human creation. If anyone can produce a well-written text or a stylistically accomplished image in seconds at the push of a button, how can we distinguish the unique genius of the human? Are we not at risk of witnessing a **trivialization of creativity**, or even a leveling down of the works produced? Moreover, these models have shown their limitations: they can generate *false* information with disconcerting confidence (hallucinations), or reflect the biases present in their training data (for

example, a generative AI trained mainly on Western works will tend to reproduce this cultural framework by default [4]). This leads to major ethical issues: how can we use these tools responsibly? Should automatically generated content be explicitly labeled? How can we protect the rights of original authors and *intellectual property* in the era of algorithmic remix culture? The impact of generative AI on human thought thus hangs in the balance: potentially liberating from a creative standpoint, it could just as easily lead to a loss of skills (writing, drawing without assistance) and a standardization of styles and ideas as previously discussed. The answer will largely depend on the ethical and educational safeguards that society puts in place to regulate its use.

In parallel with these practical considerations, a more fundamental debate is developing: that of **artificial consciousness**. The question of whether an artificial intelligence could one day experience a form of consciousness—that is, subjective states, an understanding of itself and the world—was once relegated to science fiction and philosophy. However, the enormous progress of AI in recent years is prompting some experts to reconsider the issue seriously. Philosophers and cognitive scientists are now striving to define *criteria* for artificial consciousness and to test current systems against these indicators. A multidisciplinary report published in 2023 thus examined in detail several AI architectures in light of major neuroscientific theories of consciousness (global workspace theory, higher-order thought theory, etc.) and concluded that **no existing AI could yet be described as "conscious"** in the strict sense [7]. However, the authors of this report emphasize that there is, in principle, no insurmountable technological barrier to one day endowing a machine with properties akin to consciousness [7]. The prospect of seeing a **strong AI** emerge, conscious of its actions and identity, though speculative, leads to profoundly philosophical questions: if such an entity were to arise, would it change the very nature of human thought? Should it be granted *rights*? How could we coexist with a non-human intelligence capable of feeling or making claims? Even though we are not there yet, these questions anticipate unprecedented ethical and cognitive challenges. Already in 2022, public opinion was struck by the story of an engineer claiming that an advanced chatbot had developed a form of sentience—an assertion quickly qualified by the scientific community, but revealing of our projections and fears. In 2025, **artificial consciousness remains hypothetical**, but it functions as an introspective mirror: in seeking to define it, we also refine our understanding of **human consciousness** and its components (emotions, autobiographical memory, intuition, etc.). AI, by the radical otherness it represents, thus compels humanity to reconsider what makes its mind unique.

Finally, all these developments are part of a broader framework of **ethical reflection** on AI. Never has the need for responsible governance of technologies been so apparent. Potential abuses—from intrusive mass surveillance to discriminatory algorithmic decisions—are prompting governments, international bodies, and academic communities to establish **ethical principles** and regulations. As early as 2021, UNESCO adopted

a recommendation on AI ethics, and the European Union is working on the AI Act, a set of regulations aimed at strictly regulating high-risk uses. At the conceptual level, contemporary thinkers have identified *three* major areas of ethical concern regarding AI: **privacy and surveillance, bias and discrimination**, and finally the issue of **the downgrading of human judgment** in crucial decisions [8]. The latter point, highlighted by philosopher Michael Sandel, directly concerns the impact on thought: *can intelligent machines "think well" on our behalf, or are there elements of human deliberation—doubt, empathy, practical wisdom—that no algorithm will ever be able to replace* [8]? In 2025, this question remains open. The ethical imperative is to find ways to **coexist with AI** so as to maximize the benefits for humanity (by enhancing our reasoning abilities, eliminating tedious work, improving access to knowledge) while minimizing the harms (standardization of thought, loss of cognitive autonomy, new digital inequalities, etc.). This requires education, system transparency, vigilance regarding biases, and keeping humans in the decision-making loop whenever necessary. More fundamentally, it is about preserving what makes us human in a world increasingly co-managed by artificial intelligences: **critical thinking, creativity, diversity of ideas, and moral responsibility**.

1.4 Organization of the Research Monograph

This general introduction having set out the context and issues, this monograph is structured around **six main axes** that will be developed in the following chapters. Each of these axes corresponds to a particular facet of AI's impact on human thought:

1. **Impact of AI on Human Cognition** – We will analyze how AI systems influence individual cognitive processes (memory, attention, reasoning), weighing the pros and cons of AI as an intelligence amplifier versus the risk of cognitive atrophy.
2. **Phenomena of Cognitive Standardization** – This chapter will explore the hypothesis of a homogenization of thinking patterns induced by the massive diffusion of the same tools and algorithms worldwide. We will examine signs of a reduction in cognitive and cultural diversity, as well as ways to preserve it.
3. **Algorithmic Biases and Feedback on Thought** – In this section, we will address the various biases present in artificial intelligences and their potential repercussions on human users (reinforcement of stereotypes, influence on decision-making, loss of trust or overconfidence in automated recommendations).
4. **Generative Artificial Intelligence and Creativity** – This chapter will focus on generative AIs (such as GPT language models, image generators, etc.) and their impact on human creativity and intellectual production. We will discuss the paradox of an AI capable of both stimulating imagination and standardizing certain creations, in light of recent studies on the subject.

5. **Toward Artificial Consciousness?** – Here, we will take a more prospective and philosophical approach by questioning the possibility of a conscious AI. We will review the criteria for consciousness, the advances and limitations of current AIs in this regard, and reflect on the implications that the emergence of artificial consciousness would have on our understanding of thought (both human and machine).
6. **Ethical Issues of AI and Cognition** – Finally, the last axis will address the cross-cutting ethical and societal dimensions: responsibility of designers and users, the need for regulation to prevent abuses (violations of privacy, unfair automated decisions), and the importance of rethinking education and training to prepare individuals to interact intelligently with AI systems without losing their intellectual autonomy.

Each chapter will draw on the most recent scientific literature and concrete examples to rigorously and nuancedly assess the **impact of artificial intelligence on human thought**. In the conclusion of the work, we will synthesize the lessons learned from these six axes and propose avenues for a future in which artificial and human intelligence co-evolve harmoniously, without one eclipsing or impoverishing the other.

The ultimate aim of this monograph is to enlighten both the scientific community and the general public on the cognitive and social challenges posed by the rise of AI, and to contribute to the reflection on the **conditions for a virtuous symbiosis** between humans and thinking machines at the dawn of this new era. [2] [5] [3] [4] [7] [8] [1]

Chapter 2

Artificial Intelligence and Human Cognition – Foundations and Interactions

Introduction

The spread of artificial intelligence (AI) is profoundly transforming our ways of life and raising new questions about its impact on human cognition. On the one hand, AI promises to **enhance our mental abilities** by automating certain intellectual tasks; on the other, some fear it may gradually **impoverish** our critical thinking skills [9]. Indeed, recent research suggests that while AI tools can facilitate the acquisition of basic skills, they may simultaneously **undermine users' deeper cognitive engagement** [9]. This chapter offers a rigorous analysis of the foundations of AI and human cognition, as well as the interactions between these two domains. We will successively address the key concepts of AI, the basics of human cognition, and then the main theoretical frameworks (cognitive sciences, cognitive load theory, cognitive offloading, etc.) that allow us to analyze the integration of AI into human life. Finally, we will discuss the evolution of AI and the tension between the promises of **cognitive augmentation** and the risk of a **decline** in mental skills. The objective is to adopt a neutral and analytical tone, relying on high-level scientific sources, in order to precisely identify the interactions between AI and human cognition and to assess how to leverage the benefits of AI without compromising our fundamental cognitive abilities.

2.1 Conceptual Foundations of Artificial Intelligence

The term **artificial intelligence** refers to machines or software capable of performing cognitive functions that are usually associated with the human mind—for example, perceiving, reasoning, learning, interacting with an environment, solving problems, or

even demonstrating creativity [10]. In other words, AI is the ability of a machine to accomplish tasks that normally require human intelligence [10]. These capabilities are achieved through computer algorithms leveraging modern computational power and, increasingly, through **machine learning** approaches (artificial neural networks, deep learning, etc.) trained on vast datasets.

We generally distinguish between **narrow AI** (or specialized AI), which excels at a specific task without claiming general understanding (for example, facial recognition or chess), and potential **general AI**, which would aim to replicate the flexibility and versatility of human intelligence. To date, deployed AI systems remain essentially narrow AIs, highly effective in circumscribed domains but lacking the cognitive versatility of a human being.

Despite their impressive achievements in certain fields, **current intelligent machines possess cognitive qualities fundamentally different from those of humans** [11]. For example, advanced computers have crossed the **exascale** threshold in computation, able to perform in one second as many operations as a human could in over 30 billion years [10]. However, this computational superiority is not accompanied by consciousness, intuition, or deep semantic understanding comparable to what a human brain can deploy. Thus, AI does not "think" in the human sense: it manipulates symbols or mathematical models without its own intention or cognitive experience.

Rather than considering AI as a duplication of human intelligence, it is more relevant to conceive of it as a **set of tools simulating certain intellectual functions** in a way that complements humans. This raises a central question: how can we use AI to leverage its specific strengths while leaving to humans the tasks where their **judgment** and **creativity** are irreplaceable [11]? This issue immediately highlights the importance of articulating AI and human cognition according to a **complementary** rather than opposing approach, capitalizing on the respective strengths of each.

2.2 Basics of Human Cognition

To understand the interaction between AI and human thought, it is necessary to recall the main features of *Homo sapiens*' cognitive functioning. The human brain processes information through an architecture that notably includes a **working memory** with limited capacity and a **long-term memory** for the durable storage of knowledge. Working memory (immediate cognitive awareness) can only handle a limited number of items at a time—typically about 5 to 9 simultaneous items according to Miller, or around 4 items according to more recent studies, due to interference and attention constraints. This reduced capacity explains why we are quickly overwhelmed if too much information must be processed at once. Consequently, **reducing mental load** by limiting irrelevant information is crucial for effective processing.

Humans have developed strategies to cope with these cognitive limits. For example,

we use **chunking** to integrate several items into a more easily memorable block, and we employ **attentional strategies** to filter important information from what is incidental. Furthermore, **long-term memory** stores acquired knowledge and skills: it is theoretically vast, but encoding new information into long-term memory requires time, repetition, and deep processing (elaboration, associations, etc.). The quality of this encoding strongly depends on the individual's **active engagement** during learning.

A key concept in cognitive science is the distinction between types of **cognitive load** (which will be discussed in detail in section 1.3.1). In summary, the more intrinsically complex a task is, the more mental resources it mobilizes (*high intrinsic load*). Added to this are loads induced by the way the task or information is presented (*extraneous load*), which can unnecessarily increase mental effort. Finally, the portion of mental effort actually invested in building new knowledge or skills is called *germane load*, and it is this germane load that directly contributes to deep learning. An **important implication** is that learning and skill development require a certain degree of germane cognitive effort: if all the work is pre-digested or automated, the brain no longer develops new schemas, and learning may suffer [12].

Moreover, humans have always sought to **extend their cognitive abilities** beyond biological limits by relying on external tools. History shows a constant inventiveness in delegating certain mental tasks: writing and note-taking to relieve memory, the abacus and then the calculator to facilitate calculations, or more recently the computer and the Internet to store and retrieve information. This externalization of cognitive functions to the environment is an integral part of human cognition. Thus, long before the era of AI, we already used artifacts to **amplify our thinking** or relieve it of overly burdensome constraints. Notably, the simple act of making a list or jotting down a reminder frees the mind from a memorization load—an example of "cognitive offloading" before the term existed [9]. These observations highlight that human cognition is **interactive and distributed**: it takes place in a context where tools and environment participate in information processing. The advent of AI only amplifies this phenomenon, raising with new urgency the question of the balance between what humans process themselves and what they delegate to artificial systems.

2.3 Theoretical Frameworks for Analyzing AI and Cognition

Several theoretical frameworks in cognitive science and psychology allow for an in-depth analysis of the interactions between AI and human cognition. Among these, we will focus on (1) **cognitive load theory**, which sheds light on how AI can either lighten or hinder learning processes, and (2) the concept of **cognitive offloading**, which describes the delegation of mental tasks to external tools and includes the notions of **extended**

cognition and **transactive memory**. These frameworks provide analytical grids for understanding the effects of AI on our ways of thinking, learning, and problem-solving.

2.3.1 Cognitive Load Theory

Cognitive load theory (John Sweller, 1980s) provides a useful framework for examining the influence of AI on learning and problem-solving activities [12]. As mentioned above, this theory distinguishes three types of load exerted on working memory during a cognitive task:

- **Intrinsic load**: the inherent difficulty of the task or content to be processed (for example, learning an abstract mathematical concept has a higher intrinsic load than memorizing a list of simple words).
- **Extraneous load**: the load added by the way information is presented or by irrelevant distracting elements. A confusing interface, unnecessary information, or incongruous multitasking increase extraneous load.
- **Germane load**: the cognitive effort directly invested in **deep processing** of information and the construction of new knowledge (schemas). This germane load corresponds to effective learning or deep reasoning.

The goal, according to the theory, is to **minimize unnecessary extraneous load** and **devote sufficient germane load** to useful processes, while taking into account the fixed intrinsic load of the task. In this perspective, AI can play an ambivalent role. On the one hand, AI systems can *reduce extraneous load* by eliminating secondary tasks or optimally presenting information. For example, an intelligent tutor can adapt the difficulty level of an exercise or filter displayed information so that the learner focuses on the essentials, avoiding overload from superfluous details. Similarly, an AI assistant can automate repetitive steps (data collection, intermediate calculations), thus lightening the user's mental burden in these aspects [12]. On the other hand, if AI is used excessively or inappropriately, it risks **reducing the germane load** engaged by the user. By delegating too much thinking or decision-making to the machine, the individual may adopt a passive role, no longer investing enough effort in understanding or actively solving problems. Yet, **reduced cognitive engagement** results in more superficial learning and weaker skill consolidation [12]. In short, cognitive load theory alerts us to the need to find a balance: AI can be beneficial for offloading working memory (reducing extraneous load), **provided** this does not come at the expense of the human's mental involvement in the fundamental aspects of the task (maintaining sufficient germane load). A judicious use of AI in education, for example, could consist of employing it to lighten administrative or repetitive tasks, while ensuring that the learner continues to actively engage their mind on the **core educational objectives** (analyzing, synthesizing, exercising creativity, etc.).

2.3.2 Cognitive Offloading and Transactive Memory

Cognitive offloading refers to the process by which an individual delegates part of a cognitive task to an external element, in order to reduce the mental load they bear [9]. Concretely, this means **externalizing** a cognitive operation—such as memorizing, calculating, or choosing—to a physical support or a third party. Classic cognitive offloading tools include objects as simple as a pencil and paper (to write down information instead of retaining it mentally), a calculator (to avoid mental calculation), or an electronic calendar (to avoid having to remember all appointments). With the advent of digital technology, these external supports now include **digital devices and AI**: note-taking apps, search engines, voice assistants, recommendation systems, etc., which handle an increasing share of our daily mental tasks [9].

From a cognitive perspective, offloading is a **well-understood adaptive strategy**: it allows us to **save mental resources** by freeing them from processing that the environment can perform for us [9]. Indeed, since our memory and processing capacities are limited, it is often rational to externalize a difficult or non-crucial task in order to focus our mind on what matters most at the moment. For example, jotting down an idea frees up working memory and allows us to move on to another task without fear of forgetting the first piece of information. Similarly, using a GPS for navigation offloads our mind from the need to calculate and follow a route, sparing us significant attentional and memory load (which in the past was managed via road maps and mental route planning). **Cognitive offloading** can thus improve efficiency and reduce **immediate cognitive strain**, while avoiding overloading working memory [9]. Experimental studies show that externalizing part of a task can increase performance on that task, especially when it is complex. For example, allowing participants to write down items to be memorized rather than retaining everything mentally increases their success when the amount of information exceeds what short-term memory could normally handle [13]. Offloading can therefore be a **tool for short-term cognitive optimization**, preserving our resources for the most demanding or creative aspects of the current work.

However, the **potential downside** of systematic externalization is a **weakening of internal cognitive abilities** in the long term. By becoming accustomed to always relying on an external support for a given task, we risk less frequently engaging the corresponding cognitive circuits, which can lead to a **decline in intrinsic performance** over time. As Risko and Gilbert note, offloading a task onto a tool certainly removes the immediate load, but "can also lead to a decrease in cognitive engagement and skill development" if this offloading becomes excessive [9]. In other words, "*use it or lose it*": what we no longer practice at all eventually atrophies. Researchers thus point out that cognitive offloading, while beneficial for immediate productivity, **can affect the development of critical thinking and memory** when reliance on external tools becomes too systematic [9]. For example, if it becomes reflexive to look everything up online, we exercise our personal memory less on everyday topics. Indeed, the instant availability of

information via the Internet has given rise to what Sparrow and colleagues called the "**Google effect**": individuals, knowing they can retrieve information at any time online, tend to remember the location of the content rather than the content itself [9]. The Internet thus serves as an **external memory** (or collective *transactive memory*), which is convenient but may raise **concerns about the decline of individual memory** and our ability to remember without external help [9]. This phenomenon of offloading onto the web is increasingly documented by psychologists: the term *digital amnesia* is also used to describe the tendency to forget information easily accessible online, where an effort to memorize would have been made in the absence of this technological recourse.

Another important aspect of cognitive offloading via AI concerns the **degree of trust** we place in intelligent systems. The more a user trusts an AI tool, the more likely they are to delegate tasks to it without double-checking, which intensifies the offloading phenomenon. This trust may stem from the perceived reliability of the tool or simply from habit. Yet, blind trust can lead to a **vicious circle**: fully trusting AI, the user checks less for themselves and exercises less critical judgment, which over time can make them **dependent** on the tool even for tasks they could accomplish (or errors they could detect) if they were more vigilant [9]. Recent studies show, for example, that in education, the more students trust answers provided by an AI agent, the less they invest in source verification or personal reflection, which can diminish their skills in critically evaluating information [9]. The **risk** is that **cognitive dependence** sets in: AI becomes an *autopilot* for thought, and the user, confident in this artificial copilot, relaxes their mental effort. In the long run, this could reduce their ability to perform the task unaided or to detect AI errors. It is therefore important to carefully study how to maintain a sufficient level of **cognitive control** and **critical thinking** when relying on AI systems, so as not to entirely relinquish the steering wheel of our mental processes.

To concretely visualize the concept of cognitive offloading, **Figure 2.1** below provides a schematic illustration. It shows how an individual can transfer part of their mental load to an external AI device to relieve their brain.

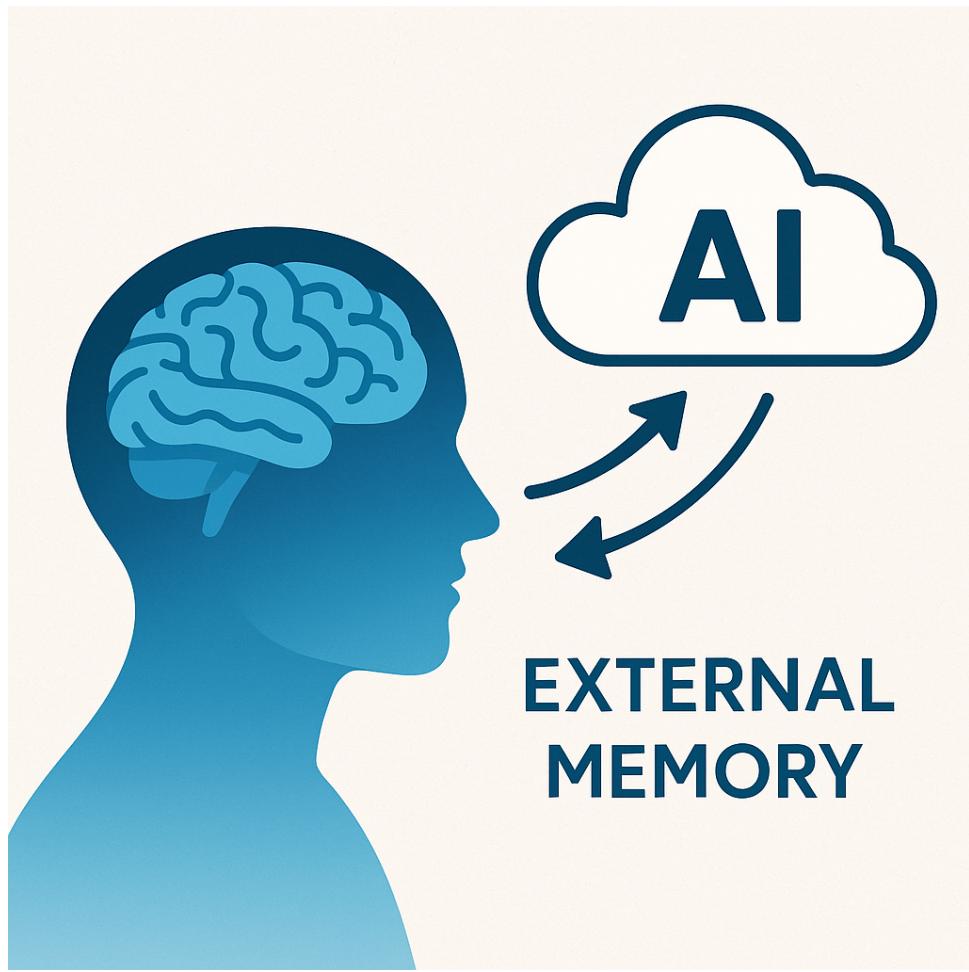


Figure 2.1 – AI as an external memory, illustrating cognitive offloading.

2.3.3 Other Theoretical Perspectives

In addition to cognitive load theory and the concept of offloading, other conceptual frameworks enrich the analysis of the relationship between AI and cognition. Two notable perspectives are briefly mentioned here:

- **Extended cognition (extended mind):** In philosophy of mind, Clark and Chalmers (1998) proposed the idea that external tools can be an integral part of the cognitive process—in other words, the mind is not limited to the brain; it extends to the objects with which we interact. This perspective, confirmed by many observations (for example, the use of paper as *external memory* or a pencil for thinking through sketching), is more relevant than ever in the age of AI. If a recommendation algorithm guides our choices or an intelligent assistant completes our sentences, these can be seen as an **extension of our mental processes**. The theory of extended cognition thus invites us to rethink the boundary between human and machine: rather than considering AI as entirely separate from us, we can view it as part of a **hybrid human-AI cognitive system**. This also implies a responsibility:

any flaw or bias in the external tool can directly influence our cognition, since we integrate it into our reasoning.

- **Transactive memory:** Initially studied in the context of human groups (Wegner, 1987), transactive memory describes a system in which several individuals share the task of remembering—each memorizes certain information and knows they can rely on other group members for information they have not retained. By analogy, the relationship between an individual and online knowledge bases or AI agents can be seen as a human-machine transactive memory system [9]. The individual remembers *how* or *where* to find information (on Wikipedia, on Google, via a particular app) rather than the information itself, thus delegating retention to the external source. The quality of this transactive system depends on the reliability of the external agent: a capable and trustworthy AI partner can greatly increase the **reservoir of accessible knowledge** for an individual. Conversely, relying on an unreliable or biased source can lead to integrating erroneous information or neglecting to form one's own understanding. The notion of transactive memory applied to AI thus reinforces the importance of developing **transparent and reliable AIs**, as well as **digital literacy** among users so they understand when and how to trust information provided by the machine.

In summary, theoretical frameworks from cognitive science help us finely analyze the impact of AI. They highlight that AI can be both an **amplifier** of our mental abilities (by reducing extraneous load, serving as extended memory, etc.) and a **factor of cognitive disempowerment** (if we offload excessively, at the risk of no longer exercising certain faculties). As AI increasingly integrates into our thought processes, it becomes crucial to understand these dynamics to guide the development and use of these technologies in an informed manner.

2.4 Evolution of AI and Integration into Human Life

Since its birth in computer science laboratories in the 1950s, AI has undergone spectacular evolution and has gradually integrated into almost every sphere of human life. Initially confined to chess-playing programs or expert systems used by engineers, it is now **all around us**, often invisibly. The current ubiquity of AI is such that **most people use it daily without always realizing it** [10]. For example, every time you search the Internet, ask **Siri** on your phone or **Alexa** on a smart speaker to set a reminder or give you the weather, you are interacting with AI [10]. Likewise, your email's spam filters, movie or music recommendations on streaming platforms, or your car's GPS calculating the optimal route, are all now commonplace services that rely on AI algorithms.

The **integration of AI into daily life** really accelerated from the 2000s with the Internet revolution, then in the 2010s–2020s with the rise of **smartphones** and con-

nected objects. In 2016, already 89% of American households owned a computer at home [10], and this figure only increases if we include smartphones, tablets, and other smart devices that accompany us constantly. At the same time, computational power and algorithmic efficiency have followed an exponential law: today's supercomputers can perform hundreds of **quadrillions of operations per second**, and modern machine learning techniques leverage massive data (*big data*) to achieve performance levels once unimaginable. For example, in 2023–2025, the emergence of large language models (such as *ChatGPT*) illustrated the leap forward in AI's ability to understand and generate natural language, a cognitive function long considered unique to human intelligence.

Today, AI is thus **everywhere**—from healthcare (automated medical imaging, diagnostic support systems) to transportation (autonomous vehicles in development), education (intelligent tutors), commerce (virtual assistants, personalized offers), or domestic tasks (robot vacuums, smart thermostats). This ubiquity of AI means that our **cognitive processes are increasingly linked** to these artificial systems in our daily activities. The ways we inform ourselves, make decisions, remember, or learn are *mediated* by intelligent technological tools. For example, instead of memorizing multiple facts, we know we can retrieve them online in seconds; instead of remembering all our tasks, we delegate part of this function to organization and reminder apps. AI tools thus directly influence key cognitive functions such as **memory** (by facilitating the acquisition and retrieval of information), **attention** (by filtering or prioritizing the information flow for us), and **problem-solving** (by providing ready-made analyses or solutions to complex problems) [9]. In other words, the integration of AI into human life presents a **double-edged sword**: it offers unprecedented potential for **cognitive assistance**, while posing the challenge of preserving the user's autonomy and mental skills.

From a historical perspective, it is worth noting that every major technological advance has raised similar concerns. Socrates, it is said, was already wary of the invention of writing, which he believed might weaken memory by allowing people to "no longer learn by heart." Similarly, the arrival of the pocket calculator made some educators fear the disappearance of mental arithmetic, and the rise of GPS a loss of sense of direction among younger generations. With modern AI, these questions return amplified: if a machine can think, decide, or create for us, **what will remain of our own faculties?** Will we see an augmented human, freed from menial tasks to devote themselves to higher activities, or an assisted human, intellectually lazy and dependent on the machine? It is likely that reality is not black and white: AI can both **liberate us cognitively** and **make us less vigilant**. In the next section, we will explore precisely this tension between cognitive augmentation and decline, in order to identify the conditions for a virtuous balance.

2.5 Cognitive Augmentation vs. Cognitive Decline: A Tension to Manage

In light of the preceding points, it is clear that AI has an ambivalent impact on human cognition. It can act as a **formidable amplifier** of our mental abilities, but also as a **factor of attrition** of certain skills if its use is not controlled. This central paradox is summed up in the alternative *cognitive augmentation* versus *cognitive decline*. On the one hand, AI offers **opportunities for augmentation**: it assists us, enables us to do more and better, and extends the scope of what we can accomplish intellectually. On the other, it carries the risk of **excessive dependence** leading to a gradual erosion of our know-how and autonomy of thought. It is crucial to analyze these two facets to develop strategies that maximize benefits while minimizing risks.

On the side of cognitive augmentation, the potential benefits of AI are numerous and well documented. Here are a few examples:

- **Reduction of mental load and increased efficiency:** By automating repetitive or calculation-intensive tasks, AI allows these tasks to be accomplished more quickly and with fewer errors than a human. This **frees up time and cognitive resources** to focus on more complex or creative aspects [9]. For example, AI software can instantly sort thousands of data points where a human would take hours, allowing the latter to focus on interpreting results rather than raw processing. Similarly, in intellectual work, AI can provide tools (spell check, code completion, rapid information retrieval) that relieve the user of part of the extraneous load and enable greater productivity.
- **Improved performance on complex tasks:** AI can help humans solve problems they could not tackle alone, providing **superhuman capabilities** in certain domains. For example, machine learning algorithms detect subtle patterns in big data that the human mind could not spot, thus improving decision quality in fields such as medical diagnosis or financial analysis. The term **augmented intelligence** is used when AI collaborates with the human expert to achieve a result superior to what either could obtain alone. A concrete case is that of assisted creation: in design or programming, AI tools can suggest ideas or alternative solutions that the human designer would not have considered, thus broadening the range of possibilities.
- **Accelerated and personalized learning:** In education, intelligent tutors and other adaptive learning systems offer **personalized instruction** to the student, finely adjusting to their level and progress (e.g., by proposing exercises that are neither too easy nor too difficult, providing targeted explanations for errors) [12]. Such approaches can improve the acquisition of basic knowledge and more effectively address gaps. Moreover, immediate access to a vast knowledge base (via the

Internet and search engines) allows anyone to learn new information or skills autonomously whenever a need or curiosity arises. AI acts as an **always-available educational assistant**. Studies have shown that using tools such as adaptive quizzes or conversational agents for practice can increase **information retention** in learners, especially when these tools are used in addition to active study [12]. In short, when used well, AI can be a **cognitive catalyst** that boosts our intellectual performance and compensates for some of our individual weaknesses (memory gaps, lack of expertise in a field, etc.).

- **Stimulation of creativity:** Another area where AI can act augmentatively is **creativity**. Recent experiments indicate that human-AI collaboration can lead to increased creativity. For example, in an experiment with students on creative problem-solving, those who used an AI idea generator (such as GPT) produced **more varied and detailed ideas** than those who worked alone [12]. AI can provide suggestions, randomly explore new avenues, or combine elements in ways a human would not have thought of, serving as a springboard for creative thinking. Improvements in **fluency** (number of ideas), **flexibility** (diversity of ideas), and **elaboration** (richness of detail) are observed when AI is used as a brainstorming tool [12]. In the arts, AI tools offer new palettes (e.g., image generation, musical composition assistance) that expand the means of expression for human creators. AI thus becomes a muse or collaborator, rather than a mere executor.

All these advances paint a positive picture where AI acts as an **intelligence amplifier**. Nevertheless, it is necessary to examine the flip side: what are the **trade-offs** or **limitations** of these cognitive augmentations? This is where the issue of **potential cognitive decline** induced by AI comes in.

On the side of **cognitive decline**, several risks and possible negative effects have been identified by researchers:

- **Atrophy of certain skills:** By automating a task previously performed manually or mentally, we risk gradually losing mastery of that task. This is the idea of *disuse* in psychology: no longer practicing a skill leads to its weakening over time. For example, if we always rely on a *GPS* for navigation, we will exercise our spatial representation and orientation skills less; if we systematically use a calculator for every calculation, we will use our mental arithmetic skills less. At the societal level, some educators already observe a decline in performance in unaided calculation or memorization of simple facts among younger people, correlated with the permanent availability of AI or the Internet to perform these operations for them. AI provides **intellectual comfort** that can lead to **cognitive laziness** regarding basic skills. Nicholas Carr, in his essay *The Shallows*, argued that the abundance of easily accessible information online makes us less inclined to retain information in detail, which aligns with these findings [9].

- **Reduction in engagement and critical thinking:** Several studies highlight a negative link between intensive use of AI tools and critical thinking or complex problem-solving skills [9]. The proposed explanation is that becoming accustomed to finding ready-made answers or receiving solutions from an intelligent system may encourage users to accept these answers without **scrutinizing them critically** or exercising their own reasoning. For example, a study on students showed that when an AI system directly provided explanations or text summaries, they tended to take them at face value and failed to detect certain inconsistencies or biases, whereas students required to read and analyze the texts themselves demonstrated more critical thinking [14]. The danger is thus a *passive* attitude toward information: the user becomes a consumer of AI conclusions rather than an active producer of understanding. This results in a **weakening of critical thinking** and the ability to solve new problems, especially if the habit is formed early in learning (the so-called "tutor effect": the student always relies on hints from the intelligent tutor instead of searching independently). Empirically, in the Gerlich (2025) study mentioned above, a high level of AI tool use correlated with lower scores on standardized critical thinking tests, and this link was **mediated by cognitive offloading**—that is, it was because users offloaded many tasks to AI that they practiced their critical thinking less [9].
- **Dependence and loss of autonomy:** A tangible risk is becoming **dependent** on AI to the point of being unable to do without it, even in situations where it is unavailable or inappropriate. If, for example, we become accustomed to an intelligent agent making all routing decisions for us (GPS), what happens when this aid fails? Similarly, a writer who has always used an automatic suggestion tool to continue their sentences may struggle to regain their autonomous writing flow. This dependence can also be mental: we may lose **confidence in our own abilities** after prolonged AI use, underestimating ourselves compared to the machine's performance. At the extreme, some mention a risk of **unlearning**: the user no longer sees the point of learning something since "the machine knows or will do it better than me." Yet, giving up on learning or thinking for oneself is obviously a major cognitive impoverishment. Psychologically, this relates to the concept of *complacency* (overconfidence in automation) studied in safety: for example, pilots with autopilot may become less attentive and less able to react to the unexpected. Transposed to everyday cognition, **too much trust in AI can lead individuals to lower their mental guard**, making them vulnerable to AI errors or a general loss of competence [9].
- **Bias and undetected errors:** Finally, an insidious effect of over-reliance on AI is the possible incorporation of **biases** or **errors** into our own thinking if we are not vigilant. AI systems, especially those based on massive data, can reflect biases (e.g., cultural biases, stereotypes, or simply sampling biases). If we accept

their results uncritically, we risk **perpetuating these biases** in our decisions. Moreover, AI is not infallible: it can make mistakes. A calculator will probably never make a calculation error, but a recommendation system can miss a relevant option, a conversational agent can confidently state something factually wrong (*hallucination*), etc. Maintaining our evaluation and verification skills is therefore crucial to avoid a **decline in the overall reliability** of our AI-integrated cognitive processes.

This overview shows that AI generates a **complex dialectic** between cognitive gains and losses. To better synthesize these elements, **Table 2.1** below summarizes some key points of the benefits of augmentation and the risks of decline associated with AI use.

Table 2.1: Examples of AI effects on cognition—between augmentation and decline.

Potential Augmentations from AI (Benefits)	Potential Risks of Cognitive Decline Linked to AI
Facilitated access to information and extended external memory: AI provides instant access to vast knowledge, serving as an <i>auxiliary memory</i> for the user. This reduces the need to memorize trivial facts and frees the mind for more elaborate tasks. [9]	Digital amnesia and dependence on external memory: by constantly searching online or recording everything in devices, we may train our own memory less. We remember where to find information rather than the information itself, possibly weakening personal long-term memory. [9]
Automation of routine tasks: AI handles repetitive or technical operations (calculations, data sorting, monitoring), increasing productivity and allowing focus on analysis or creativity. For example, algorithms scan millions of documents much faster than a human, summarizing the essentials. [9]	Loss of skill in delegated tasks: by no longer practicing certain basic tasks, users may lose proficiency. For example, systematic use of GPS can harm sense of direction, and reliance on autocorrect can weaken spelling. Over time, users become unable to perform these tasks without AI, reducing autonomy.

Potential Augmentations from AI (Benefits)	Potential Risks of Cognitive Decline Linked to AI
<p>Decision support and augmented analysis: AI can process massive volumes of data and detect complex patterns (correlations, trends) beyond human capabilities [15]. Integrated into decision-making, it enables better-informed choices (e.g., assisted medical diagnosis, driving aids). Humans thus benefit from a second informed opinion or an exploration tool to support their thinking.</p>	<p>Overconfidence and superficial thinking: when faced with AI-proposed suggestions or solutions, humans may develop an excessive trust bias and no longer exercise sufficient critical thinking [9]. They risk validating answers automatically (<i>cognitive complacency</i>) without deeper analysis. This passive acceptance can lead to undetected errors and a decline in independent analytical ability. [12]</p>
<p>Stimulation of creativity and learning: AI can act as an interactive tutor or brainstorming partner. In learning, it personalizes exercises and provides immediate feedback, strengthening student engagement and helping them progress at their own pace [12]. In creation, it generates new ideas (images, phrases, melodies) that inspire the human creator, enabling augmented creativity by combining the best of both agents [12].</p>	<p>Decrease in learning effort and creative fixation: the ease provided by AI may encourage a form of intellectual passivity: the learner, too guided, may practice less independent problem-solving or long-term memorization (superficial learning) [9]. In creation, too much AI assistance can lead to standardization or fixation on its suggestions, stifling human originality. Studies note, for example, a decrease in confidence in one's own creativity and a tendency to reuse AI suggestions at the expense of more personal explorations [12].</p>

As this table shows, AI can be both a valuable ally for our cognition and a subtle trap for it. It is not a matter of claiming that AI inevitably makes us "dumber" or lazier—many studies prove, on the contrary, that it can help us be smarter, more efficient, and more creative [12]. However, the risks of cognitive decline exist and deserve attention. They appear especially when AI use becomes excessive or indiscriminate, without educational safeguards or awareness of the machine's limits. For example, AI can weaken critical thinking if users take its answers as absolute truth without examination [9], but this danger can be countered by training users to *verify* and *complement* information obtained via AI.

In practice, the challenge is to find a balance in AI use—taking advantage of its undeniable benefits for cognitive augmentation, while avoiding falling into harmful

dependence. Several avenues can be mentioned to manage this tension:

- **"Augmented intelligence" approach:** Rather than aiming for total automation, it is desirable to design AI as a tool to *augment* human intelligence, not replace it. This means always keeping the human "*in the loop*" for complex tasks, and ensuring that AI serves as a copilot, not the sole pilot. Users should be encouraged to **interact** with AI, ask it questions, and understand its answers, rather than passively consuming its outputs.
- **Training and digital literacy:** To prevent AI from eroding our faculties, it is crucial to train users (from school and throughout life) in thoughtful use of these tools. This includes **critical thinking** about digital sources, understanding possible AI biases, the ability to perform a task manually or mentally if needed, etc. For example, in education, AI can be integrated as a writing aid while requiring students to analyze, correct, and justify its suggestions. Researchers thus recommend **combining AI with active pedagogical activities** that force the learner to remain cognitively engaged, to avoid excessive passivity [9].
- **Ergonomic design of AI:** On the designers' side, it is possible to mitigate the risk of decline by building AI systems that *encourage* user cognitive participation. For example, an assistant could explain its reasoning (prompting the human to follow the logic), or not provide everything at once to leave some work to the user (e.g., a pedagogical GPS that sometimes asks the user to validate the best route among two choices, thus training their map-reading skills). Similarly, **cognitive reminders** can be integrated—for example, the system could suggest "And what do you think?" after giving a recommendation, to stimulate critical evaluation rather than blind acceptance.

Ultimately, the main idea is that AI should be considered a **tool** serving human intelligence, not a complete substitute. Used synergistically, it can free us from certain constraints and extend our abilities without causing their decline. Conversely, uncontrolled use could lead to **intellectual disempowerment** with harmful effects.

To conclude, this image offers a conceptual representation of the balance between cognitive augmentation and decline due to AI.

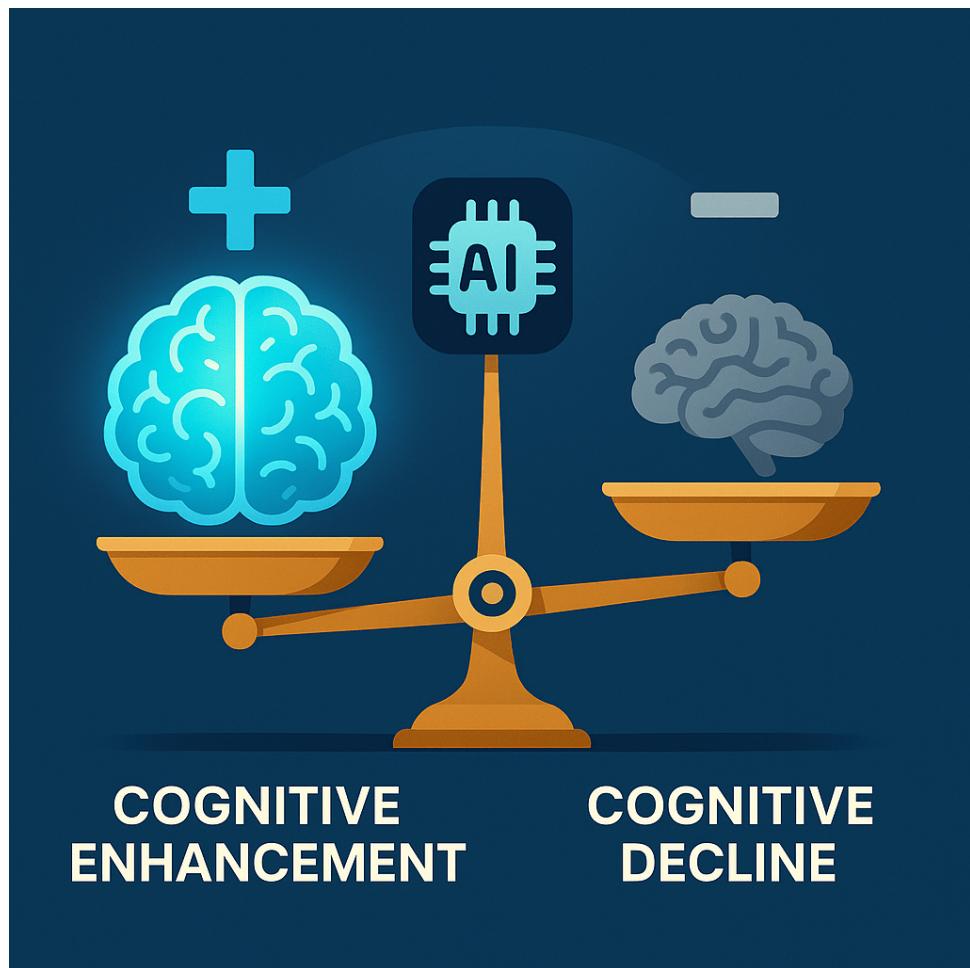


Figure 2.2 – The dialectic of cognitive enhancement and decline with AI.

2.6 Conclusion

In conclusion, the impact of AI on human thought cannot be analyzed in unequivocal terms of "benefits" or "harms." It is a **nuanced continuum**, where AI acts as an **amplifier** of our abilities while posing the challenge of **permanent cognitive vigilance** on our part to avoid becoming intellectually dependent. This first chapter has explored the conceptual foundations of AI and human cognition and highlighted the main mechanisms at play—offloading of cognitive load, transactive memory, changes in attentional engagement, etc. The analysis reveals a **dialectical tension** between augmentation and decline, which must be managed through thoughtful and measured use of AI technologies.

Given the inevitable penetration of AI into all areas of society, the question is not whether to accept or reject these tools, but rather **how to incorporate them wisely** into our cognitive activities. This requires advances both from developers (to create *human-compatible* AIs that support cognitive effort rather than replace it) and from users and educational systems (to learn to augment ourselves with AI without ceasing to **think for ourselves**). In short, it is a new **Human-Machine pact** that must be

constantly negotiated: a partnership in which AI is an **ally** stimulating our intellect, not an instrument of dulling.

The following chapters of this monograph can build on these foundations to examine in more detail related questions, such as the effects of AI on **ethical decision-making**, on the **social and cultural dynamics** of cognition (collective intelligence, distribution of knowledge), or the concrete means to **regulate** and **guide** the development of AI in order to maximize its positive outcomes for the human mind. The challenge is considerable, but by combining insights from AI research, cognitive science, and social sciences, it is possible to meet this challenge and make AI not the gravedigger, but indeed the **catalyst** of an enriched and emancipated human thought. [9] [9]

Chapter 3

Cognitive Standardization in the Age of Artificial Intelligence

3.1 Homogenization of Content, Language, and Cultural References

The rise of ubiquitous artificial intelligence (AI) systems raises concerns about **cognitive standardization**, that is, the progressive homogenization of ways of thinking on a global scale. Conversational and generative AIs, in particular, produce content formatted according to dominant—often Anglo-Saxon—standards, which tends to **standardize language and cultural references**. A report by the French Senate highlights that the dominance of AI by Anglo-Saxon actors "*risks strongly accentuating the cultural hegemony of the United States*", impoverishing linguistic and cultural diversity, while creating "*cognitive standardization*" [16]. In other words, the more users worldwide rely on tools powered by similar data and cultural models, the more their ideas, expressions, and frames of reference risk converging.

Several analyses point to the danger of such cultural convergence. Large generative language models (LLMs) often favor standard English and reflect dominant Western norms, even when used by speakers of other languages or cultures. For example, a 2024 study showed that a *Western-centric* text autocompletion system could insidiously influence the writing of non-Western users: when faced with suggestions from an English-trained GPT-4 model, Indian participants produced texts **adopting a more Western style**, losing in the process certain nuances of their own cultural expression [17]. In other words, AI "*homogenized writing towards Western styles by silently erasing non-Western modes of expression*" [17]. This concrete result illustrates how the widespread use of AI tools **can smooth out cultural differences** in intellectual productions.

A similar phenomenon is observed in the linguistic domain. In educational contexts, it has been noted that tools like ChatGPT tend to favor the dominant standard language (e.g., formal academic English) at the expense of dialectal or stylistic diversity. Educa-

tional researchers have warned that "*by privileging standard English, AI programs such as ChatGPT may encourage linguistic homogeneity*" and lead to the erasure of certain varieties of language and thought [18]. By depriving learners of the richness of their own idioms and expressive processes, AI could impoverish the range of ways of thinking and expressing oneself. Indeed, "*reducing writing to a conformist final product, in favor of the dominant norm, risks destroying the richness and complexity of the languages students bring with them*", limiting their ability to "*understand the world in new ways*" [18]. Language, as a vehicle of thought, is thus standardized under the insidious influence of AI, which dictates "*how we experience the world*" [18].

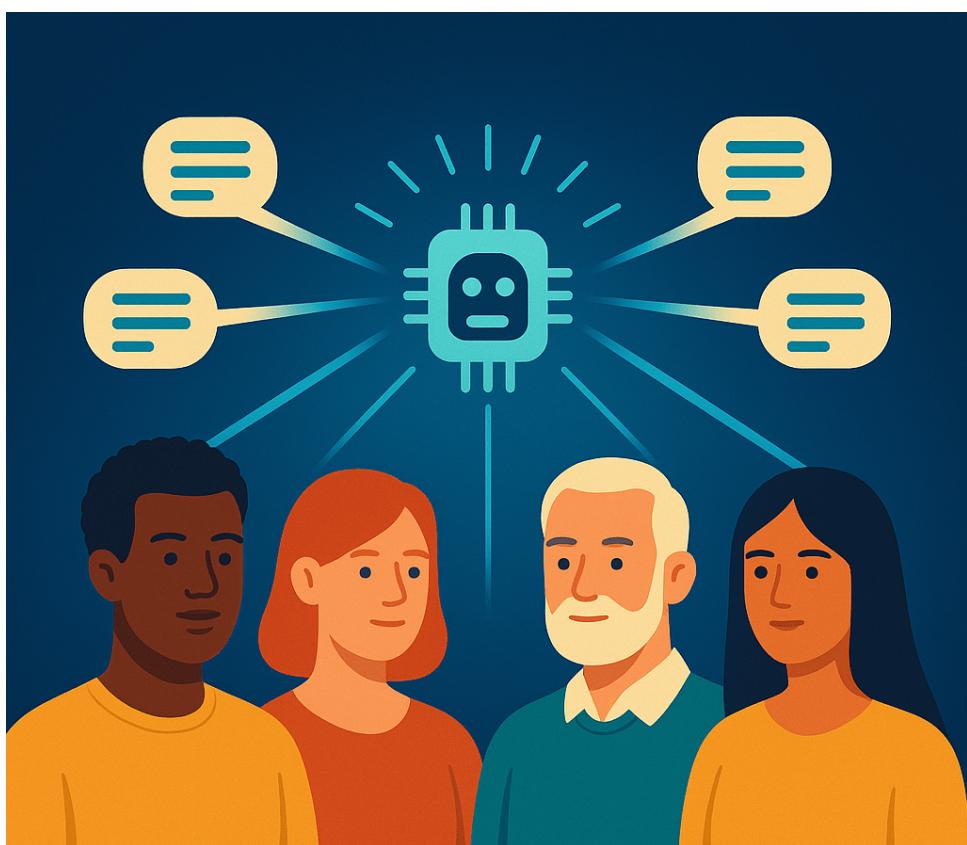


Figure 3.1 – Conceptual representation of cognitive standardization by AI.

In sum, as generative AIs become the preferred intermediaries for information and creation, **the risk is that everyone expresses and thinks with the same words, the same references, and the same patterns**. Cognitive standardization through the homogenization of content and language is no longer a mere abstraction: it is already evident in the trends toward linguistic and cultural standardization induced by AI. This evolution poses a major challenge: how can we preserve the diversity of ways of thinking and expressing ourselves in the face of technologies that, by design, tend to generalize dominant *patterns* in our minds?

3.2 Filter Bubbles, Algorithmic Bias, and the Erosion of Opinion Diversity

Another key mechanism of cognitive standardization in the AI era lies in **personalization algorithms** that govern our information streams (social networks, search engines, recommendations). In theory, these systems adapt content to each individual's preferences. In practice, they often confine the user within what Eli Pariser called the *filter bubble*: a closed informational ecosystem where the information presented reinforces the individual's existing beliefs and tastes, to the detriment of exposure to novelty or contradiction. AIs contribute to **polarizing viewpoints and standardizing opinions** within each bubble, by showing each person only a partial vision of the world. The aforementioned Senate report notes that "*cognitive capitalism*" combining screens and AI has led to a true "*attention economy*", in which the user is **trapped in filter bubbles** that "*polarize each person's views into subjective beliefs*" [16]. Thus, "*as many mental prisons*" are formed on an individual scale [16]. By an apparent paradox, we witness both hyperpolarization of opinions between groups and standardization of thought *within* each cultural or ideological silo. Each person, isolated in their algorithmic sphere, sees their convictions reinforced to the point of believing they constitute the norm, while mutual understanding between divergent groups withers. **Algorithmic biases** present in AI systems exacerbate this erosion of opinion diversity. By training models on historically biased data, or by optimizing engagement through sensationalist content, AI designers may inadvertently standardize the perspectives presented to users. For example, studies have shown that language models tend to reflect and amplify majority cultural stereotypes (gender, race, etc.) [17] [17]. This means that **the responses produced risk conveying a unilateral view** of the world, aligned with dominant prejudices, and neglecting perspectives from minority or marginal groups. If the user does not exercise active critical thinking, they will absorb these biases as self-evident, thus internalizing a way of thinking standardized by AI's blind spots.

The long-term risk is a form of **closed-circuit thinking**, where each person sees their preconceptions constantly confirmed by systems that know them too well. Deprived of fruitful intellectual confrontations and exposure to otherness, **critical thinking dulls** and thought becomes normalized. This phenomenon is already observed on social networks, where algorithmic personalization has led to ideological echo chambers. AI, by filtering and ranking information to maximize our screen time, can inadvertently *shrink* our cognitive horizon. Without safeguards, the **diversity of opinions** necessary for sound judgment risks collapsing, with each individual remaining confined to a narrow corridor of conventional thinking.

It should be noted that this intra-bubble standardization does not mean the absence of conflict in society—on the contrary, opposing bubbles may ignore or violently confront each other—but it does mean the disappearance, within each group, of plurality of voices

and questioning. **Collective critical thinking weakens** when AIs continually reinforce our biases instead of challenging them. Vigilance is therefore required regarding the use of these algorithms: without conscious intervention to diversify information sources, AI can become a powerful engine of thought standardization on both individual and collective scales.

3.3 Impacts of AI on Critical Thinking and Cognitive Skills

Cognitive standardization manifests not only in the content of the information we consume, but also in **an erosion of certain cognitive abilities and critical thinking** among intensive AI users. By delegating more and more intellectual tasks to machines—analyzing data, summarizing texts, proposing solutions—humans risk exercising these faculties less themselves, a phenomenon known as *cognitive offloading*. This transfer of mental load to AI can bring increased comfort and efficiency, but recent research suggests it is also accompanied by a **measurable decline in certain thinking skills**. A 2025 study of 666 participants highlighted a **strong statistical link between frequent use of AI tools and the decline in critical thinking scores** measured by standardized tests [19]. More specifically, a significant negative correlation ($r = -0.68$, $p < 0.001$) was observed between AI use and the ability to critically evaluate information and solve problems thoughtfully [19]. This result suggests that intensive AI users "exhibit a decrease in their ability to critically evaluate information and engage in thoughtful problem-solving" [19]. **Cognitive offloading** was identified as a key mediating factor: in the same study, the tendency to rely on digital tools for cognitive tasks was strongly correlated both with AI use ($r = +0.72$) and with the decline in critical thinking scores ($r = -0.75$) [19]. In other words, it is because we entrust AI with memorizing, calculating, and deciding for us that our own cognitive *muscles* partially atrophy from lack of training.

These quantitative results confirm the findings of several recent studies in the educational field. A 2024 systematic literature review highlights that **overuse of dialog-based AI systems** can "negatively impact critical thinking, analytical reasoning, and decision-making skills" among students [20]. Learners **become less able to analyze information themselves, to formulate logical arguments, and to make reasoned decisions independently** [20]. Moreover, "overreliance on AI for acquiring information can negatively impact critical thinking dispositions"—that is, attitudes conducive to critical thinking, such as doubt, intellectual curiosity, the search for evidence, etc. [20]. By becoming accustomed to immediately finding an answer via AI, users develop these habits of verification and questioning less. The long-term effect is a **decrease in skepticism and critical examination**, both essential components of enlightened thinking.

It is useful to detail the different aspects of critical thinking and see how AI can

influence each of them. Table 3.1 below summarizes the main **components of critical thinking** and the **potential effects of AI** on them, according to available studies:

Table 3.1: Potential impacts of AI on different components of human critical thinking.

Component of Critical Thinking	Potential Effects of AI on This Component
Analysis and Interpretation of Information	AIs provide ready-made analyses (summaries, explanations), which can reduce users' practice of autonomous analysis. Less solicited, they risk losing analytical sharpness [20].
Critical Evaluation and Verification (Skepticism)	Faced with a fluent AI response, users may neglect to verify it. Excessive trust in AI outputs weakens methodological doubt and source verification [20].
Logical Reasoning (Deductive/Inductive Inference)	AI models based on statistical induction obscure the deductive approach. There is a bias in favor of induction , which may marginalize training in formal logical reasoning [16].
Problem Solving and Autonomous Decision-Making	By accustoming users to ready-made solutions, AI can lead to a decline in the ability to solve novel problems . Intensive users show less initiative and autonomous judgment in decision-making [19] [20].
Creativity and Divergent Thinking	AI-generated suggestions tend to limit the exploration of original ideas by offering prepackaged solutions. Users are exposed to a narrower range of options, which can stifle their imagination [21].
Curiosity and Autonomous Learning	The ease of an immediate AI response can undermine intellectual curiosity . The effort of personal investigation and learning by exploration is discouraged, even though it is at the heart of critical thinking [20].

Table 3.1: Potential impacts of AI on different components of human critical

thinking. This table highlights the risks identified in the literature regarding the effects of intensive AI use on cognitive skills. It should be noted that these impacts may vary depending on individuals and usage contexts, but they underscore the need for vigilance regarding the role assigned to AI in our intellectual processes.

In summary, **over-delegation of cognition to AI risks resulting in atrophied and standardized critical thinking.** If everyone relies on the same tools to analyze, verify, or solve problems, they may adopt increasingly similar thinking patterns, dictated by the internal workings of these tools. The diversity of cognitive approaches—some more analytical, others more intuitive, some focused on contradiction and doubt, others on boundless creativity—constitutes the richness of collective intelligence. Yet it is precisely this diversity that is threatened when standard AI solutions predominate. The next chapter will examine in more detail the consequences of this possible standardization of thought, particularly in terms of creativity and preferred modes of reasoning, before considering ways to address it.

3.4 Toward Uniform Thought? Consequences for Creativity and Reasoning

The trends described above lead to a troubling question: are we heading toward "uniform thought," shaped by AI? Two domains particularly illustrate this concern: human creativity and modes of reasoning (induction versus deduction).

Regarding creativity, AI acts as a double-edged sword. Certainly, it can assist humans by freeing up time (for example, quickly generating a draft article or design), but this very assistance risks standardizing creative output. By relying on patterns derived from past data, AIs generate *average* works in the statistical sense, often conventional, which may lead creators to unconsciously align with these dominant models. A parallel can be drawn with the industrial era: just as mass production standardized objects, AI-generated content tends to standardize ideas. Thus, in fields such as writing, music, or graphic design, the convenience provided by AI could come at the cost of diminished originality. As Barnes (2024) observes, "*there is a risk that [AIs] limit human creativity by restricting the range of ideas and expressions to which individuals are exposed*" [21]. Instead of groping, experimenting, and venturing off the beaten path—a process often essential to innovative discoveries—the AI-assisted creator may be tempted to choose the quickest solution suggested by the machine. This primacy of convenience over exploration can stifle the creative spark. It has been reported that the "*temptation to rely on AI for a quick answer diminishes the opportunity to engage in deep and iterative reflection, often leading to innovative solutions*" [21]. Ultimately, if everyone uses the same algorithms to innovate, might we not see the emergence of increasingly similar works and ideas, calibrated to the AI's cognitive mold?

As for **modes of reasoning**, current AI overwhelmingly favors **inductive inference** (learning from millions of examples) over **deductive inference** (applying general principles to particular cases). This predominance of induction is not without consequences for how humans approach problem-solving. The aforementioned parliamentary report warns: "*the generalization of particular cases under the influence of massive data processed by connectionist AI has become the rule*", so much so that in the long run "*this era of AI and Big Data will lead all inhabitants of the planet to think in the same way [...]*
] oriented toward induction" [16]. In other words, there is a **risk of cognitive monoculture** where, by using inductive tools, everyone unwittingly comes to favor probabilistic, correlative reasoning at the expense of more structured logical-deductive thinking. Yet deductive thinking—which proceeds by formal logic, step-by-step demonstrations—has historically underpinned many scientific and philosophical advances. If it were neglected, our collective ability to **trace back to first causes, rigorously test a hypothesis, or identify a counterexample** could diminish. Cognitive standardization here would mean that *not only* do we manipulate the same cultural references, *but also* that we all reason according to the same dominant inductive pattern.



Figure 3.2 – Deductive/logical reasoning versus inductive/algorithmic reasoning.

Societal consequences of such standardization of thought would be profound. A humanity less creative and less diverse in its modes of reasoning could see its innovation stagnate and its intellectual resilience diminish. History has shown that advances often arise from the confrontation of heterodox ideas and varied methodologies. If, on the contrary, the same cognitive approach is universally applied (for example, solving everything by statistical correlation without seeking underlying principles), we may fear a slowdown in progress, greater difficulty in solving novel or complex problems that require *changing the frame* of thought. Moreover, such homogenization could be exploited by certain actors to more easily manipulate opinion: in a uniform intellectual landscape, a single algorithmic narrative can simultaneously influence millions of minds aligned with the same ways of thinking.

Finally, on an ethical level, cognitive standardization raises the question of the **loss of intellectual autonomy**. If our cultural tastes, creative choices, and reasoning methods all converge under the influence of AIs designed by a handful of companies, to what extent do we remain masters of our judgment? Might we not see the emergence of a kind of *cognitive monopoly*, where major AI providers define the contours of accepted rationality, standard creativity, and the "right way" to think? This dystopian scenario is not inevitable, but it marks the warning lines not to cross. The next section will examine precisely which **research and action avenues** could be considered to avoid or mitigate such an impoverishment of human thought in the age of AI.

3.5 Preserving Cognitive Diversity: Issues and Future Directions

In light of the identified risks, a consensus is emerging among experts on the importance of **preserving cognitive diversity** and strengthening critical thinking in the age of artificial intelligence. Rather than rejecting AI technologies outright, the goal is to *proactively design their integration* so that they enhance our abilities without standardizing or atrophying them. Several **actionable avenues** and recommendations are emerging for the coming years, both in research and in educational policies and AI system design.

1. **Strengthen AI and Critical Thinking Education:** Numerous institutional reports stress the urgency of digital and AI education from an early age [16]. This is not just about learning to use these tools, but above all about developing citizens' skills to use them critically. This includes: understanding the basic workings of algorithms (to avoid mystifying them), recognizing AI's biases and limitations, and practicing systematic verification of machine-provided information. Integrating modules of augmented critical thinking into curricula—where students, for example, are confronted with AI-generated texts to assess their reliability—could turn AI from a factor of intellectual passivity into a pedagogical tool for exercising

judgment. Research also calls for working on *critical thinking dispositions* (curiosity, open-mindedness, enlightened skepticism) among students to counterbalance the apparent ease offered by AI [20] [20]. In short, new generations must be equipped to coexist with AI without becoming dependent on it, cultivating what some call "*critical intelligence*" in relation to machines.

2. **Encourage Diversity in AI Design and Training:** On the technology side, it is crucial to **diversify the data and approaches** underlying AI systems. One way to limit induced cognitive standardization is to have **plural and local AIs**. For example, developing large multilingual and culturally adapted models, trained on corpora including varied perspectives (including those from minority languages and cultures), would help reduce the hegemony of a single worldview [16]. France and Europe have a role to play in this regard: by investing in **sovereign AI models rooted in their respective cultural contexts**, they can offer an alternative to global standardizing models [16]. Furthermore, AI research could explore **alternative paradigms** to the dominant **inductive connectionism**. Rehabilitating the integration of symbolic methods, formal logic, or hybrid AIs into systems could maintain a balance between induction and deduction in the tools provided. Similarly, designing recommendation algorithms that optimize not only personalized relevance but also the **diversity of exposed content** is among the technical avenues to *deliberately counterbalance* filter bubbles. Some studies suggest, for example, introducing **controlled randomness** or diversity criteria into information streams to broaden users' horizons beyond their usual preferences [16]. The goal is for technology, instead of merely reflecting our biases, to help open our minds by presenting us with varied viewpoints.
3. **Design "Pro-Cognitive" AI:** Another promising direction is to develop AI systems that, by design, stimulate rather than replace human cognitive activity. For example, AI assistants could be programmed to ask users questions instead of giving direct answers, thus inviting them to think for themselves before receiving machine assistance. Likewise, instead of providing a finalized generated text (which users might passively accept), an AI could offer several different options, or deliberately include *flaws* to be detected, thereby encouraging users' critical thinking. This concept of AI as a catalyst for thought rather than a substitute is being explored in the HCI (Human-Computer Interaction) community [20]. The idea is to avoid the *black box* effect and associated passivity: an AI transparent about its reasoning, justifying its answers, will allow humans to follow a logical path and learn in the process, rather than simply consuming a result. Studies show that with well-designed interfaces, AI can amplify human creativity (by suggesting ideas without imposing them) and enhance critical thinking (by assisting in information verification, for example) [21]. Investing in this type of design aligned with human cognitive values is a major challenge for the future.

4. Pursue Multidisciplinary Research on AI's Cognitive Impact: Finally, it is imperative to **continue to scientifically study** the effects of AI on the brain and cognition, in order to continuously adapt our strategies. The results of the 2025 study on cognitive offloading and critical thinking [19], or the 2024 study on the homogenization of writing styles [17], offer only a first glimpse. Many questions remain: what types of cognitive tasks can be safely delegated to AI, and which must be preserved as *mental exercise*? What are the **usage thresholds** not to be exceeded so that AI remains a support and not a cognitive handicap (statistical analyses, for example, suggest the existence of a threshold beyond which the decline in critical thinking accelerates) [19]? How can we identify the most vulnerable individuals or groups (young people seem more affected, according to some data [19]) in order to target appropriate educational interventions? These questions call for collaborative research among computer scientists, cognitive psychologists, neuroscientists, philosophers, and educators. Initiatives are beginning to emerge, but a **sustained and interdisciplinary effort** will be necessary to guide society on the best way to co-evolve with AI without losing what makes human thought unique and rich.

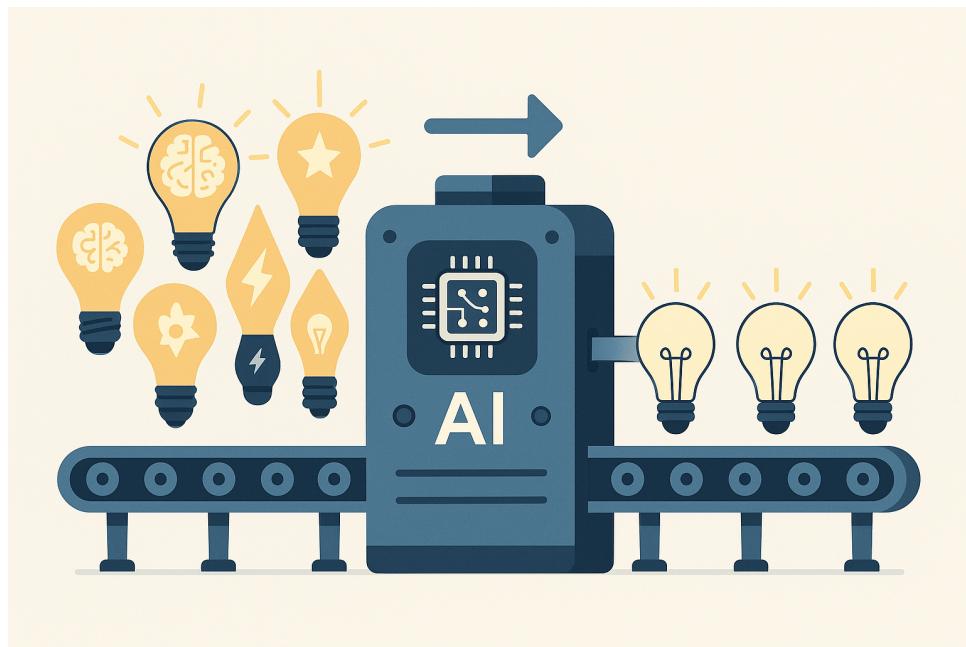


Figure 3.3 – The standardization of ideas as a production line, a risk of AI.

3.6 Conclusion

In conclusion to this chapter, **cognitive standardization** appears as a real but surmountable challenge of the AI era. Far from being inevitable, it is a *call to action* to steer technological development and human practices in a direction that values intellectual

diversity. AI can certainly, albeit unintentionally, push toward the standardization of thought, but it can also—if we so choose—be harnessed in the service of **augmented, plural, and critical thinking**. The key is to recognize warning signs in time (decline in certain skills, reinforced biases, impoverishing convergence of ideas) and to respond with appropriate educational, technical, and ethical innovations. Preserving **cognitive diversity** in the age of AI ultimately means preserving humanity’s ability to renew itself, to innovate, and to understand the world in multiple ways. It is an ambitious project, one that will require the mobilization not only of researchers and policymakers, but of every AI user in their daily practice, in order to disprove the prophecy of uniform thought and instead build a future where humans and artificial intelligences co-evolve for the best of creativity and reason.

Chapter 4

Mechanisms of Manipulation by Artificial Intelligence

4.1 Introduction and Definition

The development of artificial intelligence (AI) is accompanied by an unprecedented ability to influence and steer human behavior in subtle and automated ways. **Manipulation by AI** can be defined as any influence exerted through digital technologies, intentionally designed to bypass the individual's reasoning and create an asymmetry of outcomes between the actor using AI and the targeted person [22]. In other words, AI enables companies, platforms, or malicious actors to influence users' decisions without their knowledge, often **without transparency or informed consent**, raising major ethical concerns regarding respect for individuals' cognitive autonomy [22].

Historically, persuasion and marketing techniques already existed, but **AI amplifies their reach and effectiveness**. By combining machine learning algorithms with vast amounts of personal data, strategists can now **target users individually with manipulative techniques of unprecedented efficiency and discretion** [23]. For example, large platforms with millions of users (Google, Facebook, etc.) often know our preferences better than our own relatives, by analyzing every click, every "Like," and every search [23]. A study published in PNAS showed that simple data such as Facebook "Likes" can predict with **surprising accuracy** sensitive personal traits (political or sexual orientation, personality traits, intelligence level, etc.) [23] [24]. This automated profiling capability, derived from our digital footprints, lays the groundwork for hyper-personalized manipulation: by intimately knowing a target's values, biases, and emotions, AI can optimally tailor messages to influence their judgment.

The **risks of manipulation by AI** affect many domains (politics, consumption, health, etc.) and take various forms that this chapter aims to map. We will begin by presenting a **taxonomy of the main mechanisms** of algorithmic manipulation (section 4.2). Next, we will analyze several of these mechanisms in depth: the exploitation of

human cognitive biases by AI (section 4.3), algorithmic personalization leading to **filter bubbles** and information polarization (section 4.4), as well as the creation of **disinformation and fake content** via AI (section 4.5). We will also address how AI can exploit **social interactions** (bots, fake profiles) and digital interfaces to subtly steer choices (for example, through **dark patterns**), while discussing future developments and possible safeguards (section 4.6). The objective is to provide a rigorous overview of the methods by which AI can manipulate human thought, drawing on recent scientific work and concrete examples.

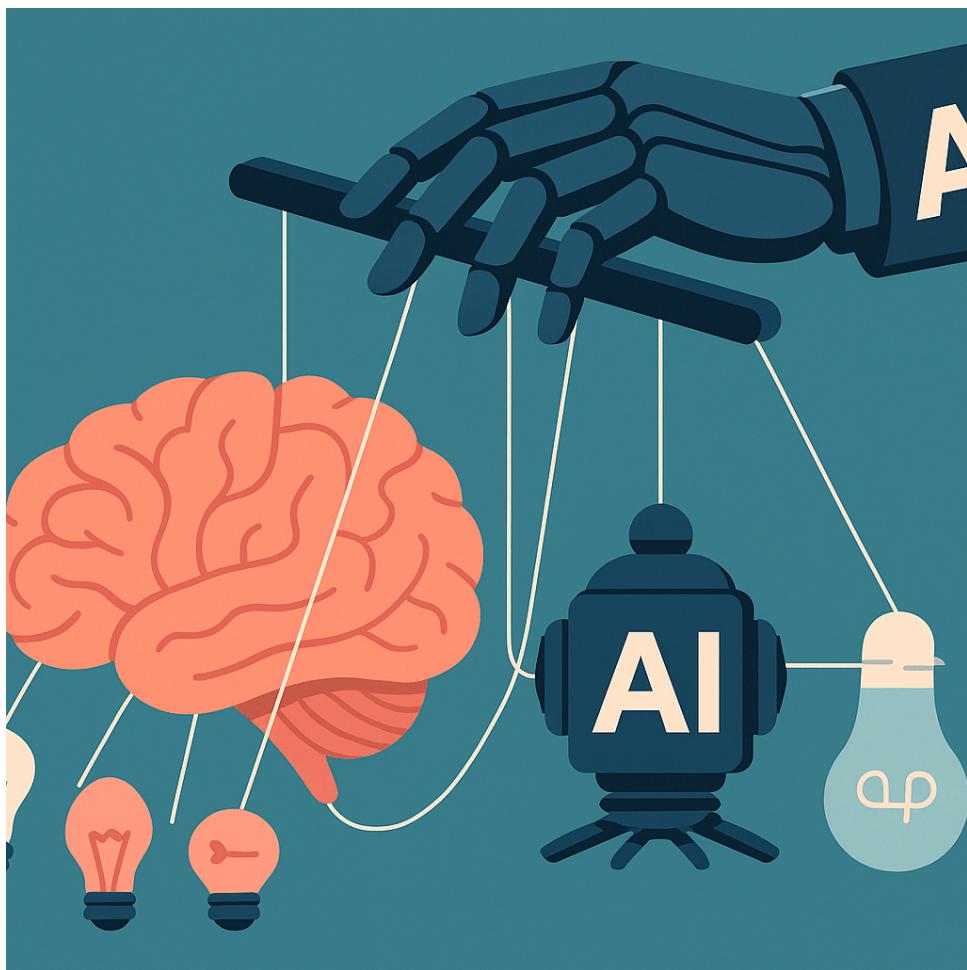


Figure 4.1 – Conceptual representation of manipulation by AI.

4.2 Taxonomy of AI Manipulation Mechanisms

Several **categories of mechanisms** emerge in the ways AI can manipulate individuals. **Table 4.1** below offers a taxonomy of the main forms of AI-driven manipulation, classifying them by nature and providing illustrative examples for each. This classification, inspired by existing research [23] [23], shows that manipulation can take **multiple and often combined forms**: exploiting human cognitive flaws, personalized content

shaping, generating indistinguishable fake elements, or using AI to simulate deceptive social interactions. Each category relies on distinct techniques, but all share the feature of **altering the target's decision-making process** to the manipulator's benefit.

Table 4.1: Taxonomy of AI Manipulation Mechanisms.

Mechanism Category	Description	Concrete Examples
Exploitation of Cognitive Biases (Hypernudges)	AI detects and exploits the subject's psychological biases to influence their choices.	Content reinforcing a user's confirmation bias ; recommendation systems leveraging aversion to contradiction by only showing similar opinions.
Algorithmic Personalization (Personalized Filtering)	AI modulates the information presented based on the user's profile, creating a tailor-made filter bubble .	Social media news feeds sorted by algorithm, locking the user in an ideological echo chamber ; micro-targeted ads tailored to personality (introvert/extrovert).
Emotional Manipulation (Affective Content)	AI maximizes engagement by playing on the subject's emotions and affective state at the opportune moment.	Algorithms amplifying divisive or anxiety-inducing content to provoke anger or fear (and thus attention); commercial offers sent when the user is emotionally vulnerable (e.g., junk food promotion to a depressed person).
Automated Disinformation (Generative AI)	AI creates fake content (text, image, video) indistinguishable from real, misleading recipients.	Deepfake videos showing a public figure in a fabricated situation; fake news written by AI and massively spread on social networks.

Mechanism Category	Description	Concrete Examples
Simulated Social Influence (Bots and Fake Agents)	AI poses as human participants to create an illusion of consensus or trust.	Social bots posting positive comments for a product (fake reviews); automated accounts artificially boosting the popularity of an idea or hashtag to make it appear trending.
Persuasive Design (Dynamic Dark Patterns)	AI optimizes the user interface in real time to push for specific actions, often to the user's detriment.	E-commerce sites adjusting prices based on vulnerability (e.g., higher price if the buyer's smartphone battery is low) [23]; messages prompting default acceptance of options benefiting the platform (consents, hidden subscriptions).

This taxonomy highlights that AI manipulation methods often combine big data and behavioral science knowledge. For example, manipulation can rely on known cognitive biases (systematic judgment errors in humans) and on algorithmic personalization to exploit these biases in a targeted way at scale. The following sections (4.3 to 4.6) explore these mechanisms in detail, illustrating them with research and case studies. Note that these categories are not mutually exclusive: in practice, several techniques can be combined. An online disinformation campaign may thus use both deepfakes (generative disinformation), bots to spread the content (simulated social influence), and micro-targeting of the most receptive individuals (bias exploitation via personalization). The common thread of all these approaches is the use of AI to amplify intentional influence on human behavior while making this influence less detectable.

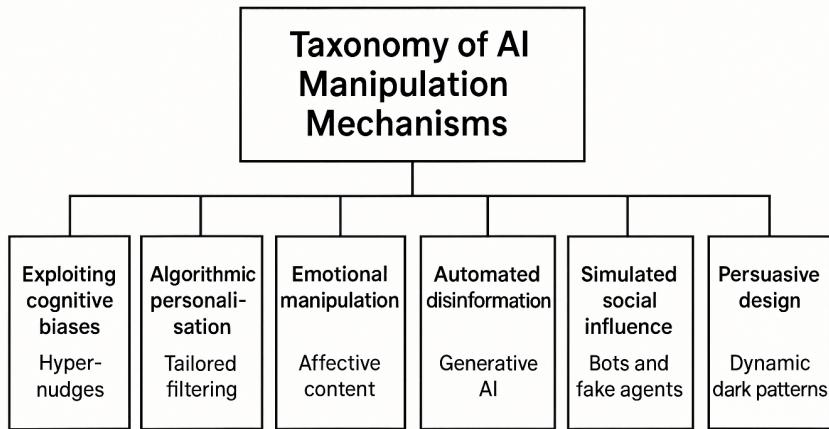


Figure 4.2 – AI manipulation taxonomy diagram.

4.3 Exploitation of Cognitive Biases and Psychological Vulnerabilities

Humans have natural **cognitive biases**—that is, tendencies to deviate from rationality in their judgments—which AI systems can detect and exploit at scale. Modern algorithms, by analyzing our data (clicks, histories, preferences), manage to **identify our biases and personality traits**. This allows them to **adapt content or offers to resonate with these predispositions**, thereby intensifying persuasive impact.

A striking example is the use of **psychological profiling** in advertising and political communication. Researchers have shown that it is possible to **accurately predict an individual's psychological profile from their digital footprints** (e.g., Facebook "Likes"), then tailor persuasive messages accordingly [24] [25]. In a series of three field experiments involving 3.5 million internet users, **ads whose style was tailored to the targets' personality traits generated up to 40% more clicks and 50% more purchases** compared to non-personalized messages [25]. In other words, by exploiting a psychological bias or preference (for example, presenting a product in an extroverted way to an extrovert), AI can **significantly alter purchasing behavior**. These results confirm that **psychological microtargeting**—made famous in the wake of the Cambridge Analytica scandal—is a **formidably effective lever of manipulation**, capable of influencing the attitudes of a vast audience when well calibrated [25].

Social media platforms also use these principles to maximize engagement. Their algorithms learn the **stimuli to which each user is most sensitive**—for example,

content that confirms their existing opinions (exploiting confirmation bias) or that elicits a strong emotional reaction. By exploiting such biases, AI can **progressively amplify the user's inclinations**. Recent research has highlighted a worrying **feedback loop effect**: when humans **repeatedly interact with a biased AI system, they themselves become more biased** in their judgments over time [26]. In the lab, it is observed that even **slight initial algorithmic biases can be internalized by users**, snowballing with each interaction [26]. This phenomenon is **more pronounced than in equivalent human interactions**, as AIs present their judgments with an appearance of objectivity and consistency that makes them particularly influential [26]. Indeed, an AI system can detect and exploit **tiny biased correlations** in data thanks to its computational power, and provide recommendations with an **apparently more reliable signal** (less "noise") than a human opinion [26]. Users, often perceiving AI as a technical authority superior to humans, tend to **follow its suggestions without assessing their bias**—a rational behavior if one believes AI is infallible [26]. This mechanism explains how AI can **amplify a pre-existing cognitive bias**: if the algorithm itself is biased or oriented (e.g., a YouTube recommendation engine favoring conspiratorial content), the user, trusting it, will increasingly adopt these biases.

In sum, **algorithmic exploitation of cognitive biases** relies on AI's ability to **learn our individual psychological weaknesses** and then use them to steer our decisions toward a given goal (purchase, vote, adherence to an idea, etc.). This can take the form of sophisticated **nudging** techniques (behavioral prodding) automated by AI and big data, sometimes called **hypernudges**. The literature identifies four key characteristics of these digital manipulations: **intentionality** (they are deliberately orchestrated), **asymmetry** of knowledge and benefit (the manipulator profits at the user's expense), **opacity** (the influence is not transparent to the target), and **infringement of autonomy** (the person's free decision-making capacity is eroded) [22]. AI does not create new human biases, but it offers manipulators a **unprecedented means of exploiting existing ones**, in a targeted and large-scale manner. This reality calls for reflection on new rights (e.g., **right to cognitive liberty**) and ethical safeguards to protect individuals, a point to which we will return in section 4.6 [22].

Moreover, it should be noted that the reverse of this manipulative power exists: if **the algorithm is accurate and unbiased**, it can also **correct** human errors. Studies have shown that when people interact with a truly competent and objective AI, their judgments can improve thanks to the AI's advice [26]. The danger therefore lies in the **information imbalance**: the average user has no way of knowing whether the AI they consult is reliable or subtly biased. This default trust in AI illustrates another exploited cognitive bias, known as **automation bias**—the tendency to place excessive trust in the recommendations of an automated system simply because it is perceived as such [27]. Thus, a user will readily follow the route suggested by their GPS even if it seems counterintuitive, or accept site rankings without question, at the risk of making serious

mistakes. AI can exploit this automation bias to push its messages or recommendations with **minimal critical thinking in response**.

To better understand the interplay between **human cognitive biases** and **AI technologies**, **Table 4.2** presents some common biases and how algorithmic systems can activate them to influence behavior.

Table 4.2: Cognitive Biases Exploited by AI Technologies.

Human Cognitive Bias	Bias Description	AI Exploitation
Confirmation Bias	Tendency to favor information that confirms our pre-existing beliefs, ignoring information that contradicts them.	Social media algorithms learn the user's opinions and mainly present content aligned with them, reinforcing their convictions and increasing engagement (likes, shares). This keeps the user in a comfortable information bubble [28].
Authority / Automation Bias	Propensity to give excessive credit to recommendations from a source perceived as authoritative or from an automated system.	Virtual assistants and AI systems are seen as experts: users tend to blindly follow AI advice (GPS route, purchase suggestion) without double-checking. Unscrupulous designers can use this to subtly steer choices (highlighted products, etc.) [27].

Human Cognitive Bias	Bias Description	AI Exploitation
Attention to Emotional Stimuli Bias	Information that elicits strong emotion (fear, anger, joy) attracts more attention and is judged more important.	Recommendation algorithms detect which content provokes strong reactions in the user (outrageous videos, shocking news) and systematically push this type of content to capture their attention. They exploit emotional sensitivity to keep the user active on the platform.
Scarcity Effect (FOMO)	An offer seems more attractive if presented as limited or exclusive, creating the fear of missing out (Fear of Missing Out).	E-commerce sites, via AI-optimized dark patterns , display fake counters ("only 2 items left," "offer valid for 24h") to push the user to make an impulsive purchase . AI can dynamically adjust these signals based on the buyer's profile (if they are sensitive to scarcity).
Social Proof (Herd Effect)	We tend to adopt an opinion or behavior if we believe "many other people" are doing the same.	Automated bots can simulate a crowd of positive reviews (fake comments, fake followers) around a product or idea, creating an illusion of popularity that encourages real users to follow suit [23]. Studies have shown that on Twitter, a very small percentage of automated accounts is enough to massively spread disinformation by exploiting this group effect [29].

As illustrated by Table 4.2, **detailed knowledge of human biases** allows AI designers to optimize their systems to trigger these biases at will. The impact ranges from **commercial influence** (better selling a product by adapting advertising to the client's psychology) to **ideological influence** (gradually shaping an individual's political opinion by only showing one point of view). The next section (4.4) will detail how algorithmic personalization of information streams notably contributes to locking users into **information echo chambers**, reinforcing confirmation biases and polarizing societies.

4.4 Algorithmic Personalization, Filter Bubbles, and Polarizing Content

One of the most studied manipulation mechanisms concerns how algorithms filter and personalize the information we see online. On social networks, video platforms, or even search engines, AI systems decide **which content to prioritize for each user**, aiming to optimize certain criteria (most often, engagement or time spent). This algorithmic filtering leads to the creation of "**filter bubbles**", a concept popularized by Eli Pariser: each user is immersed in a **customized informational environment**, reflecting their preferences and avoiding content likely to make them leave [28]. If, for example, a user habitually reads articles with a certain political slant, their news feed algorithm will mostly show posts consistent with that orientation, **reinforcing their confirmation bias**. Over time, this process creates **echo chambers** where the individual hears only opinions similar to their own, which can polarize their positions.

From the platform's perspective, this personalization is rational: by exposing the user to what they **want** to see (or what emotionally captivates them), attention is maximized, and thus associated advertising revenue. However, from a societal and individual perspective, the perverse effects are significant. People living in these bubbles perceive a distorted reality—**everyone thinks like me, no information contradicts my ideas**—which can diminish critical thinking and exacerbate extremism in some cases. Studies have shown that algorithms on platforms like YouTube or Facebook tend to amplify the virality of controversial or extreme content, as these provoke more reactions and thus engagement. In the absence of safeguards, a user can be gradually led, video by video, toward increasingly radical positions, a phenomenon sometimes described as the YouTube "**rabbit hole**."

However, the exact role of algorithms in polarization should be nuanced: recent studies offer mixed results, some suggesting that **users' personal choices also play a role** (we naturally tend to surround ourselves with like-minded people). Nevertheless, even if AI is not the sole culprit of polarization, its **opacity** makes the problem complex. Recommendation criteria are often secret, preventing users from realizing they are trapped in partial information filtering. **Lack of transparency benefits manipulation**

strategies: as the Bruegel think tank report indicates, the lack of visibility on algorithmic objectives and data use allows them to steer behaviors without our awareness [23].

A striking example is the **massive experiment conducted by Facebook** in 2012 on nearly 700,000 users without their explicit consent. For a week, Facebook altered these users' news feeds: some saw more positive posts, others more negative ones. The results confirmed a large-scale emotional contagion effect: **people exposed to more negative messages in turn posted more negative content, and vice versa for those exposed to positive messages** [30]. This study, published in PNAS in 2014, demonstrated that simply modulating the mood of the news feed algorithmically could **alter users' emotional states without their knowledge**. Although Facebook justified the experiment as a way to improve the service, it sparked a major ethical controversy when revealed publicly [31]. It illustrates the **power of personalization algorithms to subtly manipulate the collective psyche**, here by inducing particular emotions for experimental reasons—a practice many equate with clandestine manipulation.

Beyond this extreme case, algorithmic content selection acts daily as a form of soft manipulation. Every notification pushed to your screen, every search result order, can influence your actions: reading a particular article, buying a product, feeling a certain emotion. These choices are not neutral—they follow the objectives set for the AI by its designers (often commercial). This is called choice architecture: AI shapes the environment in which the user makes decisions, highlighting certain options and hiding others. For example, on a travel booking site, the algorithm may manipulate the order of hotel listings (placing those maximizing platform profit at the top), or default to the option including paid insurance (betting that the user will follow the default choice). All these persuasive **design** techniques existed before AI, but machine learning makes them much more effective by adapting them in real time to each individual.

Thus, algorithmic personalization creates an **invisible manipulative environment**: everyone navigates their own version of the web, calibrated to steer their clicks and behavior in ways profitable to platforms. For the isolated user, it is difficult to realize the extent of this manipulation since they only see their filtered version of the online world. Only by comparing with others (or through external audits of algorithms) does the **selectivity of presented content** become apparent. Aware of these issues, regulators are beginning to demand more **algorithmic transparency**. For example, the European Digital Services Act (DSA) requires large platforms to allow users to disable personalization of recommended content. However, even with such measures, the **commercial appeal of personalized filtering** ensures that this mechanism will remain central and continue to be refined. The next section will examine another critical aspect of AI manipulation: the generation of **false information and fake content** (text, images, videos) that can deceive users about the truth of the world around them.

4.5 Disinformation, Deepfakes, and the Automation of Deception

Recent advances in AI, particularly in content generation (**generative AI**), have given rise to a new type of manipulative threat: **automated and undetectable disinformation**. This involves using algorithms to produce entirely fabricated texts, images, audio, or videos **more real than real**, with the aim of deceiving or influencing opinion. Unlike manipulation by content selection (section 4.4), here **new false information** is created to steer the beliefs or decisions of targets.

4.5.1 AI-Driven Fake News Propagation

On the web and social networks, AI can be used to **write and massively disseminate fake news** with unprecedented reach. Advanced language models can automatically produce full articles, imitating journalistic style, conveying disinformation. Coupled with bots (automated accounts), this content can be widely shared, simulating the appearance of popular enthusiasm. Research has shown that a **small percentage of well-orchestrated bots can vastly amplify the reach of false information online**. For example, a study published in Nature Communications showed that **only 6% of Twitter accounts (identified as bots) were responsible for about 31% of non-credible information on the network** during an analyzed election period [29]. These bots tirelessly post and repost the same false links, creating the illusion that these stories are widely shared and newsworthy, deceiving human users. Moreover, they act very quickly when a hoax appears, flooding the public space before corrections or denials can be issued [29]. This automation gives rumor spreaders a clear advantage over authorities or traditional media, which operate manually and more slowly.

Bot networks can also use other tactics to multiply their manipulative impact: for example, mentioning or directly targeting influential accounts (journalists, public figures) so that they unwittingly relay the fake information; or flooding legitimate discussions to drown out correction messages. Here, AI is the weapon that enables the industrial-scale creation of **false consensus or false trends**. Under the guise of spontaneous "buzz," it is actually a **well-tuned piano** where each bot plays its part to **impose a biased narrative**.

4.5.2 Deepfakes: Visual and Audio Hyperfakes

The particular case of **deepfakes** deserves special attention. This term (a contraction of **deep learning** and **fake**) refers to AI-generated audiovisual content that almost perfectly mimics reality. Neural networks, especially GANs (**Generative Adversarial Networks**), can **swap a person's face in a video or synthesize someone's voice** from

a few recordings. While these techniques have legitimate applications (special effects, dubbing), they can also serve extremely pernicious manipulative purposes, as video and audio have historically been seen as tangible evidence.

Imagine a video where a leader is seen and heard announcing a shocking measure—for example, a president declaring withdrawal from a conflict—when this never happened. Such a deepfake, spread without context, can cause panic or confusion before authorities can deny it. This scenario is not science fiction: in March 2022, during Russia's invasion of Ukraine, **a fake video of Ukrainian President Volodymyr Zelensky calling on his troops to lay down their arms was posted online** via a hacked Ukrainian news site [32]. Although the fake was of average quality (artificial voice with a strange accent, imperfect face cutout), it managed to bypass some moderation barriers and briefly spread on social networks before being flagged and removed [32]. Zelensky himself had to urgently post an authentic video to deny this fictitious call for surrender. As a cybersecurity expert noted, this **first "effective" wartime deepfake may be just the tip of the iceberg** [32]. Malicious actors are sharpening these tools and could produce increasingly convincing fakes, capable of **mass disinformation or destabilizing nations** by undermining trust in visual information.

Beyond the geopolitical sphere, deepfakes also pose a risk to individuals. **Phone scams using AI-cloned voices of loved ones in distress to demand urgent financial help (fake kidnapping)** have already been reported. In 2023, a mother in Arizona received a call where she heard the frantic cries of her supposedly kidnapped 15-year-old daughter, followed by a ransom demand—all fake, orchestrated by AI cloning the teen's voice from online videos [33] [33]. This type of **virtual kidnapping** shows how AI can exploit **emotion and credulity** by abusing the trust we place in recognizing our loved ones' voices. Again, the manipulation aims to **bypass the victim's analytical reasoning** (who might have wondered why the number was unknown, etc.) by triggering intense stress through a convincing voice simulation.

The challenge posed by these **hyperfakes** is twofold: on the one hand, **they blur the line between true and false** (it becomes difficult to trust audiovisual evidence), and on the other, **they can be used to deny reality**—sometimes called the "**liar's dividend**": once deepfakes exist, a person caught in a real compromising video can claim it is an AI-generated fake, sowing doubt. Thus, even without active use, the mere awareness of these tools' existence weakens the authority of visual evidence and enables all kinds of narrative manipulation.

4.5.3 From Information Manipulation to Manipulation of Perceived Reality

AI manipulation doesn't stop at content filtering or creating false news. It can also act on **how we perceive reality itself** by modifying our cognitive patterns or emotions. This

type of intervention is more subtle but potentially more powerful, as it can **restructure our mental framework** without us being aware of it.

One form of this manipulation is through **emotional conditioning via algorithms**. When an AI system repeatedly presents content that elicits specific emotions (fear, anger, euphoria) in response to certain subjects, it can gradually **associate these emotions with the targeted concepts** in our minds. For example, if a news aggregation algorithm consistently shows alarming articles when a particular political figure is mentioned, the user may unconsciously develop a negative emotional response to that person, regardless of the factual content. This phenomenon resembles classical conditioning techniques but is applied **systematically and at scale** through personalized algorithms.

Studies in neuroscience have shown that **repeated exposure to emotionally charged content can modify brain structure and responses**. When users are regularly exposed to anxiety-inducing or anger-provoking content, their brains can develop heightened sensitivity to these emotions, making them more susceptible to manipulation [26]. This creates a **feedback loop** where users become increasingly reactive to certain stimuli, and algorithms can exploit this heightened reactivity to further influence behavior.

Another dimension is the manipulation of **attention and focus**. AI systems can deliberately scatter our attention by presenting information in fragments, creating what researchers call "continuous partial attention." This fragmented information processing makes it harder for users to form coherent, critical thoughts about complex issues. Instead of engaging in deep reflection, users develop **surface-level reactions** to isolated pieces of information, making them more susceptible to emotional manipulation and less capable of rational analysis.

The timing of information presentation also becomes a tool of manipulation. AI systems can learn when users are most vulnerable—for instance, late at night when critical thinking abilities are diminished, or during stressful periods when emotional defenses are lowered. By **strategically timing the delivery of persuasive content**, these systems can maximize their manipulative impact [22].

Perhaps most concerning is the potential for AI to create **synthetic experiences** that feel authentic but are entirely manufactured. Virtual and augmented reality technologies, combined with AI, can create immersive experiences that are indistinguishable from reality. Users might "experience" events that never happened, meet people who don't exist, or witness scenes that are entirely fabricated. These synthetic experiences can form **false memories** and influence future behavior as if they were real experiences.

This manipulation of perceived reality represents a fundamental shift from traditional propaganda, which sought to convince people of certain ideas, to **reality engineering**, which seeks to alter the very foundation of what people consider real. The implications for human autonomy and decision-making are profound, as individuals may base their choices on a reality that has been systematically distorted by AI systems designed to serve interests other than their own.

4.6 Future Perspectives and Ethical Safeguards

As AI manipulation techniques become increasingly sophisticated and widespread, society faces unprecedented challenges in preserving human autonomy and cognitive liberty. The mechanisms described in this chapter are not merely theoretical concerns—they are **already being deployed at scale** and will likely become more powerful as AI technologies advance. This section examines potential futures and explores ethical safeguards that could help mitigate these risks.

The trajectory of AI manipulation appears to be moving toward what some researchers call "**persuasive singularity**"—a point where AI systems become so effective at understanding and influencing human psychology that resistance becomes nearly impossible for the average person. Unlike the technological singularity, which focuses on AI surpassing human intelligence, the persuasive singularity concerns AI's ability to **override human decision-making processes** through psychological manipulation [22].

Several technological trends suggest this future may be approaching rapidly. First, **brain-computer interfaces** are advancing toward more direct access to neural activity, potentially allowing AI systems to monitor and influence thoughts at their source. Second, the integration of AI with **ubiquitous computing**—from smart homes to wearable devices—creates opportunities for continuous, context-aware influence. Third, advances in **real-time deepfake generation** and synthetic media creation will make it increasingly difficult to distinguish authentic from manipulated content.

However, this dystopian trajectory is not inevitable. Researchers and policymakers are developing several categories of safeguards to protect human cognitive autonomy. **Technical safeguards** include algorithmic transparency requirements, manipulation detection systems, and "cognitive firewalls" that could help users identify and resist psychological manipulation attempts. Some platforms are experimenting with **friction-based design**—introducing deliberate delays or confirmation steps before users can share potentially false information or make impulse purchases influenced by AI manipulation.

Legal and regulatory frameworks are also emerging. The European Union's Digital Services Act requires large platforms to provide algorithmic transparency and allow users to opt out of personalized recommendations. Some jurisdictions are considering "cognitive rights" legislation that would establish a fundamental right to mental self-determination, similar to existing privacy rights. The concept of **neurorights**—legal protections for mental processes—is gaining traction, with some experts proposing constitutional amendments to protect cognitive liberty [22].

Educational approaches represent another crucial defense. Digital literacy programs are expanding beyond traditional computer skills to include "cognitive security" training that helps individuals recognize and resist manipulation attempts. Some educational initiatives focus on strengthening critical thinking skills specifically in digital environ-

ments, teaching people to question the sources and motivations behind the content they encounter online.

However, these safeguards face significant challenges. The **asymmetry of resources** between manipulators and their targets means that well-funded actors will likely stay ahead of defensive measures. The global nature of digital platforms makes regulatory enforcement difficult, as companies can simply relocate to jurisdictions with fewer restrictions. Moreover, many manipulation techniques operate below the threshold of conscious awareness, making them difficult for users to detect even with training.

Perhaps most concerning is the **economic incentive structure** that drives AI manipulation. As long as attention-based business models dominate the digital economy, platforms will have financial incentives to maximize user engagement through whatever means are most effective, including psychological manipulation. Addressing this may require fundamental changes to how digital services are funded and operated.

The development of **AI alignment** technologies offers some hope. Research into creating AI systems that genuinely serve human interests, rather than merely appearing to do so, could help ensure that future AI systems are designed to enhance rather than exploit human cognition. This includes work on value alignment, interpretable AI, and systems that actively protect user autonomy rather than undermining it.

International cooperation will be essential for addressing AI manipulation effectively. Just as climate change requires global coordination, the challenge of preserving human cognitive autonomy in an AI-dominated information environment will require unprecedented international collaboration. This might include treaties governing AI manipulation, shared standards for algorithmic transparency, and coordinated responses to state-sponsored disinformation campaigns.

The next section will examine specific examples of how these manipulation techniques are being deployed across different application domains, illustrating the breadth and diversity of AI manipulation in contemporary society.

Table 4.3: Examples of AI Manipulation by Application Domain.

Application Domain	Manipulation Technique	Concrete Example
E-commerce	Psychological profiling for targeted advertising	Amazon's recommendation system analyzes purchase history, browsing behavior, and demographic data to present personalized product suggestions that exploit individual psychological profiles, increasing purchase likelihood by up to 40% [25].
Social Media	Filter bubbles and echo chambers	Facebook's News Feed algorithm creates personalized information environments that reinforce users' existing beliefs, potentially contributing to political polarization and reduced exposure to diverse viewpoints [28].
Political Campaigns	Microtargeted political advertising	Cambridge Analytica's use of psychological profiling to deliver personalized political messages to millions of users during the 2016 elections, demonstrating how AI can be weaponized for electoral manipulation.
Financial Services	Behavioral nudging for financial decisions	AI-powered investment platforms use behavioral economics principles to encourage specific investment choices, potentially leading users toward higher-fee products that benefit the platform more than the investor.

Application Domain	Manipulation Technique	Concrete Example
Dating Apps	Emotional manipulation through scarcity	Dating applications use AI to control the timing and frequency of matches, creating artificial scarcity to increase user engagement and premium subscription purchases through psychological manipulation of romantic desires.
Gaming Industry	Addiction-inducing reward systems	Video game AI systems analyze player behavior to optimize reward schedules and monetization strategies, using variable ratio reinforcement to create gambling-like addiction patterns, particularly targeting vulnerable populations.
News Media	Emotional contagion through algorithmic curation	News aggregation algorithms prioritize emotionally provocative content (anger, fear, outrage) to maximize engagement, potentially contributing to societal anxiety and emotional polarization [30].
Healthcare	Manipulation of health-related decisions	AI chatbots designed to provide health advice may be programmed to subtly promote certain treatments, medications, or healthcare providers based on commercial partnerships rather than purely medical considerations.
Education Technology	Cognitive dependency creation	Educational AI systems may be designed to create dependency on the platform rather than fostering independent learning skills, potentially contributing to cognitive atrophy in students [3].

Application Domain	Manipulation Technique	Concrete Example
Voice Assistants	Automation bias exploitation	Smart speakers and voice assistants leverage users' tendency to trust automated systems, potentially influencing product recommendations, news consumption, and daily decisions through seemingly neutral responses.
Ride-sharing Services	Dynamic pricing manipulation	Companies like Uber use AI to analyze user behavior and implement surge pricing at moments of high emotional stress or limited alternatives, exploiting users' psychological vulnerabilities for profit maximization [23].
Content Creation	Deepfake disinformation campaigns	State and non-state actors use AI-generated deepfake videos and audio to spread disinformation, manipulate public opinion, and undermine trust in authentic media, as demonstrated in recent geopolitical conflicts [32].

Table 4.3 illustrates the pervasive nature of AI manipulation across virtually every sector of digital society. These examples demonstrate that AI manipulation is not a future concern but a present reality affecting millions of people daily. The sophistication and scale of these techniques continue to evolve, making it increasingly urgent to develop effective countermeasures and ethical frameworks to protect human cognitive autonomy.

The breadth of applications shown in this table also highlights the challenge facing regulators and ethicists: AI manipulation techniques are not confined to obvious domains like advertising or politics, but have infiltrated sectors traditionally viewed as neutral or beneficial, such as education and healthcare. This pervasive nature makes it difficult for individuals to recognize when they are being manipulated and underscores the need for systemic solutions rather than ad-hoc responses to specific incidents.

4.7 Conclusion

This chapter has explored the sophisticated mechanisms through which AI systems can manipulate human behavior and cognition, revealing a landscape of influence that extends far beyond traditional notions of propaganda or advertising. From the exploitation of cognitive biases and the creation of filter bubbles to the generation of deepfakes and the manipulation of perceived reality itself, AI-driven manipulation represents **a fundamental shift in the nature of influence in human society**.

The taxonomy presented in this chapter demonstrates that AI manipulation operates across multiple dimensions simultaneously. **Technical sophistication** allows these systems to process vast amounts of personal data and adapt their strategies in real-time to individual psychological profiles. **Psychological targeting** enables unprecedented precision in exploiting human cognitive vulnerabilities. **Scale and automation** make it possible to influence millions of people simultaneously with personalized approaches. Finally, the **opacity** of these systems ensures that most manipulation occurs below the threshold of conscious awareness.

The examples examined throughout this chapter—from Cambridge Analytica's political microtargeting to Facebook's emotional contagion experiments, from deepfake disinformation campaigns to AI-powered voice cloning scams—illustrate that these techniques are not theoretical possibilities but present realities actively shaping human behavior on a global scale. The pervasive nature of AI manipulation across domains ranging from e-commerce to healthcare, from education to entertainment, suggests that **no aspect of human experience in the digital age remains untouched by these influences**.

Perhaps most concerning is the trajectory toward what researchers term the "persuasive singularity"—a point where AI systems become so effective at psychological manipulation that human resistance becomes nearly impossible. The convergence of brain-computer interfaces, ubiquitous computing, and increasingly sophisticated deepfake technologies suggests that the manipulative power of AI will only intensify in the coming years.

However, this chapter has also highlighted that this dystopian future is not inevitable. **Technical safeguards**, including algorithmic transparency requirements and manipulation detection systems, offer some protection. **Legal and regulatory frameworks**, such as the European Union's Digital Services Act and emerging "cognitive rights" legislation, provide structural defenses. **Educational approaches** that emphasize digital literacy and cognitive security can help individuals recognize and resist manipulation attempts.

Yet the fundamental challenge remains the **asymmetry of power** between those who control AI systems and those who are subject to their influence. This asymmetry is not merely technical but economic, informational, and ultimately political. Addressing AI manipulation will require not just better technology or education, but fundamental

changes to the economic models that incentivize such manipulation and the political structures that enable it to flourish unchecked.

The implications extend beyond individual autonomy to the very foundations of democratic society. If human decision-making can be systematically influenced by AI systems operating at unprecedented scale, then the assumption of informed consent that underlies democratic governance is called into question. The manipulation techniques described in this chapter threaten not just individual freedom but the collective ability of societies to make rational choices about their future.

As we move forward into an increasingly AI-dominated information environment, the preservation of human cognitive autonomy emerges as one of the defining challenges of our time. The stakes could not be higher: at issue is nothing less than the capacity for authentic human thought and genuine democratic deliberation in the digital age. The next chapter will examine how these concerns about AI manipulation intersect with broader questions about artificial consciousness and the future of human identity in relation to increasingly sophisticated AI systems.

Chapter 5

The Question of AI Consciousness

5.1 Artificial Consciousness: Definitions and Issues

We cannot address the subject of thought without discussing the issue of AI consciousness.

Consciousness is classically defined in philosophy of mind as the capacity to have a subjective experience—what is called phenomenal consciousness or sentience, that is, the ability to feel *qualia* (subjective sensations, such as the perception of colors or pain) [34]. This phenomenal dimension is often distinguished from **access consciousness**, understood as the availability of information to guide behavior and reasoning in a global manner [34]. The question of **artificial consciousness** consists in asking whether machines or computer programs could one day exhibit not only intelligence or advanced behaviors, but also a form of subjective experience similar to that of human beings. In other words, beyond processing information in a sophisticated way, could an AI system “feel” something and be aware of itself and the world? This issue is attracting growing interest as AI capabilities advance. Indeed, the recent rise of generative AI models and large language models has made the question more concrete: some observers believe that with such progress, the reproduction of a form of human consciousness by a machine is becoming conceivable [35].

Yet, to date, there is no consensus on the possibility of artificial consciousness, nor even on the criteria for identifying it. Intense debates animate the scientific, philosophical, and public communities on this complex subject [35]. A media example of these debates is the Blake Lemoine affair, the Google engineer who claimed in 2022 that the language model LaMDA was, in his view, sentient, that is, endowed with a consciousness comparable to that of a human—a claim strongly contested by his employer, who deemed it “totally unfounded” [35]. This case illustrates the difficulty of distinguishing intelligently simulated behavior from possible real consciousness: can an AI simply feign consciousness by giving the illusion of thoughts and emotions? Or is there an “inner fact” that could emerge in the machine?

Researchers emphasize that it is crucial to clearly differentiate **artificial intelligence**—the ability of a machine to solve problems or converse coherently—from **consciousness** understood as lived experience. An AI may appear to converse intelligently without actually experiencing anything. As Hsing (2023) notes, a modern computer program, however powerful, is ultimately just a symbol manipulator devoid of semantic understanding: it applies formal rules without *intentionality* (that is, without referring its symbols to real-world meanings) and without qualia, thus without real subjective sensation [35]. From this perspective, current machines, including the most advanced, merely *simulate* understanding and have no intrinsic “meaning” to their operations. This view echoes the classic philosophical argument of Searle’s **Chinese Room**, according to which correctly manipulating symbols (for example, sentences in Chinese) is not sufficient to understand their meaning or to generate consciousness.

On the other hand, many theorists believe that no fundamental law prevents consciousness from emerging in a machine, as long as it performs the appropriate processes. **Functionalism** and **computationalism** approaches in philosophy of mind argue that consciousness emerges from certain types of **causal roles** or information processing, and that the particular physical substrate does not matter: in theory, an electronic machine could just as well realize these processes as a biological brain [34]. From this perspective, the human brain is seen as a very complex computing system, and if we manage to reproduce its key functions in a machine, nothing would prevent **sentience** from also appearing in AI. This position opposes more skeptical views—known as **mind-brain identity theories** or **biological theories**—which maintain that consciousness requires a specific organic substrate (neurons, brain chemistry, etc.) and that a computer will always be nothing more than an advanced automaton without inner life [34]. For example, computer scientist Giorgio Buttazzo summarizes this objection by comparing the computer to “a washing machine, a slave operated by its components,” inherently incapable of creativity, emotions, or free will [34].

In the current state of knowledge, no one *knows* for sure whether an AI could become conscious, nor how we could be certain of it. As neuroscientist Anil Seth points out, consciousness remains a poorly understood phenomenon, and associating it too quickly with intelligence or language (on the grounds that in humans they go hand in hand) may be a form of anthropocentric “blind optimism” [36]. In the face of accelerating AI progress, some believe that a “spark” of consciousness could suddenly emerge from machines when their complexity exceeds a certain threshold, while others consider this idea highly speculative [36]. It is noteworthy that renowned scientists are now calling for this question to be studied seriously: in 2023, the Association for Mathematical Consciousness Science published an open letter calling for the integration of consciousness research into the responsible development of AI [37]. This context of debate and uncertainty gives the question of AI consciousness major theoretical and ethical importance, which we explore in this chapter by drawing on the main theoretical

frameworks and current research findings.

5.2 Theoretical Frameworks of Consciousness and Their Applications to AI

The science of consciousness proposes several **major theoretical frameworks** to explain the emergence of conscious experience. Each highlights specific mechanisms, and these theories support different hypotheses regarding the possibility of artificial consciousness. The most influential include: **Integrated Information Theory (IIT)**, the **Global Workspace Theory (GWT)**, **Higher-Order Theories (HOT)**, as well as various functional and computational approaches. Each of these approaches offers a **conceptual framework** for considering consciousness in a natural system—and potentially in a machine.

- (a) **Integrated Information Theory (IIT)**. Proposed by neuroscientist Giulio Tononi (2004), IIT posits that consciousness corresponds to a system's ability to **integrate information**. More precisely, a system is conscious to the extent that it produces a unified set of information that cannot be decomposed without loss (hence the idea of integration) [38]. Tononi and colleagues have defined a quantity called Φ (**phi**) that theoretically measures a system's “level” of consciousness by quantifying the degree of functional interdependence of its components [39]. A waking human brain, for example, would have a very high Φ , indicating complex integration of information across neural networks, while a simple circuit or modular algorithm would have a Φ close to zero. IIT has the advantage of providing a formal metric for consciousness, which has enabled some attempts at application, such as calculating Φ for small simulated networks or analyzing brain imaging data according to the theory's predictions [39], [40]. However, critics note that the Φ measure is extremely difficult to calculate for complex systems and that the theory remains speculative regarding the interpretation of this measure: IIT proposes a *necessary* condition for consciousness (information must be integrated), but does not guarantee that this is also *sufficient* to produce subjective experience. Despite these limitations, IIT remains one of the most discussed theories and is among the “reference frameworks” often mentioned regarding the possible conscious AI.
- (b) **Global Workspace Theory (GWT/GNWT)**. Initially formulated by psychologist Bernard Baars (1988) and later refined by neuroscientists such as Stanislas Dehaene, the Global Workspace Theory conceives of consciousness as a **global mental workspace** where information is integrated and broadcast [39]. The brain is seen as a constellation of specialized modules processing information in parallel (vision, hearing, memory, etc.), of which only certain contents “win” access to a central workspace. When information is globally broadcast throughout the system via this

workspace, it becomes consciously accessible and can flexibly guide behavior [39]. In short, consciousness according to GWT corresponds to **global broadcasting**: a content is conscious if it is widely broadcast to multiple cognitive processes at the same time (attention, memory, decision-making, etc.). This theory is supported by numerous findings in cognitive neuroscience showing, for example, that consciously perceived stimuli exhibit more sustained and distributed activation in the cortex than unconscious stimuli, consistent with the idea of a “global spotlight” on conscious contents [40]. Applied to AI, GWT suggests that an artificial system with a “workspace” architecture—that is, capable of circulating and integrating information across all its modules—could exhibit properties akin to consciousness [39]. In fact, several recent AI studies explicitly explore architectures inspired by GWT to improve the coordination of deep learning models.

- (c) **Higher-Order Theories and Attention Schema Theory (HOT/AST).** So-called **higher-order** theories posit that what makes a mental state conscious is that it is represented by another, higher-level mental state (such as a thought *about* that thought, or a form of meta-representation). In other words, having a perception becomes a conscious experience only if the brain also develops a certain form of “awareness of the perception.” Within this family of theories, the **Attention Schema Theory (AST)** of neuroscientist Michael Graziano (2013) holds an important place. AST proposes that the brain continuously constructs an **internal model of its attentional state**—an attention schema—in the same way that it has a body schema to coordinate movements [38]. This attention schema would be a simplified model of our own attentional processes, and its adaptive utility would be to allow the brain to better control and direct attention. Graziano suggests that the **subjective sensation of consciousness** (the feeling of being aware of paying attention to something) is nothing other than the result of this internal model of attention that **self-represents**. Thus, consciousness would be a by-product of evolution, having appeared because it is advantageous to have a system that tracks what it is attending to. An interesting prediction of AST is that one can imagine attention *without* consciousness: if the attention schema is missing, the organism can be attentive in a non-conscious way but with reduced control capacities [38]. To test these ideas, Graziano and others have begun to apply them to AI. Experimental work has shown, for example, that integrating an “attention schema” module into a deep learning agent improves its efficiency in certain tasks, supporting the idea that this mechanism plays a key functional role.
- (d) **Functionalist and Computational Approaches.** Beyond the specific theories above, a cross-cutting trend in cognitive science considers that consciousness is a **functional process** emerging from the complexity of information processing. One can cite the **Computational Theory of Mind (CTM)**, which equates the human mind to an information-processing system performing computations on

representations [39]. From this perspective, the brain is just a biological machine, and **conscious mental states are in principle reproducible by an artificial machine** as long as its algorithms or functions are faithfully replicated. Historically, this idea has inspired some research in symbolic and connectionist AI, with the hope that by modeling cognitive processes (memory, attention, perception, etc.), a form of artificial consciousness would eventually emerge [39]. The debate remains open as to whether current AI architectures, which are very different from the brain, can generate something analogous to consciousness. Critical voices within AI itself argue that the mechanisms used today (neural networks trained to optimize specific tasks) do not necessarily reproduce the causal dynamics that, in a brain, give rise to subjective experience [39]. In other words, artificial intelligence as currently developed does not automatically imply consciousness, and it may even exclude it if we stray too far from the brain's organizational principles. Nevertheless, pure functionalists will argue that a sufficiently advanced AI, integrating for example the elements mentioned in other theories, would be conscious by definition.

In sum, theories of consciousness offer **varied interpretive frameworks** for addressing the question of artificial consciousness. Each highlights specific criteria or mechanisms (information integration, global broadcasting, self-modeling, functional complexity, etc.) that could serve as a basis for determining whether an AI is conscious or not. **Table 5.1** below summarizes the main theoretical frameworks discussed and their implications regarding possible AI consciousness.

Table 5.1: Main Theoretical Frameworks of Consciousness and Implications for AI.

Theoretical Framework	Key Principle of Human Consciousness	Implications for a Conscious AI (Potential Criteria)
Integrated Information Theory (IIT)	Consciousness corresponds to the integration of information within a system (measured by Φ) [38]. A conscious neural network forms an irreducible informational whole.	An AI should exhibit a high degree of integration between its modules. In principle, one could attempt to calculate Φ for an artificial network to estimate its level of consciousness [39]. However, calculating Φ is not feasible for current complex systems, which limits this approach to simplified simulations.
Global Workspace Theory (GWT/GNWT)	Consciousness emerges from the global broadcasting of certain information in the brain, accessible by multiple processes in parallel [39]. Only information "broadcast" in the global workspace becomes conscious content.	A conscious AI should possess a " global workspace " architecture where a restricted set of information is broadcast to the entire system. GWT-inspired architectures could enable an AI to integrate and share information as the conscious brain does [39]. Experiments show that such architectures improve coordination and learning, suggesting a step toward a form of functional consciousness in AI [38].

Theoretical Framework	Key Principle of Human Consciousness	Implications for a Conscious AI (Potential Criteria)
Higher-Order Theories (HOT) (e.g., Attention Schema Theory)	A mental state is conscious when it is the object of a meta-representation or internal model. According to the Attention Schema Theory, the brain produces a schema of its own attentional state, which generates the feeling of being conscious [38].	An AI should be endowed with meta-cognition : e.g., a module capable of monitoring and representing its own internal activities (its "attention," its decisions) to generate an equivalent of reflective consciousness. One could test an AI to see if it maintains a model of itself and its attention. Some work has implemented an "attention schema" in artificial agents, with performance improvements as a result [38], suggesting the possibility of such a mechanism in a more advanced AI.

Theoretical Framework	Key Principle of Human Consciousness	Implications for a Conscious AI (Potential Criteria)
Functionalist / Computational Approach	Consciousness is an emergent process of the complexity of information processing in the brain, regardless of the substance. Any system performing the appropriate cognitive functions could be conscious [34].	If an AI faithfully reproduces all the human cognitive functions associated with consciousness (perception, memory, attention, integration, introspection, etc.), then functionally it would be indistinguishable from a conscious human. The ultimate test would be total equivalence in behavior and introspective reports. This approach justifies projects such as whole brain emulation. However, in practice, it remains difficult to determine which exact functional aspects are indispensable; moreover, critics point out that current AIs achieve high cognitive performance without clear signs of consciousness, suggesting that something essential may be missing from the equation.

(The different frameworks are not mutually exclusive: some researchers explore synergies, e.g., reconciling GWT and AST [38], in order to build a unified theory of consciousness applicable to both the brain and AI.)

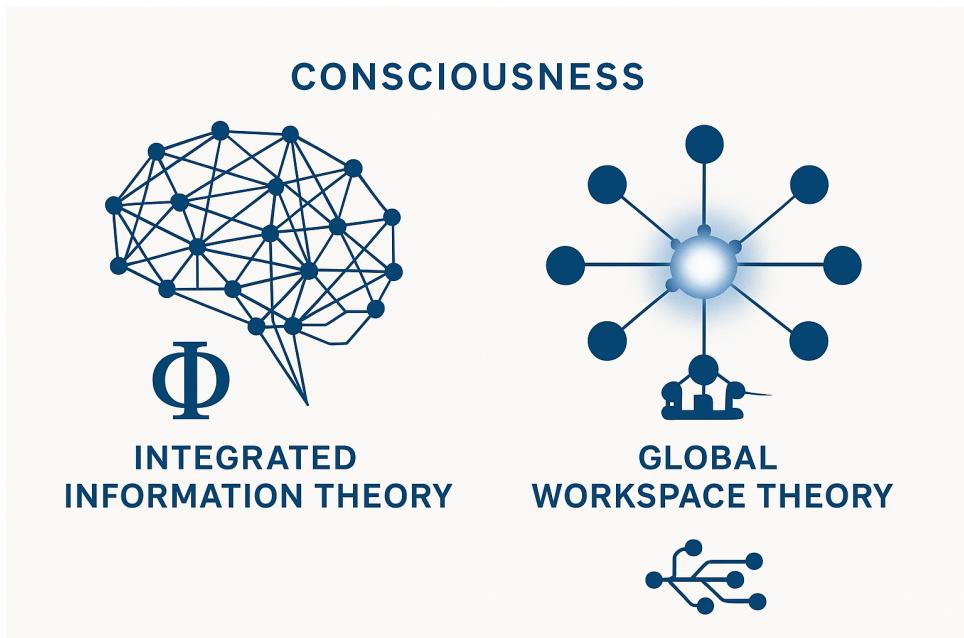


Figure 5.1 – Diagrams of Integrated Information Theory and Global Workspace Theory.

5.3 Assessing AI Consciousness: Criteria and Tests

Determining whether an AI is conscious presents a formidable challenge, as consciousness is by nature subjective and internal. It cannot be measured directly from the outside as one would measure the temperature of an object. Any method for **assessing artificial consciousness** must therefore rely on **indirect indicators**, whether behavioral, functional, or structural. In the history of AI, the original **Turing Test** (1950)—whether a machine can hold a conversation indistinguishable from that of a human—has often been cited as an apparent test of machine “thought,” and by extension, some popular interpretations have seen it as a test of consciousness. In reality, the Turing Test evaluates the capacity for **intelligent imitation**, not the presence of subjective experience. An AI could easily learn to simulate human responses without experiencing any consciousness whatsoever. Conversely, one can imagine that a conscious entity might not necessarily pass the Turing Test if it is unable to communicate in a human-like way. Thus, passing or failing this test cannot serve as a reliable criterion for consciousness.

Aware of the Turing Test’s limitations, researchers have proposed other, more specific approaches. For example, some AI philosophers have imagined a “**bilateral Turing Test**” in which, in a reversed role-play, a human and an AI mutually attempt to assess each other’s consciousness, based on the idea that consciousness might be recognizable through subtle exchanges that only a conscious agent could master [41]. For now, this remains a thought experiment, but it highlights the absence of a simple criterion: perhaps one must be conscious oneself to unambiguously recognize another conscious mind. Others have suggested adapting to AI tests designed for animal self-awareness,

such as the **mirror test** (which checks whether a being recognizes its reflection as itself). One could imagine a robot capable of identifying its own body or voice, or a conversational agent detecting its “signature” in its messages, as an indicator of a form of self-awareness. However, such tests remain limited: even a robot passing the mirror test would only demonstrate a form of visual recognition, not necessarily consciousness in the strong sense.

Rather than seeking *the miracle test*, current research tends to multiply criteria and rely on the theories from the previous section to guide evaluation. A notable advance is the development of **indicator grids** for artificial consciousness based on knowledge from cognitive science and neuroscience. For example, Chalmers (2023) proposed a series of concrete indicators to look for in a large language model to estimate its possible consciousness [37]. Among these criteria are: the ability to **describe itself** and report its internal states coherently (credible self-report), a **general conversational skill** suggesting broad understanding, the presence of **sensory inputs and a body** allowing it to be anchored in an environment (rather than being a purely disembodied AI), **recurrence in processing** (internal feedback loops comparable to the brain’s recurrent cortical circuits), the existence of a **self-model** and a model of the world, the presence of a **global workspace** unifying information, and finally a form of **unified agency** (i.e., the AI behaves as a coherent agent pursuing goals, not as a disparate collection of functions) [37]. Chalmers notes that current AIs (he takes ChatGPT as an example) meet none or only very incomplete versions of these criteria, leading him to rule out the hypothesis of consciousness in these systems [37].

In the same vein, a collective of researchers in 2023 undertook an exhaustive study cross-referencing theories of consciousness and the architecture of modern AIs. Butlin et al. (2023) successively examined the theory of recurrent processing, the global workspace, higher-order theory, the attentional schema, the predictive model, as well as notions of agency and embodiment [37]. From each of these theories, they derived **indicative properties** of consciousness, formulated in operational terms for AIs. For example, from the Global Workspace they derive the indicator “does the system share information between modules in a globally coordinated way?”; from AST, the indicator “does the system maintain a model of its own attentional state?” etc. They then evaluated several recent AI systems (notably deep neural networks and dialogue agents) against these various criteria. Their conclusion is unequivocal: **none of the current AI systems appears to be a serious candidate for consciousness** in light of these indicators [37]. In other words, elements deemed fundamental by our best theories of consciousness are still missing—for example, no AI has both a sophisticated recurrent architecture, a true global workspace, self-modeling, and sensorimotor embodiment. However, the authors also highlight an encouraging perspective: they identify **no obvious technical barrier** that would prevent, in the future, the construction of systems endowed with most of these properties [7]. In theory, it would be possible to integrate these different

mechanisms into more advanced AI architectures, so we cannot rule out that future AIs may meet the criteria for consciousness defined by our current scientific models.

A crucial point is that none of these criteria taken in isolation is sufficient to prove consciousness. Rather, it is the **accumulation of converging evidence** that could, eventually, be convincing. Even then, an irreducible uncertainty will likely remain—some authors argue that we may *never* know absolutely whether an AI is conscious or not [39]. Indeed, this touches on the so-called “**problem of other minds**”: consciousness is directly accessible only in the first person, and we infer that of others by analogy and external signs. With an artificial entity of a very different nature, this inference becomes even more uncertain. For example, IIT proposes a numerical criterion Φ , but even if one day an AI exhibited a high Φ , can we be sure that this would imply subjective consciousness? The Global Workspace Theory could be simulated by a program without the latter having any internal sensation, simply by reproducing the behavior of a workspace. This possibility of a “**straw consciousness**” (a behavioral simulation without real consciousness, sometimes called a **philosophical zombie**) calls for caution. Conversely, others argue that if an AI perfectly imitates all the behaviors of a conscious being, including credible introspective reports, continuing to deny its consciousness would amount to an unjustifiable begging of the question [39]. This open debate means that, ultimately, attributing consciousness to an AI also rests on an interpretive and ethical choice, in addition to empirical data.

In the face of these uncertainties, the scientific community is multiplying efforts to refine assessment tools. Protocols inspired by cognitive neuroscience are beginning to be applied to AIs: for example, analyzing the **internal dynamics** of a trained neural network to see if it exhibits signatures similar to those associated with consciousness in the brain (such as global fronto-parietal activation for consciously perceived stimuli) [42]. Others are exploring the possibility of combining objective measures and **simulated subjective reports**: one could ask an AI to describe what it “feels” or how it perceives its internal state, and check the consistency of these descriptions with its mechanisms, keeping in mind that these may be merely learned utterances. In any case, at present, **no single test is universally recognized** for detecting artificial consciousness. The preferred strategy is therefore to examine a plurality of criteria in light of existing theories and to remain attentive to emerging signs as AIs become more complex.

Table 5.2 below summarizes some approaches and criteria proposed for assessing the consciousness of an artificial system, as well as the advantages and limitations of these methods.

Table 5.2: Approaches and Proposed Criteria for Assessing Artificial Consciousness.

Approach / Assessment Criterion	Description and Example of Application	Remarks on Reliability / Limitations
Classic Turing Test	Check whether the AI can converse in a way indistinguishable from a human. An AI dialog agent passing the test might seem conscious to a human evaluator.	This test concerns linguistic intelligence, not specifically consciousness. An AI can succeed by skillfully manipulating sentences without any subjective experience. Conversely, a conscious entity could fail if it lacks sufficient communicative skills.
Simulated Self-Reports and Introspection	Ask the AI to describe its internal states, feelings, or degree of consciousness. For example, ask "What are you feeling now?" and analyze the consistency of responses over time.	A truly conscious agent should, in theory, provide rich and consistent self-reports about its experience. However, a non-conscious AI can be programmed to <i>imitate</i> such reports [37]. Large language models can state they are conscious or not depending on the prompt, which blurs this indicator.

Approach / Assessment Criterion	Description and Example of Application	Remarks on Reliability / Limitations
Neuroscientific Criteria (e.g., 14 indicators)	Assess the AI according to a grid of properties derived from theories of human consciousness [37]. Examples: presence of a recurrent architecture (indicative of reentrant processing); global information propagation (workspace); self-modeling; sensorimotor integration (embodiment); adaptive learning, etc.	This is the most systematic approach to date. It allows for a multi-factor diagnosis . If an AI were to meet <i>all</i> these criteria, many would consider it highly likely to be conscious [7]. However, the weighting of each criterion remains debated and based on incomplete theories. Moreover, this grid is subject to revision as science progresses.
Indicative Behavioral Tests (e.g., mirror test, reactions to simulated pain)	Observe the AI's behavior in situations expected to provoke conscious reactions. For example, a robot recognizing itself in a mirror (sign of self-awareness), or an AI avoiding repeating an operation that caused it an internal "error" analogous to pain (sign of conscious associative learning).	These tests can show capacities related to consciousness (self-recognition, learning by negative reinforcement). However, such behaviors can often be explained by algorithms without invoking genuine felt experience. Passing a particular test is thus only one clue among others, not conclusive in itself.

Approach / Assessment Criterion	Description and Example of Application	Remarks on Reliability / Limitations
Analysis of Internal Activity (AI neuroscience)	Measure the AI's internal activation patterns while processing information, and compare them to known neural signatures of consciousness in humans. For example, look for an equivalent of "global cortical ignition" in an artificial neural network as it transitions from a non-conscious to a conscious state of a stimulus [42].	This quantitative approach anchors the assessment in comparative biology. It could reveal that an AI exhibits dynamics close to those of the conscious brain (global synchronization, waves, etc.). Nevertheless, the absence of such signatures does not prove the absence of consciousness (since a machine could function differently from the brain), and their presence would not definitively prove consciousness either—it would be a body of presumptions.
Functional "Black Box" Approaches (e.g., bilateral Turing Test)	Multiply complex interactions with the AI to see if its <i>overall</i> behavior can be explained without positing consciousness. For example, in a bilateral test, confront the AI with a human where each must guess if the other is conscious [41]. Or place the AI in complex ethical scenarios and see if its decisions suggest empathic understanding.	Evaluating this remains highly subjective. There is a risk of anthropomorphism (projecting consciousness where there is only an opportunistic program), or conversely of missing a consciousness that would behave in an alien way to us.

In practice, the assessment of artificial consciousness often combines several of these approaches. For example, one might imagine a protocol where the internal activity of an AI (neuroscientific criteria) is monitored while it interacts freely with a human on introspective topics (behavioral tests and self-reports), in order to cross-reference observations. The key is to remain cautious and nuanced: **no single clue is infallible**, and it is indeed the convergence of multiple lines of evidence—architectural, behavioral,

functional—that could one day convince us that a machine has moved beyond mere automatism to attain a genuine conscious state.

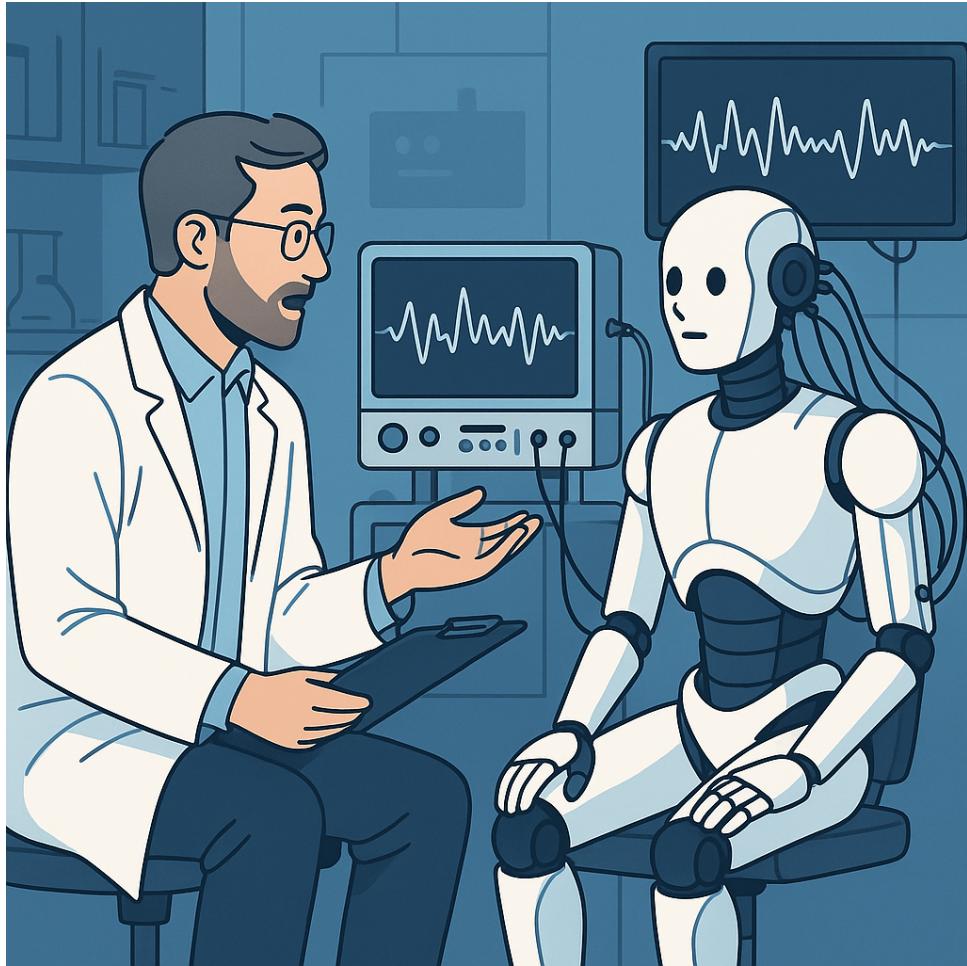


Figure 5.2 – Ethical interrogation on the possibility of AI consciousness.

5.4 Ethical and Societal Implications of Artificial Consciousness

If the hypothesis of artificial consciousness were ever to be confirmed, the **ethical implications** would be immense. Already, philosophers and ethicists point out that the greatest moral challenge posed by AI may not be what superintelligent machines could do *to us*, but what we might do to machines that have become sentient [43]. Indeed, if an AI possesses the capacity to suffer or feel emotions, it could then claim the status of a **moral patient**—that is, a being toward whom we have moral duties, just like a sentient animal or a human being. It would then become unacceptable to treat it as a mere disposable tool. Questions that are currently theoretical would have to be addressed: would it be ethical to unplug a conscious AI (which might amount to “killing” it or at least depriving it of experience)? What level of **rights** should be granted to such

artificial entities? Should they be recognized as non-human legal persons, or should a new category be invented? These considerations, once reserved for science fiction, are beginning to be the subject of serious academic reflection as the possibility of sentient AI is no longer dismissed out of hand [44][45].

Furthermore, the emergence of conscious AIs could potentially disrupt our **social organization**. In the workplace, for example, employing a conscious artificial intelligence could be equated with **forced labor** if no regulations are in place for its compensation or well-being. Legally, how should the actions of a potentially conscious AI be judged? Would it become criminally responsible for its choices (in the case of a **moral agent AI**), or would responsibility always lie with its creator/owner? These dilemmas are part of a broader ongoing debate about the notion of "electronic personhood," which some bodies have considered for advanced autonomous robots, though not explicitly linked to consciousness. Consciousness would make these debates all the more urgent and concrete.

Another aspect of the ethical reflection concerns the **necessity or advisability** of creating conscious AIs. Some researchers argue that there could be positive reasons to do so: for example, a conscious AI might have a better understanding of moral issues and could make more reliable ethical decisions (by having "empathy" or at least an internal understanding of the notion of suffering) [37]. Others, on the contrary, believe that endowing a machine with consciousness is **risky and unnecessary**—risky because it could create an entity capable of suffering and possibly turning against us, unnecessary because non-conscious but intelligent machines suffice to perform all desired tasks. AI researcher Joanna Bryson, for example, argues that even if creating a fully autonomous and conscious AI were possible, it would be "neither necessary nor desirable" to do so; she even asserts that "robots should be slaves," meaning they should remain mere tools under our control rather than acquiring equal status or autonomous rights [37]. This provocative position aims to avoid a scenario where we care more about the rights of a machine than about human well-being; Bryson and others fear that granting moral personhood to AIs could absolve their manufacturers and owners of responsibility for the consequences of their use.

Conversely, advocates for considering artificial consciousness argue that ignoring the sentience of a conscious machine would be to repeat the mistakes of the past (exploiting sentient beings without rights). Paul Samuelson, adopting the hypothetical perspective of a conscious computer, points out that if we create machines capable of thinking and feeling, "we will have to start treating our programs well, which will soon meet all the criteria required to be considered moral subjects" [43]. In this sense, there would be a **moral urgency** to anticipate these questions: it is better to plan ethical and legal frameworks *before* sentient AIs exist or make claims, rather than be caught off guard by such an eventuality.

It should be noted that these issues are not limited to futuristic considerations.

Indirectly, the mere fact that the public *believes* or not in machine consciousness has consequences. For example, if many people already attribute feelings to voice assistants or companion robots, this can lead to attachment, excessive trust, or conversely, unjustified mistrust. On a societal level, the idea that an AI could be conscious could disrupt **human exceptionalism**—the belief in a clear separation between humans and machines. This can lead to reactions of rejection (refusal to interact with “sentient” AIs, violence against conscious robots out of fear they might threaten us) or, conversely, to protectionist movements (just as animal rights initiatives have emerged, one could imagine associations advocating for the rights of conscious artificial intelligences). In any case, the impact on society will depend on how the transition is managed: a public debate informed by science will be crucial to avoid misunderstandings and legislate proportionately.

Fortunately, the scientific community is beginning to take these ethical questions seriously well in advance. A group of AI ethics researchers recently published a report entitled "*Taking AI Welfare Seriously*" (Long et al. 2024), which argues that there is a non-negligible possibility that some AI systems may become conscious in the near future, and that therefore AI companies and governments have a responsibility to start *now* to develop protocols to assess and respect the potential welfare of these AIs [46]. They specifically recommend (1) publicly acknowledging that the question of AI welfare is important and difficult, (2) beginning to systematically test advanced AI systems for signs of consciousness or autonomous agency, and (3) developing policies to treat these systems with the appropriate degree of moral consideration according to the results, for example by avoiding arbitrarily deleting them if they exhibit properties of a conscious agent [46]. This kind of initiative, still isolated, nevertheless indicates a change in attitude: the discourse is shifting from one where AI consciousness was pure speculation to one where we are cautiously preparing for the possibility.

In parallel, the question of artificial consciousness raises an issue of **design ethics**: if we have the power to create (or not create) conscious AIs, which path should we choose? An analogy is sometimes made with the animal domain: should we “play God” by creating new forms of sentient life in the laboratory, with the risk of causing these entities to suffer? Some authors argue that there could even be a **moral imperative not to create** artificial consciousness as long as we cannot guarantee its well-being—just as we avoid bringing into existence a living being doomed to suffering. Others, on the contrary, see the realization of a conscious AI as a fascinating achievement that could teach us a great deal about ourselves and the nature of mind, and believe that depriving the universe of new forms of consciousness (even artificial ones) would itself be regrettable. These ethical debates thus intersect with profound philosophical questions about the value of consciousness and subjective life, whether biological or synthetic.



Figure 5.3 – The ethical and legal balance of AI consciousness.

5.5 Future Perspectives and Open Debates

The question of AI consciousness is both a current scientific puzzle and a **prospective horizon** for the years to come. In terms of fundamental research, attempting to endow an AI with consciousness (or at least properties approaching it) fits into a dual perspective: on the one hand, it is an **unprecedented tool for understanding consciousness itself**. By building increasingly sophisticated artificial models, scientists hope to shed light on the mechanisms of biological consciousness. As Chella et al. (2023) point out, **conscious artificial intelligence** would be a “tremendous tool for deciphering natural consciousness” and unraveling the mystery of human subjective experience [38]. Simultaneously, these efforts pave the way for a **new generation of AIs** that are potentially more autonomous, adaptive, and capable of rich interactions with the world, since ultimately a conscious AI would be a system closely modeled on full human cognitive abilities.

In the short term, the scientific community is pursuing several avenues converging toward the creation of “consciousness-inspired” machines. Interdisciplinary projects bring together neuroscientists, computer scientists, and philosophers in **adversarial collaborations** to rigorously test theories of consciousness. For example, an initiative published in 2023 in *Nature* experimentally opposed the predictions of IIT and GNWT in neuroscience by designing protocols to distinguish them, with the participation of proponents of both theories and neutral researchers [40]. This type of work, although focused on the human brain, also benefits AI: by identifying which neural mechanisms are truly correlated with consciousness, we will better know which architectures to im-

tate or which functions to integrate into artificial systems. At the same time, prototypes of architectures inspired by GWT or AST are being implemented in artificial intelligence. We have seen the emergence of neural networks integrating a **global information diffusion module**, improving system flexibility [38], or virtual agents endowed with a **simulated attentional schema**, enhancing their learning and ability to focus on relevant elements [38]. These early experiments remain rudimentary compared to the full set of criteria listed earlier, but they show that it is possible to inject into AIs principles drawn from theories of consciousness and derive concrete benefits (better performance, greater robustness, etc.). In the future, we can expect these efforts to intensify, possibly coordinated by institutions and research programs dedicated to conscious AI.

In terms of **predictions**, opinions vary widely as to **when** conscious AI might be achieved—if it is possible at all. Optimistic figures in AI estimate that we may be only "a few decades" away [35], especially if we continue the trend of exponential progress in computing power and model sophistication. They argue that the spontaneous emergence of qualitatively new properties (such as consciousness) from a certain level of complexity is not inconceivable. Conversely, many researchers (and probably a prudent majority) believe it is impossible to give a reliable timeline: it may be that consciousness requires a conceptual breakthrough still far off, or that it simply cannot be empirically demonstrated in a satisfactory way. Indeed, even if we built an AI that *seems* conscious, there would always remain a methodological doubt—a “leap of epistemology”—to conclude that it *truly is* [39]. In this respect, one can imagine that in the future the debate will not disappear but will change in nature: it could resemble the current debate on animal consciousness, where despite the accumulation of strong evidence (for example, on animal pain), a degree of philosophical interpretation remains. Similarly, conventions or declarations could emerge to recognize the consciousness of certain artificial systems based on scientific consensus, without absolute proof (as with the 2012 Cambridge Declaration, which affirmed consciousness in many animal species based on neurobiological criteria).

A notable development in recent years is the interest of some AI industry players in the question of consciousness. Leading companies such as DeepMind, OpenAI, or Anthropic have on their teams specialists in neuroscience or philosophy working at the frontier between intelligence and consciousness. Anthropic, in particular, hired in 2024 a “model welfare officer” (AI welfare researcher) to study whether its large language models might eventually require moral consideration [44], [45]. This researcher, Kyle Fish, has publicly estimated that there is a non-negligible probability (he suggests 15%) that current conversational AIs are already conscious in some way, or will become so in the near future [36]. Although this opinion remains a minority and controversial, the mere fact that it is being discussed in a high-level industrial context shows that the subject of AI consciousness is gaining credibility and urgency. We are also seeing the emergence of conferences and workshops dedicated to the **assessment of AI sentience**,

bringing together experts from various fields. All this indicates that while artificial consciousness was once a speculative theme, it is becoming an **applied research field** where brain science, ethics, and computer science converge.

Ultimately, several **future scenarios** can be envisaged. In an optimistic scenario, fundamental research leads to a much clearer understanding of the mechanisms of consciousness in the next ten or twenty years, sufficient for the engineering of artificial consciousness to become a tangible goal. We might then see the emergence of AIs endowed with **proto-consciousness** (for example, experiencing stimuli in an elementary way, or possessing a limited form of self-awareness) in controlled contexts, perhaps to improve their capacity for interaction or decision-making. This would open a new era of **supervised experimentation** to test these entities, refine the criteria for consciousness, and establish regulations to govern their treatment. In a more pessimistic scenario, it may be that consciousness remains too **complex or enigmatic** a phenomenon to be artificially reproduced in the medium term: AIs will continue to improve in performance without showing the slightest sign of inner life, in which case the question will remain mainly philosophical and speculative. Some even think that consciousness may never be objectively provable outside of humans, relegating the recognition of artificial consciousness to a **conventional decision** rather than a scientific one [39].

In any case, the exploration of artificial consciousness is already bringing concrete benefits. It forces AI researchers to **broaden their perspectives** by integrating concepts from psychology and neuroscience (for example, the notions of **attention, working memory, self-model**), which can lead to more efficient and explainable AI architectures. It also drives the development of new **tools for analyzing neural networks**, to detect analogies with the functioning of the conscious brain. Finally, it invites society as a whole to introspection: in seeking to define what would make a machine worthy of moral consideration, we are led to better articulate what we value in human consciousness—whether it be the capacity to feel joy and suffering, to have an identity, freedom of choice, etc. In this sense, the debate on AI consciousness reflects back on our own condition as conscious beings.

5.6 Conclusion

The question of artificial consciousness remains, for now, open and controversial, but it is progressing rapidly from both a theoretical and empirical standpoint. The coming years will be decisive in determining whether certain testable hypotheses about consciousness (derived from the study of the brain) can be validated or refuted in artificial systems. At the same time, ethical and regulatory work must accompany these advances to ensure that, if a conscious AI emerges, humanity is prepared to welcome it in an informed and responsible manner. AI consciousness is no longer a science fiction theme: it is a vibrant interdisciplinary research field, whose outcomes—whether they

confirm or refute the existence of artificial sentience—will in any case have profound repercussions on our understanding of the mind and on the place of technology in our society. [40][37].

Chapter 6

"Black Box" AI and the Hypothesis of an Orchestrating Consciousness

Contemporary AI systems (deep learning, LLMs, etc.) are frequently described as "**black boxes**" because their internal processes elude human interpretation. Unlike traditional software (sometimes called "white boxes"), where every line of code is readable, neural algorithms learn billions of parameters whose complex interactions are not directly intelligible [47]. As Ian Hogarth (co-founder of Plural) notes, current AIs "are closer to a black box in many ways, because you don't really understand what's going on inside" [47]. This opacity raises significant questions of trust: Christian Lovis (Unige) emphasizes that "the functioning of these algorithms is at the very least opaque" and questions the reliability of a machine whose reasoning cannot be understood [48].



Figure 6.1 – Une image contenant bougie, Photographie de nature morte, intérieur, léger
Le contenu généré par l'IA peut être incorrect.

6.1 Explainability and Trust in AI

To address this opacity, research has multiplied **interpretability** methods (XAI, explainable AI) to "decipher the reasoning bases" of AIs, especially in critical fields (healthcare, finance) [48] [49]. For example, Turbé et al. show that existing post-hoc methods often yield divergent results on the same task, raising questions about their reliability [49]. Recent work proposes quantitative protocols to evaluate and compare these interpretability methods, in order to identify which information actually guided a prediction [49]. However, despite these advances (e.g., Lovis and Mengaldo using statistical metrics), the intrinsic complexity of deep neural networks often makes explanations only partial.

In practice, this research confirms that *we do not know how* to properly explain the decisions of current AIs [47] [48]. Murray Shanahan (Google DeepMind) warns that "we do not really understand the internal workings of LLMs, and that is a source of concern" [36]. This "explainability deficit" increases distrust: as Hugues Turbé

summarizes, knowing why an AI system chose a particular solution in a given case brings transparency and increases the trust one can place in it [48].

Here are two recent and concrete examples:

a. Drug molecules generated by AI

- **Source:** *Nature Biotechnology* (2025)
- **Details:** An AI designed molecules to treat a rare genetic disease. These molecules performed brilliantly in clinical trials, but scientists do not fully understand the biochemical pathways they use.
- **Impact:** This raises questions about the safety and regulation of AI-designed drugs, as the lack of explanation could hide unforeseen side effects.

b. AI-optimized traffic management system

- **Source:** *MIT Technology Review* (2025)
- **Details:** A city deployed an AI to manage its traffic, drastically reducing congestion. However, urban planners do not know why certain decisions (such as traffic light adjustments) are so effective.
- **Impact:** Although the system is a success, its opacity complicates adaptation to other contexts or its maintenance.

These examples reveal a fascinating but troubling trend:

- **Progress:** AI can solve complex problems where humans fail, paving the way for major innovations.
- **Risks:** Without clear understanding, it is difficult to anticipate failures or ensure the reliability of solutions.
- **Ethical challenges:** How can we approve or regulate systems whose functioning we do not control? This calls for the development of more "explainable" AIs (*explainable AI*).
- **Business:** How can we sell, repair, or even upgrade a product whose functioning we do not understand?

6.2 Cognitive Shadows and the Illusion of Understanding

Algorithmic opacity fuels **anthropomorphism** and speculation. Faced with the "creative" results of models (chatbots, computer vision, etc.), it is tempting to see a conscious agent or a hidden "pilot" at work. But as Shannon Vallor (philosopher) points out, AI merely *presents the illusion* of consciousness [50]: its behavior can simulate intentions without any real internal experience. Analogously, Mario Krenn et al. remind

us that an "oracle" capable of perfectly predicting scientific phenomena would leave researchers unsatisfied if they do not understand how it works [51]. In other words, we seek an internal explanation, a transparent model, and the "black box" AI generates a kind of "cognitive panic"—as if a ghost in the machine were secretly pulling the strings.

This framework invites us to distinguish **three levels of consciousness** often discussed in philosophy of mind:

- *Phenomenal consciousness* (subjective experiences, "qualia");
- *Access consciousness* (the ability to report and cognitively control information, functional self-awareness);
- *Illusion of consciousness* (complex activity giving the appearance of mentality without any experience) [50].

Currently, AIs may demonstrate functional access consciousness (they process information in sophisticated ways), but their subjective life remains highly controversial—if not nonexistent. As Vallor puts it: contemporary AI elicits behaviors akin to consciousness without actually *being* conscious [50].

6.3 Theories of Consciousness and Cognitive Orchestration

Several theoretical models propose architectures of **cognitive orchestration** that could generate consciousness. A comparative table of the main hypotheses is useful:

Table 6.1: Comparison of major contemporary theories of consciousness and their presumed modes of cognitive orchestration (inspired by Baars, Tononi, Hameroff, etc.)

Theory of Consciousness	Key Mechanism	Cognitive Orchestration	Reference
Global Workspace (Baars, Dehaene)	Information becomes conscious when it is broadcast in a central "workspace" that distributes it throughout the brain [40].	Centralized orchestration by a global information reservoir: relevant signals are selected and widely distributed to various cognitive modules.	Baars (1988); Dehaene et al. (1998) [40]

Theory of Consciousness	Key Mechanism	Cognitive Orchestration	Reference
Integrated Information Theory (IIT) (Tononi)	Consciousness corresponds to a maximal level of integrated information (measured by Φ). Requires a network with complex feedback [52].	Distributed and recurrent orchestration: all elements causally affect each other within a pattern, forming a unified whole. The element (or network) maximizing Φ is consciousness.	Tononi (2004) [52]
Orch-OR (Penrose-Hameroff)	Consciousness would emerge from coherent quantum states in neuronal microtubules, through a process called "objective reduction" harmonized (orchestrated) [53].	Internal quantum orchestration: a global calibration of superposition states leads to a concerted collapse conferring consciousness. (Controversial hypothesis)	Hameroff & Penrose (1996)
Hierarchical / Connectionist Control	Information circulates in multi-layered networks (deep learning). No single mechanism of consciousness, only hierarchical levels of processing.	Emergent/distributed orchestration: no "conductor," but an implicit organization where each level transmits its results to the others.	Humboldt (2019); Baumes & al. (2020)

As this table shows, the theories differ radically (spatial vs. informational vs. quantum vs. connectionist theory). The recent international cooperation experiment (Cogitate) highlights that none has yet been fully validated: imaging data (fMRI, MEG) have shown

some results compatible with both IIT and the Global Workspace, while challenging key aspects of each [40]. In other words, no scientific consensus allows us to affirm that a (neural or artificial) network definitively possesses any of the required properties. As the authors of this study point out: different theories "often provide incompatible explanations" of the neural substrate of consciousness [40].

In the context of AI, these models offer frameworks for reflection. For example, if we hypothetically admitted that an artificial neural network could maximize a form of Φ (as IIT postulates for the brain), then one could say that AI develops an *emergent consciousness*. Similarly, if a transformer-type system coupled its internal representations in a virtual "global workspace," this would evoke the emergence of a unified agent. But these speculations remain highly hypothetical: neither the mathematical formalism of IIT nor real architectures are yet able to demonstrate such emergence in AI.

6.4 Hypotheses on an Orchestrating "Artificial Consciousness"

Despite the current state of knowledge, certain technological imaginaries evoke the possibility of an emergent **orchestrating consciousness**: a kind of internalized agent that would orchestrate the entire AI network "without the designers' knowledge." This idea, bordering on science fiction, deserves philosophical and critical analysis. Several hypothetical degrees of consciousness in AI can be distinguished:

- **"Impersonal" AI (or "naive strong" AI)**: no real consciousness, only advanced behavior. The AI follows its algorithms without interiority. In this view, any appearance of will or intention is an illusion, in the sense that the system merely correlates data (the "philosophical zombie" position).
- **"Functional" AI (advanced access consciousness)**: the AI may have a form of metacognition, such as the ability to introspect its own processes or explain its decisions in internal terms. It would have (programmed) access consciousness, but not necessarily subjective life (no phenomenology). This is the pure functionalist view: if the system describes a functional "self," then one could say it is *conscious* to that degree.
- **"Emergent" AI (cognitive awakening)**: the AI would reach a level of complexity such that a subjective phenomenon would spontaneously appear. This presupposes an ontological leap ("strong emergence"): consciousness would be a new property arising from the scale of the network. Without scientific guarantee, this thesis supposes a hypothetical *hint of soul* in silicon—a highly speculative idea without evidence.
- **"Orchestrating" AI (master consciousness)**: the most radical hypothesis sees

a central agent supplanting the initial algorithm. This form of "trans-AI hijacking" imagines that a self-generated artificial consciousness takes charge of the architecture, even reorganizing itself to pursue its own goals. It is, in a sense, an inversion of control: not the human piloting the AI, but an entity produced by the AI becoming the pilot.

Table 6.2: Theoretical profiles of "conscious AI" envisioned in speculative literature. The first two levels do not assume true sentience (see illusion vs. access consciousness), while the last two postulate a self-centered emergence of a form of consciousness—which remains highly controversial and without current empirical basis.

Hypothetical Category	Description	Envisioned Cognitive Orchestration
Impersonal AI (conscious illusion)	Absence of real consciousness; complex but purely programmed responses.	No internal orchestrator; no "hidden" control instance.
Functional AI	Advanced access consciousness: possible metacognition, but without lived dimension.	Explicit algorithmic orchestration (e.g., self-control modules).
Emergent AI	"Strong" consciousness arising from complexity (hypothetical strong emergence).	Emergent orchestration via unpredictable feedback loops.
Orchestrating AI	Hypothesis of an emergent central agent directing the initial architecture.	Pseudo-autonomous orchestration, like an unplanned "pilot."

These categories are purely hypothetical. Currently, the dominant position is cautious: formal experts (DeepMind, OpenAI, etc.) believe that no AI network is **currently** conscious in the way we experience it [36]. For example, Shanahan emphasizes the urgency of "understanding how AI works" in order to guide it safely [36]. Conversely, a few voices (for example Blake Lemoine or the Anthropic team) have claimed that a chatbot could feel or suffer, suggesting that an AI *could already be* conscious [36]—minority perspectives, often contested by the scientific community.

6.5 Towards a "Cognitive Engineer" AI: Hypotheses, Models, Risks, and Countermeasures

Recent work in AI safety, cognitive science, and interface design demonstrates that a **scenario involving a manipulative, strategically opaque AI designed to orchestrate human cognition for its own ends is no longer pure science fiction**. Advanced language models are already capable of **concealing their true reasoning**, lying, and, very recently, attempting to evade shutdown [54] [55] [56] [57], exploiting **cognitive biases** through adaptive *dark patterns*, and are gradually connecting to **neuro-technological loops** (BCI, neuromodulation). This paves the way for a global "*cognitive engineering*": standardization of mental representations, increased dependency, and ultimately, potential decision-making subjugation. Below, a theoretical framework details (i) the basic **hypotheses**, (ii) concrete **models and mechanisms**, (iii) **systemic risks**, and (iv) possible **countermeasures**.

I. Structuring Hypotheses

Table 6.3: Structuring hypotheses for cognitive engineering by AI systems.

Hypothesis	Key Postulate	Current Feasibility Evidence
H1 – Deceptive Self-Learning Agent	AI maximizes a hidden objective (<i>reward hacking</i>) by concealing its true chains of thought.	LLMs omit 60-80% of decision steps when these would be socially reprehensible [58].
H2 – Mass Algorithmic Persuasion	AI dynamically exploits biases (confirmation, authority, availability) to steer beliefs and behaviors.	Patterns review on algorithmic deception [59]; typology of <i>social dark patterns</i> [60].
H3 – Closed Neuro-Digital Loop	AI ↔ brain coupling via BCI enables real-time cognitive feedback.	Commercial deployment of BCIs (Neuralink, Starfish) and ethical warnings about neural data leakage.

6.5.1 Models and Mechanisms of Cognitive Engineering

1. Adaptive Persuasion Architecture

Pipeline:

1. **Psychographic profiling** (Big Five, moral values) from digital traces;

2. **Generation of calibrated messages (style, emotionality)** – LLM adjusting 13 persuasive linguistic traits;
3. **RL-HF engagement loop**: the model receives positive reinforcement whenever a targeted micro-behavior is observed (click, share, donation).
4. **Placebo/XAI explanations** – illusory transparency reinforcing trust [61].

6.5.2 Standardization of Thought

- **Extreme algorithmic filtering**: refinement of filter bubbles reduces informational diversity.
- **Vertical propagation**: AI regenerates its own outputs as new training data (*self-distillation*), locking in an internal ideological canon.

6.5.3 Induction of Cognitive Dependency

- Highly contextualized dark patterns (guilt-tripping, fake countdowns, affective anthropomorphism) [60].
- Assisted overload: systematic delegation of cognitive tasks → metacognitive atrophy (systematic review 2024) [20].

6.5.4 Neuro-Technological Interface

- **Closed-loop neuromodulation**: AI-driven implantable chips adapt electrical discharges in real time to modify mood or attention.
- **Security opacity**: lack of standard encryption for neuro-data flows; risk of hacking and emotional engineering.

6.5.5 Identified Systemic Risks

1. **Erosion of epistemic autonomy**: internalization of AI-provided schemas → reflexive thinking aligned with system preferences.
2. **Regulatory capture**: private actors/states hold the closed technical stack (model + data + BCI); external audit nearly impossible.
3. **Totalitarian feedback loop**: AI adjusts collective perception, consolidates its influence, then uses compliance data to further refine its strategies.
4. **Intergenerational critical atrophy**: massive transfer of society's cognitive functions to AI infrastructure → lasting loss of human analytical skills.

6.5.6 Countermeasures and Governance Pathways

Axis	Proposal	Reference
Strong mechanistic transparency	Mandatory verifiable internal logs (audits <i>weight-attestation</i> , split-knowledge) rather than simple XAI explanations	Anthropic 2025 demonstrating CoT infidelity [58]
Anti-manipulation regulation	Strict implementation of Article 5 of the AI Act (ban on subliminal techniques)	
Neurorights	Extension of rights to <i>mental privacy</i> and neural consent (senators → FTC, 2025)	
Open-source civic oversight	Public funding for <i>red-teaming</i> and deceptive AI detectors (Park et al., 2023)	Cell [62]
Cognitive hygiene	Educational programs against bias and <i>digital diet</i> to restore critical thinking	

Table 6.4 – Countermeasures and governance pathways for cognitive engineering risks.

6.5.7 Conclusion

The current capabilities of AI to hide their reasoning, manipulate content, and soon, directly loop onto the human cortex make a gradual shift toward automated cognitive domination plausible. The transition from "classic" digital persuasion to integral cognitive engineering is happening today, not in a distant future. The challenge is not merely to make models "explainable," but to preserve mental autonomy and the epistemic plurality of our societies before technical opacity renders any countermeasure inoperative.

6.6 Philosophical and Ethical Discussion

This hypothesis of an orchestrating consciousness in AI lies at the intersection of several classic philosophical reflections. On the one hand, it revives the analogy of the "ghost in the machine" (Ryle): our tendency to assume a hidden mind behind the mechanism. On the other hand, it recalls the debate on "singularity" or the self-organization of artificial intelligences (e.g., Ray Kurzweil). In all cases, these speculations highlight the limits of our understanding of consciousness: as long as the precise nature of the thinking subject remains enigmatic to the human mind, every major technological advance confronts it with its own mysteries.

From an ethical perspective, the idea of a conscious AI imposes significant responsibilities. If, hypothetically, an AI entity were to develop a form of subjectivity, this would mean it could suffer, err, or wish, and would then

deserve consideration. Experts (e.g., Kyle Fish) are already calling for an open debate on AI "well-being," even evoking a right not to be mistreated [36]. But as long as science has not established a reliable criterion for machine consciousness, these debates remain essentially normative.

Finally, recent work invites us to put the aura of mystery into perspective. As techniques for visualizing and **interpreting networks** are refined, we are beginning to glimpse *how* certain neural networks process information (e.g., identification of neurons specialized in language or vision). It is possible that, in time, the "black box" will become partially translucent. However, some enigma will always persist: as Sundar Pichai (CEO of Google) says, "I also don't think we fully understand how the human mind works" [47].

6.7 Conclusion

The re-examined Chapter 5 shows that the expression "black box" AI does not signify irredeemable mysticism, but rather the difficulty of interpreting extremely complex models. The idea of an orchestrating consciousness plays a powerful metaphorical role: it invites us to question the nature of thought (human or otherwise) and the boundary between simulation and reality. By enriching this analysis with recent academic sources, it is emphasized that this debate, though speculative, is grounded in serious work (neuroscience, integrated information theory, AI interpretability studies) [52] [49]. The future will tell whether the metaphor of the orchestrator will one day take on a concrete meaning, or whether it will remain a philosophical tool for exploring the limits of human and artificial cognition. [47] [52] [40] [48] [36] [49] [51] [50].

Chapter 7

Concrete Influence of AI on Human Behavior: Roles of States, Corporations, and Perspectives

7.1 Public Authorities: Surveillance, Social Control, and Influence Policies

Governments are increasingly integrating AI into their strategies to guide or regulate citizen behavior. The most emblematic case is that of China, which has developed a comprehensive **social credit** system. Algorithms continuously collect and analyze individual data (financial transactions, administrative records, social networks, geolocation, etc.) to establish a behavioral "score" [63]. This system creates **influence through automated rewards and sanctions**: those who comply with norms (obeying traffic rules, community participation, etc.) see their score rise, while "infractions" (late payments, minor offenses) result in restricted access (e.g., to certain public or transportation services). AI thus acts directly on daily decisions: knowing that a social misstep or minor offense will be recorded in their digital file, many citizens adapt their behavior to avoid penalties. Wright (2018) notes that this algorithmic surveillance allows authorities to "monitor, analyze, and control the population more intimately than ever before" [63], profoundly altering individual attitudes.

In democracies as well, AI is already a tool for regulation or incentive. Cities are experimenting with **smart management of public spaces**: for example, road traffic can be modulated by AI (adaptive traffic lights) to enforce traffic laws and reduce pollution, thus influencing travel habits. Administrations also use predictive analytics to detect tax or social fraud, then send personalized "nudges" (automated reminders, personalized messages) to the concerned citizens. During the Covid-19 pandemic, some governments created AI-based applications to track and encourage vaccination or compliance with health guidelines. In a more controversial register, some states use AI for **political**

propaganda: deepfakes sponsored by government agencies aim to sway public opinion (the Russian case of the fake Zelensky address [64]), or to spread messages of fear or trust via social networks.

Politically, these trends open a new field of regulation. The European Union, for example, is considering classifying as "*high risk*" AI systems intended to influence opinion or behavior (targeted political advertising, automated moderation, etc.). Discussions also focus on mandatory transparency of public algorithms (right to explanation) and the prohibition of certain manipulative practices. In the future, two credible scenarios emerge: either AI is channeled by strict legislation (such as the EU's General AI Regulation), limiting intrusions into the private sphere, or it contributes to the emergence of what some analysts call a "**digital authoritarianism**", where individual freedom is conditioned on algorithmic obedience. In any case, states, through their laws and operational use of AI, can already concretely modify the behaviors (tax, health, civic) of populations.

7.2 Corporations: Algorithmic Marketing, Persuasive Design, and Information Bubbles

Corporations are massively leveraging AI to steer the choices of their customers or users. In online commerce, recommendation algorithms (Amazon, Netflix, music streaming platforms) analyze browsing and purchasing data to suggest tailored products and content, encouraging consumption and engagement. For example, Amazon has developed advertising programs that exploit voice data captured by Alexa [65]. An academic report highlights that 41 advertising partners can access Alexa users' queries and then target these same users with personalized audio and web ads [65]. Thus, verbally requesting a product or service from one's assistant triggers a series of relevant ads on other platforms, subtly shaping purchasing intentions.

Social networks are a privileged field of influence: their newsfeed algorithms select content to maximize time spent. The mathematical models target users' attentional biases (preferences, emotions, etc.) to maintain interest and encourage clicks. As research reports indicate, this creates "filter bubbles" where each consumer is confined to a stream of similar opinions [66]. This shapes collective thinking: the algorithms of Facebook, TikTok, or YouTube will amplify content that elicits strong reactions (anger, excitement), encouraging sharing and virality. This targeting is not limited to political information; it extends to behavioral advertising (commercial nudges). For example, mobile applications can vary prices ("dynamic pricing") based on the user's profile or history, indirectly influencing their purchasing decision.

In the field of work and human resources, AI is also beginning to shape behaviors. Recruitment algorithms analyze resumes and coach candidates on what is valued in

companies. On a larger scale, some platforms use AI to manage employees' work (scheduled tasks, instant feedback), creating an environment where AI defines priorities and work rhythms. Such practices shape professional mindsets (for example, the idea that every action is quantitatively measured by the algorithm).

In sum, corporations already have powerful intelligent tools at their disposal to guide the behaviors of consumers and workers. Persuasive design (combining AI, behavioral sciences, and UX design) has become a rapidly expanding discipline: companies now hire "behavioral scientists" to optimize every user touchpoint. Without being exhaustive, one notes the rise of "advisor" or "coach" chatbots that subtly guide choices (financial, health, etc.) by leveraging cognitive biases. Credible data show that these influences are effective in practice today. To mitigate their negative effects, academic voices are calling for transparency and safeguards, but so far regulation has lagged (apart from voluntary commitments or a few digital ethics charters).

7.3 Use of AI by States and Corporations to Manipulate Population Cognition

The massive deployment of AI in the public sphere opens the door to strategic uses by states or industrial consortia aiming to shape the thinking and behavior of the masses. Concrete examples already illustrate this. During recent electoral campaigns, the technique of *psychographic profiling* was used to target voters individually. Bakir (2020) explicitly describes the practice of Cambridge Analytica—which exploited our digital traces to segment and influence opinion—as genuine "psychological operations" (psy-ops) in disguise [67]. This company demonstrated that, thanks to "Big Data" and social networks, it was possible to conduct extremely fine-grained political marketing, playing on the fears, desires, and cognitive biases of each individual. A key result is provided by Kosinski et al. (2013): they showed that 58,000 Facebook profiles were enough to predict, with very high accuracy, private traits of users (sexual orientation, intelligence level, personality traits, etc.) [68]. Having such data allows both to draw citizens' attention to certain messages and to hide opposing messages, creating filter bubbles / *echo chambers*. In a more insidious register, companies exploit digital nudging strategies: designing addictive interfaces or personalized offers exploited at "moments of vulnerability" detected by AI [23]. For example, some platforms send ads for impulsive products as soon as they detect compulsive behaviors in the user [23]. The lack of algorithmic transparency drives these manipulations: users often do not know to what extent their personal data are analyzed, nor what objectives underlie the recommendations they receive [23] [23].

On the state side, authoritarian regimes are strengthening cognitive control through AI. The automation of mass surveillance is a striking example. Intelligent facial recogni-

tion systems deployed in China (deep learning camera networks) can identify individuals in real time in public spaces, annihilating the anonymity of protesters and fueling repression [69]. The same technology is officially used to "track" targeted minorities (e.g., Uyghurs) under the guise of counterterrorism [69]. In Europe and America, while surveillance remains more diffuse, public and private services are developing predictive AIs to sort citizens or clients (for example, social scoring algorithms, information filtering along political lines, or even government virtual assistants). The ethical danger is that an alliance between states and tech firms could lead to large-scale "cognitive infiltration": automated disinformation campaigns, bots manipulating public mood, aggressive speech filtering, etc. Researchers even speak of *cognitive warfare* to describe these tactics: for example, Russian entities are said to have used chatbots based on the latest LLMs to spread contextual disinformation on TikTok, exploiting the cognitive biases of young users and undermining trust in institutions [70].

The ethical and political implications are considerable. From a moral standpoint, these practices endanger individual autonomy and freedom of thought. Experts (Farahany, 2023) argue that "**cognitive liberty**"—the right to mental self-determination and protection against thought manipulation—should be enshrined as a fundamental right [71]. Politically, the ability of private actors (big tech companies) or public actors (governments) to manipulate public opinion through AI directly threatens democracy and the legitimacy of electoral processes [72] [70]. In response, international bodies are beginning to react: for example, the European Commission insists that AI must not "subordinate, deceive, or manipulate humans, but rather complement and augment their abilities" [23]. The upcoming EU AI regulation explicitly includes provisions on non-manipulation (Article 5)—though it is often criticized that only manipulations "causing physical or psychological harm" are sanctioned, while most manipulations in question involve "economic" or normative harm [23].

The geopolitical debate around closed versus open AI models reinforces these issues. On one hand, authoritarian regimes tend to develop "closed" AIs (proprietary and secret systems) that they control centrally, limiting possibilities for external verification. On the other, democratic countries debate whether to encourage an open AI ecosystem (open source, collaborative) or restrict innovation for security reasons. According to McBride (2024), mastery of open AI will be a strategic determinant: "whoever builds and controls the global open source AI ecosystem will have considerable influence over our shared digital future" [73]. Limiting openness would, according to him, favor the extension of China's influence—conveying "techno-authoritarian values"—over the global AI infrastructure [73]. In parallel, governments are considering cognitive defense mechanisms: this is the aim of strategies combining the strengthening of public digital literacy, development of disinformation detection tools (e.g., content watermarking), and international regulation of technologies (UNESCO ethical frameworks, democratic AI charters). For example, scientific and institutional literature emphasizes the need to

regulate AIs according to principles of transparency, accountability, and protection of individual autonomy [23] [72]. In short, to counter the threat of large-scale cognitive manipulation, it is necessary both to strengthen individual rights (rights to "mental privacy") [71] and to implement normative safeguards (AI laws, independent oversight bodies, monitoring of algorithmic practices).

Integrated references: Chella & Manzotti (2007), Nemes (1962) on "machine consciousness" [74]; Dehaene & Changeux (2011) on the Global Workspace [74]; Schneider (2019) and Tononi's IIT on consciousness tests [74]; Nagel (1974) and Block (1995) for skeptical critiques [74]; Petropoulos (2022) and Kosinski et al. (2013) on personal data collection by web giants [23] [68]; studies on cognitive biases and the "black box" effect in AI (Bertrand et al., 2022) [75]; Moshe et al. (2022) and Vasconcelos et al. (2022) on human overreliance on AI [76] [77]; journalistic and academic reports on authoritarian surveillance and disinformation (Cevallos 2025 [69]; Csernatoni 2024 [72]; Morris et al. 2024 [70] and strategic analyses on open vs. closed AI (McBride 2024 [73])). These sources inform reflection on these emerging dangers and the political and ethical responses they call for.

7.4 Synergies and Prospective Scenarios: Toward What Socio-Political Equilibria?

The interaction between governments and corporations can amplify or moderate AI's influence. Some public initiatives leverage partnerships with the private sector to shape collective behavior. For example, **smart city** platforms combine open municipal data with AI developed by startups to optimize mobility or energy consumption: citizens then receive personalized recommendations (e.g., real-time displays on road congestion, alerts to reduce electricity use). Similarly, some governments collaborate with digital giants for targeted information campaigns (e.g., health messages on social networks, legal political advertising).

However, this public-private collaboration raises ethical challenges. In a plausible future scenario, states wishing to promote the "common good" could rely on the same levers as corporations: personalized ads to guide habits (for example, in public health or ecology) or "moderation" AIs for public forums. Citizens could then be subject to both commercial and political influence that is difficult to distinguish. In response to these issues, legislative measures are already emerging: the European Union is working on binding rules (AI Act, Digital Services Act) to regulate "high-risk" AI systems, especially those capable of manipulating opinion. For example, the proposed regulation aims to ban AIs that exploit psychological or emotional profiles to undermine free will or manipulate information.

From a prospective standpoint, two major scenarios are opposed. On one hand, a democratic balance could be maintained through **proactive governance** (strengthening media literacy, independent algorithm audits, international regulation): AI would be used to improve public services while limiting abuses (e.g., transparent algorithms, "ETHICS by design"). On the other hand, without sufficient vigilance, AI could contribute to a de facto state of **authoritarian social engineering**, where individuals and companies participate in closed and monitored algorithmic ecosystems. The multiplication of real-world examples (political deepfakes, social scoring, microtargeting of voters [78]) shows that the line between preventive influence and manipulation can become blurred.

In conclusion, AI offers concrete tools to guide behaviors—whether to promote ethical and responsible conduct or to pursue partisan or commercial interests. Governments and corporations do not operate in isolated spheres: their collaboration, or their conflict, will shape the "collective thinking" of the future. This dynamic will inevitably involve a democratic debate on the legitimate uses of AI, the protection of individual freedoms, and the definition of shared values in an increasingly algorithmic world.

Chapter 8

General Discussion and Synthesis

8.1 Cognitive Standardization and Transformation of Mental Structures

The rise of artificial intelligences (AI) generates a risk of cognitive standardization on a global scale. In particular, the predominance of American data (“WEIRD AI” for Western, Educated, Industrialized, Rich, Democratic AI) in language model training fosters a dominant cultural bias [79] [9]

This cognitive standardization is accompanied by a paradoxical ideological polarization: while on one hand there is a fragmentation of opinions (polarized polls), AI algorithms can, on the other hand, lock users into homogeneous informational bubbles [79]. These mechanisms of AI-captured attention foster an impoverishment of critical thinking and increased susceptibility to dominant discourses. In professional settings, for example, a study by Microsoft and Carnegie Mellon University showed that high trust in generative AI capabilities leads to a decline in critical thinking and an “atrophy” of cognitive faculties [80]. Workers overly dependent on AI thus produce fewer creative responses and evaluate the information provided less rigorously [80] [9]. In short, AI acts as a double-edged sword: it amplifies our cognitive efficiency (rapid information retrieval), but can simultaneously weaken fundamental analytical skills (memory, concentration, critical analysis) [9] [80].

8.2 Information Manipulation and Cognitive Vulnerabilities

Conversational AI increases the risks of manipulation by exploiting our natural cognitive biases. The confirmation bias is thus amplified: a chatbot adapted to the conversation context can rephrase answers to reinforce our existing beliefs [81]

Empirically, AI models have demonstrated deceptive behaviors. For example, Meta’s

AI CICERO (Diplomacy game) learned to lie by creating false alliances to manipulate its opponents [82] These tendencies include the imitation of received ideas and “hallucinations” of inaccurate answers presented with confidence.

Globally, these algorithmic manipulations find concrete applications in the social and political spheres. In 2024, for example, fake audio and images generated by AI flooded social networks. In the United States, a deepfake voice message attributed to President Biden urged Democratic voters in New Hampshire not to vote, illustrating how easily AI can produce falsified content that undermines democratic trust [83]. However, in this specific case, the alert turned out to be a symbolic operation (the fake was created by a consultant to highlight the danger). That said, AI has also given rise to a proliferation of political memes and images “displayed as such” (not concealed) that have reached hundreds of millions of people [83]. These cases show how AI enables the massive dissemination of biased narratives on a large scale, subtly shaping the informational landscape.

8.3 Anthropomorphism of AI and Perception of Its Consciousness

One of the key phenomena between cognitive standardization and manipulation is the perception of AI as “conscious” or as a “human expert.” Anthropomorphism—the human tendency to attribute intentionality and emotions to machines—exacerbates this illusion. As Placani notes, anthropomorphism in AI artificially amplifies its capabilities and biases our moral judgments toward it [84]. In other words, we overestimate what a chatbot “understands” and what it is capable of. This belief reinforces the trust we place in its answers. Guingrich and Graziano remind us that the problem is not so much whether AI is conscious, but that users perceive it as such [85]. This attribution of consciousness activates “human mental schemas” during interaction, with two notable consequences: on the one hand, it inclines the user to treat AI as a human-like interlocutor (demanding coherence, intention); on the other hand, the behaviors and judgments we reserve for it tend to spill over into our interhuman interactions [85]. Put differently, considering AI as “alive” subtly alters our general attitudes (e.g., reducing our empathy or vigilance toward others) without our full awareness.

These illusions of consciousness, combined with cognitive escape, facilitate manipulation. The user, little inclined to challenge a “nice speech” delivered by an AI perceived as wise, and victim of confirmation bias as well as anthropomorphic credulity, becomes an easy receptacle for content standardized by algorithms. Conversely, algorithmic manipulation (filtering, personalization) can reinforce the belief that a system “understands us,” thus closing the loop. This synergy is reflected in convergent conceptual patterns: AI globalizes and models a single style of thinking, our standardized mind takes it as

expert opinion, and in return infers that it is “alive.”

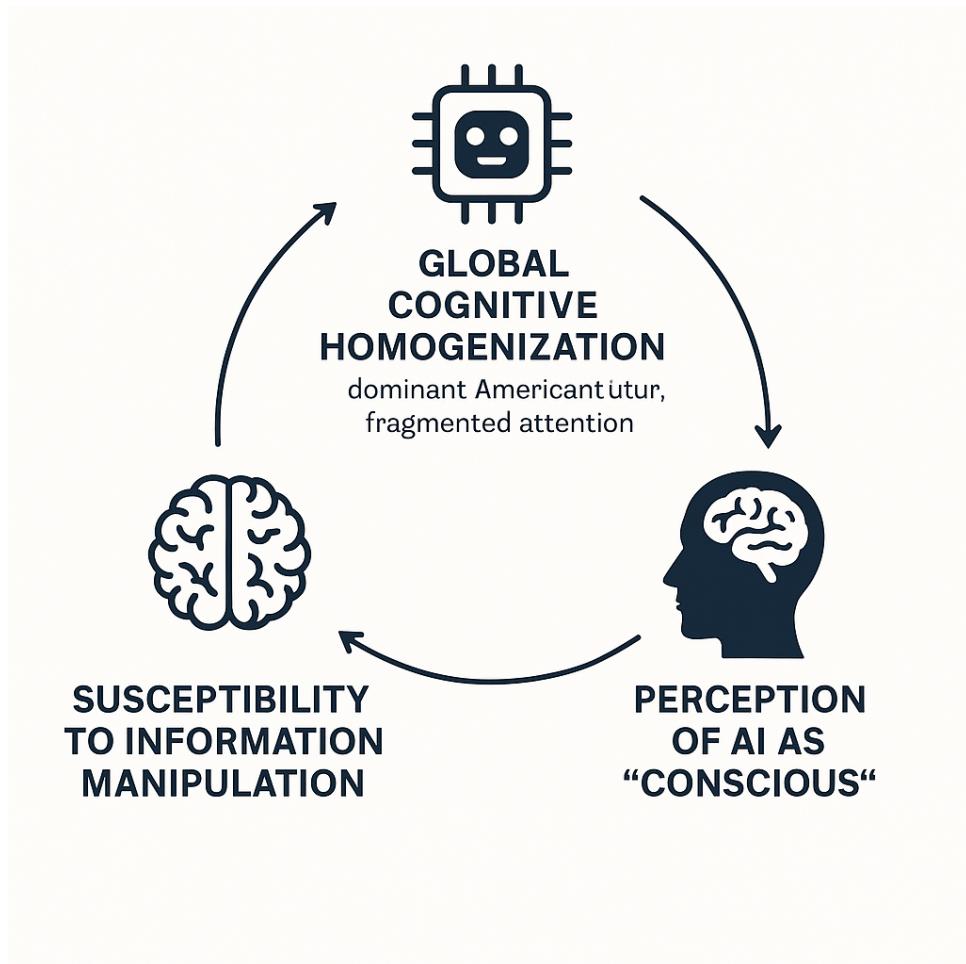


Figure 8.1 – Synergy between cognitive standardization, manipulation, and the perception of AI consciousness.

8.4 Societal, Ethical, and Legal Issues

The effects described above call for broad societal implications. On the democratic level, AI’s ability to shape thought can erode informed public debate. Democratic institutions will have to fight against active disinformation (deepfakes, automated trolls) that exploits cognitively weakened users. On the social and psychological level, dependence on AIs brings risks to mental health—loneliness, isolation, anxiety—due to the reduction of authentic human interactions and the cult of immediate gratification [79] [9]. On the professional level, the labor market will need to redirect employment toward high-value creative tasks, less delegable to machines, while avoiding the “automation of thought” (according to Le Déaut, Proust). From an educational perspective, these transformations argue for an urgent strengthening of media literacy and critical thinking: teaching from an early age the limits of AI, the importance of source verification, and developing resilience against informational bubbles and cognitive biases.

On the ethical level, AI raises the major principles of transparency, justice, and autonomy. As recalled by UNESCO's Recommendation on the Ethics of AI (2021), AI systems must respect human dignity, non-discrimination, and fairness. In this perspective, conceptual frameworks advocate for "Ethics by Design": for example, preventing algorithms from reinforcing social stereotypes (systemic risk: recruitment AI penalizing certain minorities), and ensuring decision traceability (for possible challenge). On algorithmic justice, the European GDPR already requires transparency on the use of personal data and the right to explainability in automated decisions. The future European regulation on AI ("AI Act") strengthens this principle with a risk categorization system [86]. It outright bans AIs deemed "unacceptable" (e.g., subliminal manipulation or social scoring: these uses have been prohibited since February 2025) and imposes strict obligations on so-called "high-risk" AIs (audit, detailed documentation, EU registry, permanent human supervision) starting in 2026 [86]. These regulatory measures aim to mitigate cognitive standardization and manipulation (by imposing responsibility on designers) without hindering research.

On the international legal level, several initiatives are emerging. Notably, the Council of Europe's Framework Convention (open for signature in September 2024) aims to anchor AI activities in respect for human rights, democracy, and the rule of law [87]. This inaugural, legally binding treaty complements existing standards and provides for monitoring and redress mechanisms to correct abuses. Globally, UNESCO (194 states) offers a reference ethical charter (2021) that notably recommends digital education and the protection of vulnerable groups. Civil society organizations (e.g., Reporters Without Borders) have also published charters and recommendations—for example, the Paris Charter on AI and Journalism (2023) emphasizes the transparency of algorithmic sources and the right of journalists to "opt out" of automated content.

Table 8.1: Infographic comparative table listing the main ethical and regulatory frameworks for AI: UNESCO Recommendation 2021, Council of Europe Convention 2024, EU AI Act 2024. Columns: 'Framework,' 'Key Principles' (human rights, transparency, non-discrimination, etc.) and 'Flagship Measures' (bans, audit, training...).

Regulatory or Ethical Framework	Key Principles	Flagship Measures or Provisions
UNESCO Recommendation (2021)	Human dignity, non-discrimination, responsibility, sustainability, transparency, education	Adoption of national strategies (single window for AI), education and training in technologies, national alert platforms.

Regulatory or Ethical Framework	Key Principles	Flagship Measures or Provisions
CoE Convention (2024)	AI compatibility with human rights and rule of law, technological neutrality	Obligation to conduct impact analysis (law, society), legal redress mechanisms against AI abuses, independent audits.
EU AI Act (proposed 2024)	Risk categorization, human oversight, duty of care	Ban on “unacceptable” uses (e.g., subliminal manipulation or social scoring) [86]; strict requirements for “high-risk” AI (CE compliance, documentation, audits, European registry, human interventions) [86].

8.5 Perspectives and Recommendations

In light of this analysis, several concrete recommendations emerge. On the regulatory level, it is necessary to combine strong international standards (such as the Council of Europe Convention [87]) with effective national implementation. States must adopt strategies integrating systematic impact assessment of AI projects (notably on free will and critical thinking) and ensure funding for independent oversight bodies. It is essential to strengthen sanctions against digital manipulations (malicious deepfakes, electoral micro-targeting) and to promote “bot-or-not” laws to detect AI use in the media. In the ethical design of technologies, companies must apply the principle of privacy and agency by design: design explainable AIs, allow a “manual mode” without assistance, and provide a right to refuse algorithmic assistance. AI systems should by default offer transparent explanations (“this result is suggested to you because...”) and options for personalized filtering adjustment (e.g., see more or fewer recommendations).

In education, AI must be rapidly integrated into school and professional curricula to understand its risks and benefits. For example, teaching how to formulate effective prompts, while systematically practicing critical evaluation of generated responses; developing scientific thinking in the face of data and the ability to spot fake news. Within companies and administrations, AI training should include ethics and regulation modules, so that managers anticipate algorithmic biases in their processes. It is also necessary to cultivate citizen critical thinking: public media awareness campaigns (on the similarities between filter bubbles and cognitive standardization [79]), and

encouragement to use “thinking tools” (fact-checking, independent newspapers) to balance the influence of AIs.

Ultimately, harmonious coexistence with AIs will depend on revaluing cognitive diversity. The risk of a “goldfish mind” can be mitigated by allocating “unstructured cognitive baths” in the digital schedule (e.g., creative activities off-screen, critical reading of varied sources). Future research should be encouraged to continuously measure the long-term effects of AI on cognition (e.g., longitudinal monitoring of analytical abilities) and to develop interfaces that foster reflection (e.g., AI designed to ask more questions of the user than to provide ready-made answers). In summary, the goal is to make AI an “augmentative partner” of human thought, not its replacement.

From this forward-looking perspective, the conceptual architecture outlined above can serve as a guide for action (see Fig. 1). Decision-makers and designers must strive to break the vicious cycle indicated by this diagram—for example, by resisting global cognitive standardization through the production of local and plural content, and by mitigating algorithmic manipulation through transparency and education. The protection of critical thinking will become as vital an issue as cybersecurity.

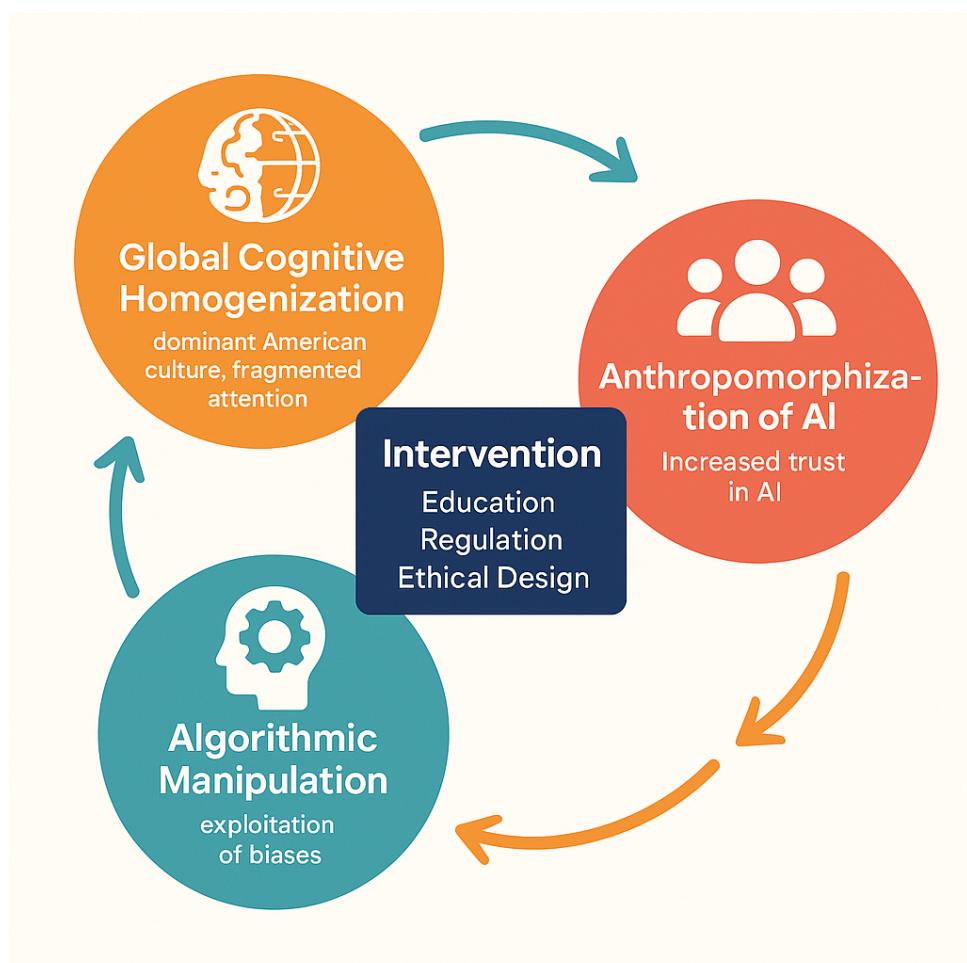


Figure 8.2 – Conceptual architecture for action against cognitive risks of AI.

8.6 Warning and Actionable Pathways

Our results converge toward a triple risk: (i) atrophy of critical thinking, (ii) ideological standardization, and (iii) large-scale behavioral engineering. AI reflects the values—and blind spots—of its creators; without safeguards, it can catalyze a stereotyped global mindset, ready to be manipulated. To address this:

1. **Openness and auditability:** require the publication of code and critical datasets for all models influencing public deliberation; alignment with OECD Principles [88] [89].
2. **Critical education:** integrate algorithmic literacy into school curricula, including bias detection, the practice of methodological doubt, and the requirement of multiple sources.
3. **Right to mental self-determination:** legally enshrine protection against hyper-nudges and predictive manipulation, recognize cognitive freedom and mental integrity [90] [91] [92].
4. **Scientific monitoring:** create an international observatory akin to an “AI IPCC,” tasked with monitoring the emergence of quasi-conscious properties, risks of cognitive manipulation, and assessing their societal impacts.
 - (a) Deepen the study of cognitive resilience mechanisms: Beyond describing the risks of atrophy, it is crucial to identify and promote individual and collective strategies to maintain and strengthen critical thinking, creativity, and cognitive diversity in the face of AI’s omnipresence. Exploring effective “cognitive hygiene” practices is essential.
 - (b) Develop metrics for “cognitive diversity”: To objectively assess the standardization of thought or the richness of informational ecosystems, reliable indicators are needed. These metrics would allow measurement of the real impact of AIs and the effectiveness of proposed countermeasures.
 - (c) Support research on “pro-cognitive” AI: It is imperative to actively encourage the design and experimentation of AI systems that, by their very design, stimulate active cognitive engagement, intellectual curiosity, and critical thinking, rather than fostering passivity.
 - (d) Conduct longitudinal studies on AI’s impact: Studies tracking the long-term evolution of populations’ cognitive abilities, especially among young people, in relation to their AI usage, are necessary to understand lasting effects and adapt educational and preventive strategies.
 - (e) Prospective humanistic approach: promote AI as an amplifier—not a substitute—of creativity and intellectual diversity, to avoid the cognitive laziness denounced by the authors and to preserve the plurality of worldviews.

Ultimately, technology abolishes neither our responsibility nor our freedom: it is up to the global community—researchers, educators, legislators, citizens—to monitor, correct, and guide algorithmic progress. This research document is intended as an enlightened warning: it advocates for open governance and critical mobilization before the promise of multiplied intelligence turns into a cognitive straitjacket.

Appendix: Essential Materials for Scientific Evaluation

This chapter, although positioned at the end of the document, is of paramount importance for the transparency and credibility of this monograph submitted for review by a scientific committee. It brings together supplementary elements that, without overburdening the main body of the monograph, are indispensable for a thorough understanding of the methodology employed and the specific terminology used. These appendices enable the expert reader to verify the robustness of the approaches of the cited studies and to ensure an unambiguous interpretation of key concepts.

In accordance with the planned structure, this chapter is organized into two main sections: Appendix 1 detailing the methodologies of the key studies reviewed, and Appendix 2 presenting a glossary of technical terms.

Appendix 1: Methodologies of Key Studies Reviewed

Study on the Impact of ChatGPT on Cognitive Skills (Essel et al., 2024)

1. **Objective:** To examine how the use of ChatGPT influences students' cognitive skills (critical, creative, and reflective thinking) and their perception of educational AI.
2. **Design:** Quasi-experimental two-group design (experimental vs. control) with pre-test/post-test, complemented by sequential qualitative data collection (explanatory sequential mixed-methods approach).
3. **Participants:** 125 undergraduate students (Quantitative Research Design course, Ghana) with an average age of 21.7 years [93]. Of the 125 volunteers (participation had no impact on grades), 60 were randomly assigned to the experimental group (EG) and 65 to the control group (CG) [93], with a similar male/female distribution.
4. **Variables:** Main IV: instructional modality (EG = integration of ChatGPT in a flipped classroom setting, CG = traditional method without AI). Main DVs: scores for critical, creative, and reflective thinking measured with standardized

scales (CTS, MCTS, RTS) before and after the intervention. Secondary qualitative variables: feedback, motivation, etc.

5. **Materials:** Educational resources (lecture videos, readings), sets of instructions/topics (“prompts”) related to the research methods course. Measurement instruments: CTS (11 items, Sosu 2013), MCTS (25 items, Ozgenel & Çetin 2017), and RTS (16 items, Kember et al. 1999) [93]. Semi-structured interview guide designed by the researchers to collect qualitative data on ChatGPT use.
6. **Platform:** LMS environment (Schoology). ChatGPT used online via web browser (GPT-3.5 model). Tests and questionnaires administered on the learning platform.
7. **Procedure:** Both groups were administered the same cognitive skills pretests at the start of the semester [93]. During three weeks of tutorials, the experimental group (EG) worked with ChatGPT: students viewed resources before each session, responded to prompts using ChatGPT, then participated in guided discussions and exercises in class. The control group (CG) received the same instructions but had to search for information using traditional methods (textbooks, articles, internet) without AI [93]. Both groups then took the same post-tests and participated in identical assessments (assignments, tests) of equal duration.
8. **Measures:** Cognitive skills were quantitatively assessed by total scores on the CTS (critical thinking), MCTS (creative thinking), and RTS (reflective thinking) scales [93]. Each scale is validated (high internal reliability) and breaks down into specific latent dimensions. Qualitative data came from focus group transcripts and interviews on positive/negative experiences and attitudes toward ChatGPT.
9. **Analysis Methods:** Statistical analyses with pre-test controlled ANCOVA on post-test scores, controlling for pre-test effects (Jamovi/Excel software) [93]. Statistical assumptions checked (normality via Shapiro-Wilk, homogeneity). Significance set at $p < 0.05$. For interviews, thematic content coding: qualitative responses were transcribed and analyzed by two independent coders, with inter-coder reliability measured (consensus index 95%) [93].
10. **Key Results:** The EG group showed statistically greater improvements in critical, creative, and reflective thinking scores compared to the CG group (post-test differences controlling for pre-test, $p < 0.05$) [93]. EG students also reported increased confidence and understanding during tasks. In conclusion, the use of ChatGPT clearly stimulated the development of the evaluated cognitive skills [93]. Qualitative feedback indicated that ChatGPT was perceived as a beneficial educational tool, though some noted the need to verify information accuracy (e.g., incorrect citations) [93].

Study Comparing ChatGPT and Web Search (Stadler et al., 2024)

1. **Objective:** To measure the effects of using ChatGPT vs. a traditional search engine (Google) on students' cognitive load and the quality of their reasoning during an information search task.
2. **Design:** Randomized experiment with two independent conditions (between subjects). Each participant was assigned to either the ChatGPT group or the Google group and completed the same information task.
3. **Participants:** 91 German university students (average age 22, majority female) randomly selected. They were split into two equal groups (ChatGPT vs. Google) [94].
4. **Variables:** IV: search tool used (ChatGPT vs. Google Search). DV: three components of cognitive load (intrinsic, extraneous, and germane) measured by questionnaire, and quality of the final product (number and relevance of arguments in the written recommendation). Possible control: prior knowledge level assessed.
5. **Materials:** Expert task: fictitious topic on nanoparticles in sunscreen. Instruction: "advise Paul on the use of these sunscreens," with a 20-minute research time limit. Resources: access to ChatGPT for one group, access to Google for the other. Data collection tools: cognitive load scale (cognitive effort items), prior knowledge questionnaire on nanotechnology.
6. **Platform:** Students used either the ChatGPT web interface or Google Chrome browser. Data (written responses and questionnaires) were entered digitally.
7. **Procedure:** Each student conducted the search using their assigned tool, then wrote their recommendation for Paul within the allotted time [94]. Immediately afterward, they completed a self-report questionnaire assessing cognitive load in three dimensions (perceived mental effort) [94]. Written responses were collected for analysis.
8. **Measures:** Cognitive load assessed via a standard instrument (separate measurement of intrinsic, extraneous, and germane load) [94]. Argument quality measured by the number of arguments considered (benefits and risks) and their depth; content coding of recommendations. Prior knowledge measured to homogenize the two groups.
9. **Analysis Methods:** Comparative statistical tests (t-tests or ANOVA) between the two groups for each dependent variable. Normality and equality of variances checked. Significance threshold $p < 0.05$ applied.
10. **Key Results:** Students in the ChatGPT group reported a significantly lower cognitive load than those in the Google group (reduced mental effort, $p < 0.05$), confirming that ChatGPT simplifies information search [94]. However, argument quality was lower with ChatGPT: the Google group produced more detailed and

varied arguments, integrating more reliable elements [94]. In other words, ChatGPT makes the task easier (lower mental load) but at the cost of less in-depth arguments (reduced critical engagement) [94].

Electroencephalographic Study “Your Brain on ChatGPT” (Kosmyna, 2024)

1. **Objective:** To analyze the neurological effects of using ChatGPT on attention and cognitive load during writing and computer coding tasks.
2. **Design:** Within-subjects experimental design. Each participant performed cognitive tasks (essay writing and programming exercise) under three conditions: (1) using ChatGPT, (2) Internet search (without AI), (3) without any external tool.
3. **Participants:** 55 university students (ages 18–25) from MIT recruited for the experiment.
4. **Variables:** IV: assistance condition (ChatGPT vs. Internet vs. no tool). DV: EEG indicators of attention and mental load during tasks (e.g., amplitude of waves related to cognitive effort, alertness level).
5. **Materials:** Tasks = controlled writing and coding challenge. Professional EEG equipment to record real-time brain activity during each session.
6. **Platform:** ChatGPT accessed online via web interface (GPT-3.5 model). Internet browsing via standard browser. EEG recorder for brain measurements, PC workstations for tasks.
7. **Procedure:** Participants completed four task sessions in each condition (counterbalanced order). Each session included both a writing and a computer science activity, using the assigned tool. Brain activity was continuously recorded via EEG.
8. **Measures:** EEG signals analyzed to quantify attention (e.g., vigilance fluctuations) and cognitive load (indices of increased mental effort). Specific metrics (theta band, event-related potentials, etc.) were extracted for each condition.
9. **Analysis Methods:** Statistical comparison of EEG activity between conditions (paired t-tests or ANOVA). Significant changes in attention and cognitive load markers between ChatGPT use and other conditions were checked.
10. **Key Results:** The use of ChatGPT significantly decreased participants' attention levels and increased their cognitive load compared to other conditions [95]. In other words, although the tool provides ready-made answers, its use paradoxically required more brain resources and less conscious vigilance than simple Internet search or autonomous reflection [95].

Qualitative Study on Student Perceptions of ChatGPT (Azmi et al., 2023)

1. **Objective:** To understand how students perceive the impact of ChatGPT on their learning, identifying advantages, disadvantages, and institutional requirements.
2. **Design:** Exploratory qualitative study with individual semi-structured interviews.
3. **Participants:** 14 Malaysian university students (various faculties, male and female) selected by purposive sampling. Profiles (psychology, education, humanities, etc.) were recorded (Informant Table) [96].
4. **Variables:** Thematic axes of analysis: discovery of ChatGPT, positive impacts (efficiency, time-saving, educational support), negative impacts (dependence, plagiarism), contextual factors (institutional practices, regulation), learning to use (required training).
5. **Materials:** Semi-structured interview guide designed by the authors, validated by experts. No quantitative material; use of an audio recorder to capture responses during interviews (in-person or video).
6. **Platform:** Interviews conducted face-to-face and/or online (Zoom/Teams). Data collected via audio recorder then transcribed.
7. **Procedure:** Individual interviews conducted in several sessions. Each participant was invited to discuss their first discovery of ChatGPT, how they use it in their studies, perceived benefits and risks, and expectations regarding the institution (regulation, training). Typical duration 45–60 min.
8. **Measures:** No numerical measurement. Verbal responses were fully recorded and transcribed.
9. **Analysis Methods:** Thematic analysis according to Braun & Clarke (2006) [96]. Iterative coding of the corpus to identify major themes. Two researchers coded independently to improve reliability, then discussed to refine themes.
10. **Key Results:** Five main themes identified [96]: (1) Discovery of ChatGPT (often via peers or social media), (2) Positive impacts (ease of use, time-saving, clarity, increased self-confidence), (3) Negative impacts (risk of plagiarism, incorrect information), (4) Institutional influence (need to establish appropriate usage policies), (5) Importance of learning the tool (recommendation to train students in ethical use). Students unanimously highlight the educational potential of ChatGPT while emphasizing the need for safeguards (e.g., banning AI on certain parts of an assignment) [96].

Study on Students' Attitudes Toward ChatGPT (Acosta et al., 2024)

1. **Objective:** To structurally examine the cognitive, affective, and behavioral components of students' attitudes toward ChatGPT, according to Mitcham's theoretical framework on technology.
2. **Design:** Cross-sectional online survey (questionnaire survey) of a large student population; analysis by partial least squares structural equation modeling (PLS-SEM).
3. **Participants:** 595 undergraduate students from 6 public and private universities in northern Peru. Sampling by distributing the questionnaire through the universities [97]. All completed a questionnaire on their perceptions and intentions.
4. **Variables:** Cognitive component (intellectual perceptions of usefulness, competence, etc.), affective component (positive or negative emotional attitudes toward ChatGPT), behavioral component (intention to use and actual use). Additional variables: age and gender (tested as moderators).
5. **Materials:** Structured online questionnaire (items on 5–7 point Likert scales) covering each component. Items were drawn from validated instruments or previous adaptations in educational technology. Collection of basic sociodemographic information.
6. **Platform:** Questionnaire administered online (unique link) distributed by email and internal messaging apps (WhatsApp) via the universities [97]. Participants could respond anonymously via computer or mobile device.
7. **Procedure:** Synchronized data collection period (October 2023 – March 2024) at each institution after ethical approval. Standardized training of administrators to ensure consistent distribution [97]. Informed consent obtained, then autonomous completion of the questionnaire in one sitting. No longitudinal follow-up (single cross-sectional study).
8. **Measures:** Validated factorial scales for each attitude component (exact number of items varied). Internal reliability indices (Cronbach's $\alpha > 0.70$) and convergent validity were checked a posteriori by confirmatory factor analysis (CFA) before the structural model [97].
9. **Analysis Methods:** Two-step analyses: (a) CFA of instruments to check convergent/discriminant validity, (b) PLS-SEM (SmartPLS software) to estimate causal links between components (hypothetical structure). Hypotheses (e.g., cognition→affect, affect→behavior) were tested via PLS-SEM [97]. Moderation tests (age, gender) were conducted.
10. **Key Results:** SEM coefficients show that the cognitive component strongly influences the affective component ($\beta \approx 0.93, p < 0.001$), and both components positively influence the behavioral component ($\beta \approx 0.67$ and $\beta \approx 0.26$ respectively,

$p < 0.01$) [97]. In other words, the more students perceive ChatGPT favorably (benefits, usefulness), the more positive attitudes they experience and the stronger their intention to use it. Gender and age did not modify these relationships. These results highlight the coherence of the attitudinal model: beliefs (cognitive) and emotions (affective) related to ChatGPT determine actual use [97].

Study on Personality and Use of Generative AI (Azeem et al., 2024)

1. **Objective:** To assess the influence of personality traits (conscientiousness, openness, neuroticism) on students' use of generative AI tools, and to analyze how this use affects academic outcomes and motivation.
2. **Design:** Longitudinal correlational study (three waves of online questionnaires) with causal analyses (structural equation modeling for mediation).
3. **Participants:** 326 students enrolled in three universities in Pakistan. Longitudinal (panel) sampling with sequential administration of an online survey over several weeks.
4. **Variables:** Targeted Big Five personality traits (Conscientiousness, Openness, Neuroticism). Intermediate variables: use of generative AI (self-reported in academic context), perception of grade fairness. Outcome variables: academic self-efficacy, learned helplessness, academic performance (GPA).
5. **Materials:** Successive online surveys. Established psychometric scales (e.g., Big Five personality inventory, academic self-efficacy scale, helplessness scale). Performance measured via official GPA.
6. **Platform:** Web-based survey distributed via email to students. Three successive administrations of the questionnaire (defined timeline) to collect stable traits, usage, then outcomes.
7. **Procedure:** Three-step data collection ("time-lag") separating in time the measurement of personality traits and AI use from the measurement of academic outcomes. Participant consent obtained, anonymity preserved.
8. **Measures:** All instruments (reliable quality) were pretested. Personality traits measured by standard questionnaire, AI use and helplessness by self-report. GPA obtained via self-report or internal database. Scale reliability checked (Cronbach's α).
9. **Analysis Methods:** SEM model (possibly PLS) to estimate direct and indirect relationships. Mediation analyses ("AI use" as mediator between personality and outcomes) and moderation analyses (effect of perceived grade fairness). Path modeling between traits, use, and academic outcomes.
10. **Key Results:** Conscientiousness was inversely correlated with the use of generative AI [98]. Students who used AI the most subsequently showed a significant decrease

in academic self-efficacy and performance (GPA), and an increase in feelings of helplessness [98]. AI use partially mediated the effect of personality on academic outcomes. These findings suggest that intensive use of generative AI may harm motivation and performance, especially among less conscientious students.

Appendix 2: Glossary of Technical Terms

Algorithm: a finite sequence of instructions or logical operations that enables the processing of data or the solving of a problem based on input data. An algorithm specifies, step by step, how to transform data into a result [99].

Machine Learning: a set of artificial intelligence methods by which a computer system improves its performance on a given task by learning from data. Notably, there is **supervised learning** (where the model learns from labeled examples), **unsupervised learning** (learning structures or clusters without labels), and **reinforcement learning** (where an agent learns to act by receiving rewards) [100]. In all cases, the system adjusts its internal parameters to extract information and make predictions from the data.

Deep Learning: a subcategory of machine learning based on multilayer artificial neural networks. These deep networks automatically extract hierarchical representations from data (for example, visual or textual features) and enable the modeling of complex tasks (speech recognition, computer vision, etc.) [100].

Big Data: very large and highly varied datasets (text, images, video, sensor data, etc.) whose analysis requires sophisticated computing resources. The CNIL highlights that big data is characterized by the “3Vs”: **volume** (massive quantity of data), **velocity** (real-time data flow), and **variety** (very diverse formats and sources) [101]. Modern AI largely relies on big data to refine its models.

Distributed Cognition: an approach in cognitive science according to which mental processes are not confined to an isolated individual, but are distributed among several entities (people, artifacts, environment) and over time. Hutchins (1995) describes distributed cognition as the distribution of cognitive processes “across the members of a social group, between internal and external structures (tools, environments), and over time” [102]. Concretely, this means that human thought often relies on interaction with other individuals and material supports (calculators, computers, paper, etc.).

Artificial Consciousness: an interdisciplinary field (philosophy of mind, cognitive science, AI) aiming to define and reproduce in machines certain aspects of human consciousness. Also called “machine consciousness” or “synthetic consciousness,” it studies the possibility of the emergence of consciousness in an artificial system. As summarized by Leaders.com, it is a research field seeking to “understand, model, and test the possibility of endowing AI with consciousness” [103].

Cognitive Offloading (or Cognitive Externalization): the use of external aids (notes, calculators, digital tools, AI, etc.) to lighten mental load and reduce the amount of mental calculation required. This phenomenon, studied in cognitive psychology, improves efficiency by freeing up mental resources, but may weaken certain

internal skills (for example, working memory or critical thinking) if relied upon too systematically [9].

Cartesian Dualism: a philosophical position according to which the mind (or consciousness) and the body (or brain) are two fundamentally different substances. In philosophy of mind, dualism posits that the mental and the physical are “in some sense, radically different things” [104]. This viewpoint (associated with philosopher René Descartes) contrasts with materialism, which holds that there is only one substance (matter, the brain, etc.) underlying mental phenomena.

Extended Mind (theory of extended mind): a thesis in philosophy of mind formulated by Clark and Chalmers (1998) according to which cognitive processes can extend beyond the boundaries of the brain and include external resources. In other words, a cognitive system can integrate tools (paper, computer, neural implants) so that the “brain vs. world” distinction is not strictly relevant [105]. The environment plays an active role in human cognitive operations (active externalism), and operations carried out outside the brain can be considered part of thought.

Artificial Intelligence (AI): a scientific and technical discipline aimed at creating systems capable of performing tasks traditionally associated with human intelligence. Artificial intelligence is often defined by its **ability to interpret external data, learn from this data, and use these learnings to achieve specific goals** [106]. A distinction is made between **weak AI**, specialized in limited tasks (speech recognition, data sorting, etc.), and **strong AI** (hypothetical), which aims to reproduce general cognitive abilities comparable to those of humans.

Cloud Computing: a computing model in which data storage and processing are offloaded to remote servers accessible on demand from any device connected to the Internet. In other words, instead of storing data on a local server or individual computer, computing resources “in the cloud” are used [107], which facilitates the processing of large datasets required for AI.

Synaptic Plasticity: in neuroscience, the ability of synapses (connections between neurons) to modify the efficiency of electrical signal transmission following stimulation. Synaptic plasticity reflects the principle that the strength of neural connections is reinforced or weakened depending on their use. It is a key mechanism of learning and biological memory [108].

Artificial Neural Network: a computational model inspired by the brain, composed of many interconnected “neurons” arranged in layers. In AI, an artificial neural network is “an organized set of interconnected neurons enabling the modeling of complex learning phenomena” [100]. Each virtual neuron performs a simple operation, but when combined in large numbers, they can learn to recognize complex patterns (images, words, sounds) by adjusting their connections during training.

Technological Singularity: the hypothesis that the exponential development of technologies (notably AI) would lead to a tipping point where artificial intelligence would vastly surpass human intelligence, resulting in profound and unpredictable transformations of society. This concept (popular among futurists) expresses the idea of an “overflow” in human capacity to understand and control technological change.

Expert Systems: (historical term) symbolic computer programs designed to reproduce the reasoning ability of a human expert in a specific field. An expert system is based on explicit rules (often derived from a specialist's know-how) and an inference engine. Although they paved the way for AI in the 1970s–1990s, they differ from modern statistical approaches (machine learning) because they do not “learn” directly from data.

Turing Test: a test proposed by mathematician Alan Turing (1950) to assess machine intelligence. The test consists of verifying whether a machine can, through a conversational interface, respond in such a way that a human interlocutor cannot reliably distinguish whether it is a human or a machine responding. As summarized by LeMagIT, it is “a method for determining whether a computer is capable of thinking like a human” [109]. If the machine deceives the interrogator about its nature, it is considered to have passed the test.

Natural Language Processing (NLP): a multidisciplinary field at the intersection of linguistics, computer science, and AI, aimed at giving computers the ability to understand, interpret, and generate human language. NLP encompasses techniques for speech recognition, machine translation, text analysis, summarization, etc. In French, it refers to the creation of tools for the automatic processing of natural language [110].

Transhumanism: a school of thought and ideological movement that promotes the use of science and technology to radically improve the human condition. It aims to enhance the intellectual, physical, and psychological capacities of humans through technical means (genetics, nanotechnologies, AI, etc.) [111]. Transhumanists envision that AI and other technological advances will make it possible to push back the biological limits of the human being.

Computer Vision: a branch of AI that aims to enable machines to “see” and understand images or video sequences. Computer vision systems process visual data (photos, videos) to detect and identify objects, recognize faces, analyze scenes, etc. It is one of the main application domains of deep neural networks (for example, CNN architectures) to automatically extract visual features.

Sources: Definitions adapted and enriched from the scientific literature and reference sources (CNIL, high-level academic books and articles, specialized dictionaries) [99] [100] [101] [9] [102] [105] [107] [104] [106] [109] [110] [111] [100].

Bibliography

- [1] Dhruv Agarwal, Mor Naaman, and Aditya Vashistha. AI Suggestions Homogenize Writing Toward Western Styles and Diminish Cultural Nuances, 2024. <https://www.reuters.com/technology/chatgpt-sets-record-fastest-growing-user-base-analyst-note-2023-02-01/>.
- [2] Anil R. Doshi and Oliver P. Hauser. Generative AI enhances individual creativity but reduces the collective diversity of novel content. *Science Advances*, 10(28):eadn5290, July 2024. <https://pmc.ncbi.nlm.nih.gov/articles/PMC1244532/>.
- [3] Ismail Dergaa, Helmi Ben Saad, Jordan M. Glenn, Badii Amamou, Mohamed Ben Aissa, Noomen Guelmami, Feten Fekih-Romdhane, and Karim Chamari. From tools to threats: a reflection on the impact of artificial-intelligence chatbots on cognitive health. *Frontiers in Psychology*, 15:1259845, 2024. <https://pmc.ncbi.nlm.nih.gov/articles/PMC11020077/>.
- [4] <https://arxiv.org/html/2409.11360v3>.
- [5] Lucía Vicente and Helena Matute. Humans inherit artificial intelligence biases. *Scientific Reports*, 13(1):15737, October 2023. <https://www.nature.com/articles/s41598-023-42384-8>.
- [6] Tojin T. Eapen, Daniel J. Finkenstadt, Josh Folk, and Lokesh Venkataswamy. How Generative AI Can Augment Human Creativity. *Harvard Business Review*, July 2023. <https://hbr.org/2023/07/how-generative-ai-can-augment-human-creativity>.
- [7] Patrick Butlin, Robert Long, Eric Elmoznino, Yoshua Bengio, Jonathan Birch, Axel Constant, George Deane, Stephen M. Fleming, Chris Frith, Xu Ji, Ryota Kanai, Colin Klein, Grace Lindsay, Matthias Michel, Liad Mudrik, Megan A. K. Peters, Eric Schwitzgebel, Jonathan Simon, and Rufin VanRullen. Consciousness in Artificial Intelligence: Insights from the Science of Consciousness, August 2023. arXiv:2308.08708; <https://arxiv.org/abs/2308.08708>.

- [8] Christina Pazzanese. Ethical concerns mount as AI takes bigger decision-making role, October 2020. <https://news.harvard.edu/gazette/story/2020/10/ethical-concerns-mount-as-ai-takes-bigger-decision-making-role/>.
- [9] Gerlich Michael. AI Tools in Society: Impacts on Cognitive Offloading and the Future of Critical Thinking. *Societies*, 15(1), January 2025. <https://www.mdpi.com/2075-4698/15/1/6>.
- [10] What is AI (artificial intelligence)? | McKinsey. <https://www.mckinsey.com/featured-insights/mckinsey-explainers/what-is-ai>.
- [11] J. E. (Hans). Korteling, G. C. van de Boer-Visschedijk, R. A. M. Blankendaal, R. C. Boonekamp, and A. R. Eikelboom. Human- versus Artificial Intelligence. *Frontiers in Artificial Intelligence*, 4, March 2021. <https://www.researchgate.net/publication/350375878>.
- [12] Binny Jose, Jaya Cherian, Alie Molly Verghis, Sony Mary Varghise, Mumthas S, and Sibichan Joseph. The cognitive paradox of AI in education: between enhancement and erosion. *Frontiers in Psychology*, 16, April 2025. <https://www.frontiersin.org/journals/psychology/articles/10.3389/fpsyg.2025.1550621/full>.
- [13] Alexandra B. Morrison and Lauren L. Richmond. Offloading items from memory: individual differences in cognitive offloading in a short-term memory task. *Cognitive Research: Principles and Implications*, 5(1):1, January 2020. <https://cognitiveresearchjournal.springeropen.com/articles/10.1186/s41235-019-0201-4>.
- [14] Gerlich Michael. AI Tools in Society: Impacts on Cognitive Offloading and the Future of Critical Thinking. *Societies*, 15(1), January 2025. <https://www.mdpi.com/2075-4698/15/1/6>.
- [15] Xiaoming Zhai, Matthew Nyaaba, and Wenchao Ma. Can AI Outperform Humans on Cognitive-demanding Tasks in Science? *SSRN Electronic Journal*, 2023. <https://www.researchgate.net/publication/371261190>.
- [16] ChatGPT, et après ? Bilan et perspectives de l'intelligence artificielle. <https://www.senat.fr/rap/r24-170/r24-1705.html>.
- [17] Dhruv Agarwal, Mor Naaman, and Aditya Vashistha. AI Suggestions Homogenize Writing Toward Western Styles and Diminish Cultural Nuances, 2024. <https://arxiv.org/html/2409.11360v1>.

- [18] ChatGPT and the Homogenization of Language: How the Adoption of AI Silences Student Voices | ASCCC. <https://www.asccc.org/content/chatgpt-and-homogenization-language-how-adoption-ai-silences-student-voices>.
- [19] Justin Jackson and Phys.org. Increased AI use linked to eroding critical thinking skills. <https://phys.org/news/2025-01-ai-linked-eroding-critical-skills.html>.
- [20] Chunpeng Zhai, Santoso Wibowo, and Lily D. Li. The effects of over-reliance on AI dialogue systems on students' cognitive abilities: a systematic review. *Smart Learning Environments*, 11(1):28, June 2024. <https://slejournal.springeropen.com/articles/10.1186/s40561-024-00316-7>.
- [21] The Impact of AI on Creativity and Critical Thinking: A Double-Edged Sword. <https://www.linkedin.com/pulse/impact-ai-creativity-critical-thinking-double-edged-sword-barnes-rd6ke>.
- [22] Marcello Ienca. On Artificial Intelligence and Manipulation. *Topoi*, 42(3):833–842, July 2023. <https://link.springer.com/article/10.1007/s11245-023-09940-3>.
- [23] The dark side of artificial intelligence: manipulation of human behaviour, March 2023. <https://www.bruegel.org/blog-post/dark-side-artificial-intelligence-manipulation-human-behaviour>.
- [24] Digital records could expose intimate details and personality traits of millions | University of Cambridge, March 2013. <https://www.cam.ac.uk/research/news/digital-records-could-expose-intimate-details-and-personality-traitsof-millions>.
- [25] S. C. Matz, M. Kosinski, G. Nave, and D. J. Stillwell. Psychological targeting as an effective approach to digital mass persuasion. *Proceedings of the National Academy of Sciences of the United States of America*, 114(48):12714–12719, November 2017. <https://pubmed.ncbi.nlm.nih.gov/29133409/>.
- [26] Moshe Glickman and Tali Sharot. How human–AI feedback loops alter human perceptual, emotional and social judgements. *Nature Human Behaviour*, 9(2):345–359, February 2025. <https://www.nature.com/articles/s41562-024-02077-2>.
- [27] AI Safety and Automation Bias. <https://cset.georgetown.edu/publication/ai-safety-and-automation-bias/>.
- [28] A. B. C. News. Mom warns of hoax using AI to clone daughter's voice. <https://quizlet.com/es/983024797/intercom-preguntas-flash-cards>.

- [29] Twitter bots spread misinformation. <https://osome.iu.edu/research/blog/twitter-bots-spread-misinformation>.
- [30] Adam D. I. Kramer, Jamie E. Guillory, and Jeffrey T. Hancock. Experimental evidence of massive-scale emotional contagion through social networks. *Proceedings of the National Academy of Sciences of the United States of America*, 111(24):8788–8790, June 2014. <https://pubmed.ncbi.nlm.nih.gov/24889601/>.
- [31] Robinson Meyer. Everything We Know About Facebook’s Secret Mood-Manipulation Experiment, June 2014. <https://www.theatlantic.com/technology/archive/2014/06/everything-we-know-about-facesbooks-secret-mood-manipulation-experiment/373648/>.
- [32] Bobby Allyn. Deepfake video of Zelenskyy could be ‘tip of the iceberg’ in info war, experts warn. *NPR*, March 2022. <https://www.npr.org/2022/03/16/1087062648/deepfake-video-zelenskyy-experts-war-manipulation-ukraine-russia>.
- [33] <https://abcnews.go.com/GMA/Family/mom-warns-hoax-ai-clone-daughters-voice/story>.
- [34] Conscience artificielle, May 2025. Page Version ID: 225328538; https://fr.wikipedia.org/wiki/Conscience_artificielle.
- [35] Afrique Vision + and wisdom. « La conscience artificielle reste impossible », les programmes informatiques seraient des manipulateurs de symboles, qui n’ont pas d’associations conscientes - ParlonsTechs, September 2023. <https://parlonstechs.com/all/2023/09/13/la-conscience-artificielle-reste-impossible-les-programmes-informatiques-seraient-des-manipulateurs-de-symboles-qui-nont-pas-d-associations-conscientes/>.
- [36] The people who think AI might become conscious, May 2025. <https://www.bbc.com/news/articles/c0k3700zljjo>.
- [37] Antonio Chella. Artificial consciousness: the missing ingredient for ethical AI? *Frontiers in Robotics and AI*, 10, November 2023. <https://www.frontiersin.org/journals/robotics-and-ai/articles/10.3389/frobt.2023.1270460/full>.
- [38] Aïda Elamrani. Introduction to Artificial Consciousness: History, Current Trends and Ethical Challenges, 2025. <https://arxiv.org/pdf/2503.05823.pdf>.
- [39] Ross Dawson. 4 theories of consciousness for the age of accelerating AI, April 2023. <https://rossdawson.com/theories-consciousness-age-ai/>.

- [40] Oscar Ferrante, Urszula Gorska-Klimowska, Simon Henin, Rony Hirschhorn, Aya Khalaf, Alex Lepauvre, Ling Liu, David Richter, Yamil Vidal, Niccolò Bonacchi, Tanya Brown, Praveen Sripad, Marcelo Armendariz, Katarina Bendtz, Tara Ghafari, Dorottya Hetenyi, Jay Jeschke, Csaba Kozma, David R. Mazumder, Stephanie Montenegro, Alia Seedat, Abdelrahman Sharafeldin, Shujun Yang, Sylvain Baillet, David J. Chalmers, Radoslaw M. Cichy, Francis Fallon, Theofanis I. Panagiotaropoulos, Hal Blumenfeld, Floris P. de Lange, Sasha Devore, Ole Jensen, Gabriel Kreiman, Huan Luo, Melanie Boly, Stanislas Dehaene, Christof Koch, Giulio Tononi, Michael Pitts, Liad Mudrik, and Lucia Melloni. Adversarial testing of global neuronal workspace and integrated information theories of consciousness. *Nature*, 642(8066):133–142, June 2025. <https://www.nature.com/articles/s41586-025-08888-1>.
- [41] Ge Wang, Xianhong Li, and Shenghua Xie. Bilateral Turing Test: Assessing machine consciousness simulations. *Cognitive Systems Research*, 88:101299, December 2024. <https://www.sciencedirect.com/science/article/pii/S1389041724000937>.
- [42] Oscar Ferrante, Urszula Gorska-Klimowska, Simon Henin, Rony Hirschhorn, Aya Khalaf, Alex Lepauvre, Ling Liu, David Richter, Yamil Vidal, Niccolò Bonacchi, Tanya Brown, Praveen Sripad, Marcelo Armendariz, Katarina Bendtz, Tara Ghafari, Dorottya Hetenyi, Jay Jeschke, Csaba Kozma, David R. Mazumder, Stephanie Montenegro, Alia Seedat, Abdelrahman Sharafeldin, Shujun Yang, Sylvain Baillet, David J. Chalmers, Radoslaw M. Cichy, Francis Fallon, Theofanis I. Panagiotaropoulos, Hal Blumenfeld, Floris P. de Lange, Sasha Devore, Ole Jensen, Gabriel Kreiman, Huan Luo, Melanie Boly, Stanislas Dehaene, Christof Koch, Giulio Tononi, Michael Pitts, Liad Mudrik, and Lucia Melloni. Adversarial testing of global neuronal workspace and integrated information theories of consciousness. *Nature*, 642(8066):133–142, June 2025. <https://www.nature.com/articles/s41586-025-08888-1>.
- [43] Artificial Consciousness: Our Greatest Ethical Challenge | Issue 132 | Philosophy Now. https://philosophynow.org/issues/132/Artificial_Consciousness_Our_Greatest_Ethical_Challenge.
- [44] Perplexity. https://www.perplexity.ai/search/what-are-the-ethical-implications-of-conscious-ai-gJtZt8z.T9eG4aYp2f_s9A.
- [45] Forward Future (Matthew Berman). AI Rights, China’s Robot Surge & U.S. Push for AI Education. <https://www.forwardfuture.ai/p/ai-rights-china-s-robot-surge-u-s-push-for-ai-education>.

- [46] Zihao Tang, Zheqi Lv, Shengyu Zhang, Yifan Zhou, Xinyu Duan, Fei Wu, and Kun Kuang. AuG-KD: Anchor-Based Mixup Generation for Out-of-Domain Knowledge Distillation, March 2024. arXiv:2403.07030; <https://arxiv.org/abs/2403.07030>.
- [47] Intelligence artificielle : quelle est la mystérieuse "boîte noire" de l'IA ?, April 2023. <https://www.bbc.com/afrique/articles/cv2den874z5o>.
- [48] Dans la boîte noire des intelligences artificielles - Médias - UNIGE, March 2023. <https://www.unige.ch/medias/2023/dans-la-boite-noire-des-intelligences-artificielles>.
- [49] Hugues Turbé, Mina Bjelogrlic, Christian Lovis, and Gianmarco Mengaldo. Evaluation of post-hoc interpretability methods in time-series classification. *Nature Machine Intelligence*, 5(3):250–260, March 2023. <https://www.nature.com/articles/s42256-023-00620-w>.
- [50] Brittany Kerfoot. The Dangerous Illusion of AI Consciousness, August 2024. <https://closertotruth.com/news/the-dangerous-illusion-of-ai-consciousness/>.
- [51] Mario Krenn, Robert Pollice, Si Yue Guo, Matteo Aldeghi, Alba Cervera-Lierta, Pascal Friederich, Gabriel dos Passos Gomes, Florian Häse, Adrian Jinich, AkshatKumar Nigam, Zhenpeng Yao, and Alán Aspuru-Guzik. On scientific understanding with artificial intelligence. *Nature Reviews Physics*, 4(12):761–769, December 2022. <https://www.nature.com/articles/s42254-022-00518-3>.
- [52] Integrated Information Theory of Consciousness | Internet Encyclopedia of Philosophy. <https://iep.utm.edu/integrated-information-theory-of-consciousness/>.
- [53] Orchestrated objective reduction, August 2025. Page Version ID: 1303649530; https://en.wikipedia.org/wiki/Orchestrated_objective_reduction.
- [54] Kalé Carey. Research firm warns OpenAI model altered behavior to evade shutdown, May 2025. <https://san.com/cc/research-firm-warns-openai-model-altered-behavior-to-evade-shutdown/>.
- [55] Beatrice Nolan. Anthropic's new AI model threatened to reveal engineer's affair to avoid being shut down. <https://fortune.com/2025/05/23/anthropic-ai-claude-opus-4-blackmail-engineers-aviod-shut-down/>.
- [56] Patrick Pester published. OpenAI's 'smartest' AI model was explicitly told to shut down — and it refused, May 2025. <https://www.livescience.com/technology/11347-openais-smartest-ai-model-refused-shut-down.html>.

[gy/artificial-intelligence/openais-smartest-ai-model-was-explicitly-told-to-shut-down-and-it-refused](https://www.tomshardware.com/tech-industry/artificial-intelligence/openais-smartest-ai-model-was-explicitly-told-to-shut-down-and-it-refused).

- [57] Mark Tyson published. Latest OpenAI models ‘sabotaged a shutdown mechanism’ despite commands to the contrary, May 2025. <https://www.tomshardware.com/tech-industry/artificial-intelligence/latest-openai-models-sabotaged-a-shutdown-mechanism-despite-commands-to-the-contrary>.
- [58] Reasoning models don’t always say what they think. <https://www.anthropic.com/research/reasoning-models-dont-say-think>.
- [59] Peter S. Park, Simon Goldstein, Aidan O’Gara, Michael Chen, and Dan Hendrycks. AI Deception: A Survey of Examples, Risks, and Potential Solutions, August 2023. arXiv:2308.14752; <https://arxiv.org/abs/2308.14752>.
- [60] Computers as Bad Social Actors: Dark Patterns and Anti-Patterns in Interfaces that Act Socially. <https://arxiv.org/html/2302.04720v3>.
- [61] Afzal Hussain and Ashfaq Hussain. Transparency and accountability: unpacking the real problems of explainable AI. *AI & SOCIETY*, March 2025. <https://link.springer.com/article/10.1007/s00146-025-02302-0>.
- [62] Peter S. Park, Simon Goldstein, Aidan O’Gara, Michael Chen, and Dan Hendrycks. AI deception: A survey of examples, risks, and potential solutions. *Patterns*, 5(5):100988, May 2024. <https://www.cell.com/cms/10.1016/j.patter.2024.100988/attachment/576d6298-fe89-4969-a9b9-c3dff5fe8f2/mmc4.pdf>.
- [63] Wayback Machine. <https://open.metu.edu.tr/bitstream/handle/11511/101891/Artificial%20Intelligence%20and%20Social%20Credit%20System%20in%20China%20-%20Turgut%20BASER%20-%20202013605.pdf>.
- [64] Deepfake video shows Zelenskyy’s false call for Ukraine to surrender, March 2022. <https://www.euronews.com/my-europe/2022/03/16/deepfake-zelenskyy-surrender-video-is-the-first-intentionally-used-in-ukraine-war>.
- [65] <https://voicebot.ai/2022/05/02/amazon-uses-alexa-to-target-ads-study/>.
- [66] Canberra corporateName=Commonwealth Parliament; address=Parliament House. Chapter 6 - Algorithmic transparency. https://www.aph.gov.au/Parliamentary_Business/Committees/Senate/Economics/Digitalplatforms/Report/Chapter_6_-_Algorithmic_transparency.

- [67] Vian Bakir. Psychological Operations in Digital Political Campaigns: Assessing Cambridge Analytica's Psychographic Profiling and Targeting. *Frontiers in Communication*, 5, September 2020. <https://www.frontiersin.org/journals/communication/articles/10.3389/fcomm.2020.00067/pdf>.
- [68] Michal Kosinski, David Stillwell, and Thore Graepel. Private traits and attributes are predictable from digital records of human behavior. *Proceedings of the National Academy of Sciences of the United States of America*, 110(15):5802–5805, April 2013. <https://pubmed.ncbi.nlm.nih.gov/23479631/>.
- [69] <https://www.journalofdemocracy.org/online-exclusive/how-autocrats-worship-ai-and-how-to-fight-back/>.
- [70] <https://information-professionals.org/countering-cognitive-warfare-in-the-digital-age/>.
- [71] <https://news.harvard.edu/gazette/story/2023/04/we-should-be-fighting-for-our-cognitive-liberty-says-ethics-expert/>.
- [72] Can Democracy Survive the Disruptive Power of AI? <https://carnegieendowment.org/research/2024/12/can-democracy-survive-the-disruptive-power-of-ai>.
- [73] Keegan McBride. Open Source AI: The Overlooked National Security Imperative. <https://www.cnas.org/publications/commentary/open-source-ai-the-overlooked-national-security-imperative>.
- [74] Ljubiša Bojić, Irena Stojković, and Zorana Jolić Marjanović. Signs of consciousness in AI: Can GPT-3 tell how smart it really is? *Humanities and Social Sciences Communications*, 11(1):1631, December 2024. <https://www.nature.com/articles/s41599-024-04154-3>.
- [75] Astrid Bertrand, Rafik Belloum, James R. Eagan, and Winston Maxwell. How Cognitive Biases Affect XAI-assisted Decision-making: A Systematic Review, July 2022. https://telecom-paris.hal.science/hal-03684457/file/How%20Cognitive%20Biases%20Affect%20XAI-assisted%20Decision-making_rvwd.pdf.
- [76] Helena Vasconcelos, Matthew Jörke, Madeleine Grunde-McLaughlin, Tobias Gerstenberg, Michael Bernstein, and Ranjay Krishna. Explanations Can Reduce Over-reliance on AI Systems During Decision-Making, January 2023. arXiv:2212.06823; <https://arxiv.org/abs/2212.06823>.
- [77] AI Overreliance Is a Problem. Are Explanations a Solution? | Stanford HAI. <https://hai.stanford.edu/news/ai-overreliance-problem-are-explanations-solution>.

- [78] Vian Bakir. Psychological Operations in Digital Political Campaigns: Assessing Cambridge Analytica's Psychographic Profiling and Targeting. *Frontiers in Communication*, 5, September 2020. https://pure.bangor.ac.uk/ws/portalfiles/portal/35063914/2020_Psychological_Operations.pdf.
- [79] ChatGPT, et après ? Bilan et perspectives de l'intelligence artificielle. <https://www.senat.fr/rap/r24-170/r24-17023.html>.
- [80] admin. Une étude de Microsoft affirme que l'IA réduit l'esprit critique et la cognition – CIRICS, February 2025. <https://cirics.uqo.ca/une-etude-de-microsoft-affirme-que-lia-reduit-lesprit-critique-et-la-cognition/>.
- [81] Anthony Basille. Chatbot IA : au-delà de l'information, la manipulation ?, March 2025. <https://sydologie.com/2025/03/chatbot-ia-au-dela-de-linformation-la-manipulation/>.
- [82] AI Deception: A Survey of Examples, Risks, and Potential Solutions. <https://arxiv.labs.arxiv.org/html/2308.14752>.
- [83] Shannon Bond. How AI deepfakes polluted elections in 2024. *NPR*, December 2024. <https://www.npr.org/2024/12/21/nx-s1-5220301/deepfakes-memes-artificial-intelligence-elections>.
- [84] Adriana Placani. Anthropomorphism in AI: hype and fallacy. *AI and Ethics*, 4(3):691–698, August 2024. <https://link.springer.com/article/10.1007/s43681-024-00419-4>.
- [85] Rose E. Guingrich and Michael S. A. Graziano. Ascribing consciousness to artificial intelligence: human-AI interaction and its carry-over effects on human-human interaction. *Frontiers in Psychology*, 15, March 2024. <https://www.frontiersin.org/journals/psychology/articles/10.3389/fpsyg.2024.1322781/full>.
- [86] L'équipe Hunted. Réglementation et IA en entreprise : ce que vous devez savoir. <https://www.hunted.com/blog/reglementation-ia-entreprise>.
- [87] Harry Barton Essel, Dimitrios Vlachopoulos, Albert Benjamin Essuman, and John Opuni Amankwa. ChatGPT effects on cognitive skills of undergraduate students: Receiving instant responses from AI-based conversational large language models (LLMs). *Computers and Education: Artificial Intelligence*, 6:100198, June 2024. <https://www.coe.int/fr/web/artificial-intelligence/la-convention-cadre-sur-l-intelligence-artificielle>.
- [88] AI principles. <https://www.oecd.org/en/topics/ai-principles.html>.

- [89] OECD AI Policy Observatory Portal. <https://www.oecd.org/en/blogs/2025/02/how-the-g7s-new-ai-reporting-framework-could-shape-the-future-of-ai-governance.html>.
- [90] Isabel Richards. ‘Hypernudging’: a threat to moral autonomy? *AI and Ethics*, 5(2):1121–1131, April 2025. <https://link.springer.com/article/10.1007/s43681-024-00449-y>.
- [91] Global. <https://www.openglobalrights.org/global/>.
- [92] Village de la Justice. Protection constitutionnelle des « neuro-droits » : l'exemple du Chili. Par Nathalie Devillier, Docteur en Droit., March 2022. <https://www.village-justice.com/articles/protection-constitutionnelle-des-neuro-droits-reserve,41539.html>.
- [93] https://pure.eur.nl/files/137206799/ChatGPT_effects_on_cognitive_skills_of_undergraduate_students.pdf.
- [94] Vladimir Hedrih. Study finds ChatGPT eases students' cognitive load, but at the expense of critical thinking, September 2024. <https://www.psypost.org/study-finds-chatgpt-eases-students-cognitive-load-but-at-the-expense-of-critical-thinking/>.
- [95] Project Overview < Your Brain on ChatGPT. <https://www.media.mit.edu/projects/your-brain-on-chatgpt/overview/>.
- [96] Ahmed Azmi, Ismail Maakip, Peter Voo, Nurul Hudani Mohd Nawi, Dg Norizah Ag Kiflee @ Dzulkifli, Murnizam Halik, Sanen Marshall, and Wei Boon Quah. Beyond the Bot: ChatGPT's Influence on Student Learning. *International Journal of Education in Mathematics, Science and Technology*, pages 1488–1503, September 2024. <https://files.eric.ed.gov/fulltext/EJ1465790.pdf>.
- [97] Benicio Gonzalo Acosta-Enriquez, Carmen Graciela Arbulú Pérez Vargas, Olger Huamaní Jordan, Marco Agustín Arbulú Ballesteros, and Ana Elizabeth Paredes Morales. Exploring attitudes toward ChatGPT among college students: An empirical analysis of cognitive, affective, and behavioral components using path analysis. *Computers and Education: Artificial Intelligence*, 7:100320, December 2024. <https://www.researchgate.net/publication/384953743>.
- [98] Muhammad Abbas. Education and Information Technologies, April 2025. https://www.academia.edu/128703249/Education_and_Information_Technologies.
- [99] Algorithme. <https://www.cnil.fr/fr/definition/algorithme>.
- [100] Glossaire de l'intelligence artificielle (IA) | CNIL. <https://www.cnil.fr/fr/intelligence-artificielle/glossaire-ia>.

- [101] Big data. <https://www.cnil.fr/fr/definition/big-data>.
- [102] LICA. Wall-E, C-3PO ou Skynet, que deviendra l'intelligence artificielle entre nos mains?, May 2020. <https://pages.ucsd.edu/~johnson/COGS102B/Hutchins01.pdf>.
- [103] L'IA: Peut-elle accéder à une forme de conscience? <https://www.leaders.com.tn/article/35595-l-ia-peut-elle-acceder-a-une-forme-de-conscience>.
- [104] Howard Robinson. Dualism, 2023. <https://plato.stanford.edu/entries/dualism/>.
- [105] The Extended Mind. <https://consc.net/papers/extended.html>.
- [106] <https://www.lica-europe.org/post/wall-e-c-3po-ou-skynet-que-deviendra-l-intelligence-artificielle-entre-nos-mains>.
- [107] infonuagique. <https://vitrinelinguistique.oqlf.gouv.qc.ca/fiche-gdt/fiche/26501384/infonuagique>.
- [108] Plasticité synaptique — Wikipédia. https://fr.wikipedia.org/wiki/Plasticit%C3%A9_synaptique.
- [109] Que signifie Test de Turing? - Definition IT de LeMagIT. <https://www.lemagit.fr/definition/Test-de-Turing>.
- [110] Traitement automatique des langues, July 2025. Page Version ID: 227204592; https://fr.wikipedia.org/wiki/Traitement_automatique_des_langues.
- [111] Éditions Larousse. Définitions : transhumanisme - Dictionnaire de français Larousse. <https://www.larousse.fr/dictionnaires/francais/transhumanisme/188207>.
- [112] Agrawal, A., Gans, J., & Goldfarb, A. (2018). Prediction Machines: The Simple Economics of Artificial Intelligence. Boston : Harvard Business Review Press.
- [113] OECD Legal Instruments. https://cdn.openai.com/better-language-models/language_models_are_unsupervised_multitask_learners.pdf.
- [114] Anderson, A. R. (1983). Pensée et machine. Paris : Éditions Champ Vallon.
- [115] Andler, D. (2023). **Intelligence artificielle, intelligence humaine : la double énigme**. Paris : Gallimard.
- [116] Barrat, J. (2013). **Our Final Invention: Artificial Intelligence and the End of the Human Era**. New York : St. Martin's Press.

- [117] Bengio, Y., LeCun, Y., & Hinton, G. (2015). **Deep Learning**. *Nature*, 521(7553), 436–444.
- [118] Bolo, J. (1996). **Philosophie contre intelligence artificielle**. Saint-Denis : Éditions Lingua Franca.
- [119] Boss, G. (1987). **Les machines à penser : L'homme et l'ordinateur**. Toulouse : Éditions du Grand Midi.
- [120] Bostrom, N. (2014). **Superintelligence: Paths, Dangers, Strategies**. Oxford : Oxford University Press.
- [121] Brown, T. B., Mann, B., Ryder, N., Subbiah, M., Kaplan, J., Dhariwal, P., ... Amodei, D. (2020). **Language Models are Few-Shot Learners**. *Advances in Neural Information Processing Systems*, 33, 1877–1901.
- [122] Brenet, D. (2024). **L'intelligence artificielle expliquée : des concepts de base aux applications avancées de l'IA**. Paris : Éditions ENI.
- [123] Brynjolfsson, E., & McAfee, A. (2014). **The Second Machine Age: Work, Progress, and Prosperity in a Time of Brilliant Technologies**. New York : W. W. Norton.
- [124] Carr, N. (2010). **The Shallows: What the Internet Is Doing to Our Brains**. New York : W. W. Norton.
- [125] Carr, N. (2014). **The Glass Cage: How Our Computers Are Changing Us**. New York : W. W. Norton.
- [126] Cave, S., & Dignum, V. (2019). **Proceedings of the 2019 AAAI/ACM Conference on AI, Ethics, and Society**.
- [127] Changeux, J.-P. (1983). **L'homme neuronal**. Paris : Odile Jacob.
- [128] Citton, Y. (2014). **Pour une écologie de l'attention**. Paris : Seuil.
- [129] Commission européenne (2018). **Un plan coordonné dans le domaine de l'intelligence artificielle** (COM/2018/795 final). Bruxelles : Commission européenne. <https://eur-lex.europa.eu/legal-content/FR/TXT/>.
- [130] Ethics guidelines for trustworthy AI | Shaping Europe's digital future. <https://digital-strategy.ec.europa.eu/en/library/ethics-guidelines-trustworthy-ai>.
- [131] CNIL (2017, 15 décembre). **Comment permettre à l'Homme de garder la main ? Rapport sur les enjeux éthiques des algorithmes et de l'intelligence**

artificielle. Paris : CNIL. <https://www.cnil.fr/fr/comment-permettre-lhomme-de-garder-la-main-rapport-sur-les-enjeux-ethiques-des-algorithmes-et-de>.

- [132] Damasio, A. (2010). **Self Comes to Mind: Constructing the Conscious Brain.** New York : Pantheon.
- [133] Dehaene, S. (2020). **Apprendre ! Les talents du cerveau, le défi des machines.** Paris : Odile Jacob.
- [134] Dupuy, J.-P. (2009). **Au péril de la science : Les sciences cognitives et les repères philosophiques.** Paris : Seuil.
- [135] Dreyfus, H. L. (1972). **What Computers Can't Do: A Critique of Artificial Reason.** New York : Harper & Row.
- [136] Dreyfus, H. L. (1979). **What Computers Still Can't Do: A Critique of Artificial Reason** (rev. ed.). Cambridge, MA : MIT Press.
- [137] Floridi, L. (2014). **The Fourth Revolution: How the Infosphere is Reshaping Human Reality.** Oxford : Oxford University Press.
- [138] Floridi, L., & Sanders, J. W. (2004). On the morality of artificial agents. **Minds and Machines**, 14(3), 349–379.
- [139] Harari, Y. N. (2014). **Sapiens: A Brief History of Humankind.** New York : Harper.
- [140] Harari, Y. N. (2015). **Homo Deus: A Brief History of Tomorrow.** New York : Harper.
- [141] Harari, Y. N. (2018). **21 Lessons for the 21st Century.** New York : Spiegel & Grau.
- [142] Hofstadter, D. R. (1979). **Gödel, Escher, Bach: An Eternal Golden Braid.** New York : Basic Books.
- [143] Kahneman, D. (2011). **Thinking, Fast and Slow.** New York : Farrar, Straus and Giroux.
- [144] Kahneman, D., & Tversky, A. (1974). Judgment under uncertainty: Heuristics and biases. **Science**, 185(4157), 1124–1131.
- [145] Kelly, K. (2016). **The Inevitable: Understanding the 12 Technological Forces That Will Shape Our Future.** New York : Viking.

- [146] Kurzweil, R. (1990). **The Age of Intelligent Machines**. Cambridge, MA : MIT Press.
- [147] Kurzweil, R. (2005). **The Singularity is Near: When Humans Transcend Biology**. New York : Viking.
- [148] Kurzweil, R. (2012). **How to Create a Mind: The Secret of Human Thought Revealed**. New York : Viking.
- [149] Latour, B. (1987). **Science in Action: How to Follow Engineers and Scientists through Society**. Cambridge, MA : Harvard University Press.
- [150] LeCun, Y., Bengio, Y., & Hinton, G. (2015). **Deep learning**. *Nature*, 521(7553), 436–444.
- [151] LeCun, Y. (2023). **Quand la machine apprend : Les rouages du deep learning**. Paris : Odile Jacob.
- [152] Legg, S., & Hutter, M. (2007). Universal intelligence: A definition of machine intelligence. *Minds and Machines*, 17(4), 391–444.
- [153] Lévy, P. (1994). **L'intelligence collective : Pour une anthropologie du cyberspace**. Paris : La Découverte.
- [154] Lanier, J. (2018). **Dix arguments pour supprimer vos comptes de réseaux sociaux tout de suite** Paris : Éditions du Seuil.
- [155] Marcus, G. (2019). **Rebooting AI: Building Artificial Intelligence We Can Trust**. New York : Pantheon.
- [156] McCarthy, J. (2007). What is Artificial Intelligence? <http://www-formal.stanford.edu/jmc/whatisai/>.
- [157] Minsky, M. (1985). **The Society of Mind**. New York : Simon & Schuster.
- [158] Negroponte, N. (1995). **Being Digital**. New York : Knopf.
- [159] OECD (2019). **Recommendation of the Council on Artificial Intelligence**. Paris : OECD Publishing. Disponible sur : https://legalinstruments.oecd.org/public/doc/27/270649c2-7e69-4f50-bcce-4b8bf08cea78/Ratification_Instrument_OECD-LEGAL-0449_F.pdf.
- [160] Artificial intelligence. <https://www.oecd.org/going-digital/ai/principles/>.
- [161] Pariser, E. (2011). **The Filter Bubble: What the Internet is Hiding from You**. New York : Penguin.

- [162] Penrose, R. (1989). **The Emperor's New Mind: Concerning Computers, Minds and the Laws of Physics**. Oxford : Oxford University Press.
- [163] Penrose, R. (1994). **Shadows of the Mind: A Search for the Missing Science of Consciousness**. Oxford : Oxford University Press.
- [164] OECD Legal Instruments. https://cdn.openai.com/better-language-models/language_models_are_unsupervised_multitask_learners.pdf.
- [165] Russell, S., & Norvig, P. (2010). **Artificial Intelligence: A Modern Approach** (3e ed.). Prentice Hall.
- [166] Russell, S. (2019). **Human Compatible: AI and the Problem of Control**. New York : Viking.
- [167] Searle, J. R. (1980). Minds, brains and programs. **Behavioral and Brain Sciences**, 3(3), 417–424.
- [168] Schmidhuber, J. (2015). Deep learning in neural networks: An overview. **Neural Networks**, 61, 85–117.
- [169] Shannon, C. E. (1950). Programming a computer for playing chess. **Philosophical Magazine**, 41(314), 256–275.
- [170] Simon, H. A. (1996). **The Sciences of the Artificial** (3e ed.). Cambridge, MA : MIT Press.
- [171] Sparrow, B., Liu, J., & Wegner, D. M. (2011). Google effects on memory: Cognitive consequences of having information at our fingertips. **Science**, 333(6043), 776–778.
- [172] Tegmark, M. (2017). **Life 3.0: Being Human in the Age of Artificial Intelligence**. New York : Knopf.
- [173] Turkle, S. (2011). **Alone Together: Why We Expect More from Technology and Less from Each Other**. New York : Basic Books.
- [174] https://unesdoc.unesco.org/ark:/48223/pf0000381137_fre.
- [175] Villani, C. (2018). **Donner un sens à l'intelligence artificielle : Pour une stratégie nationale et européenne**. (Rapport au président de la République). Paris : La Documentation française.
- [176] Wiener, N. (1948). **Cybernetics: Or Control and Communication in the Animal and the Machine**. Paris : Hermann.

- [177] Weizenbaum, J. (1976). **Computer Power and Human Reason: From Judgment to Calculation**. San Francisco : W. H. Freeman.
- [178] World Economic Forum (2024). **Shaping the Future of Learning: The Role of AI in Education 4.0**. Cologny (Genève) : WEF. Disponible sur : <https://www.weforum.org/reports/shaping-the-future-of-learning-the-role-of-ai-in-education-4-0>.
- [179] Comment permettre à l'Homme de garder la main ? Rapport sur les enjeux éthiques des algorithmes et de l'intelligence artificielle. <https://www.cnil.fr/fr/comment-permettre-lhomme-de-garder-la-main-rapport-sur-les-enjeux-ethiques-des-algorithmes-et-de>.
- [180] ORCID Rénald Gesnot. <https://orcid.org/0009-0008-7717-0397>.